

OS: An Overview



Operating Systems
Wenbo Shen

Why are we studying OS?

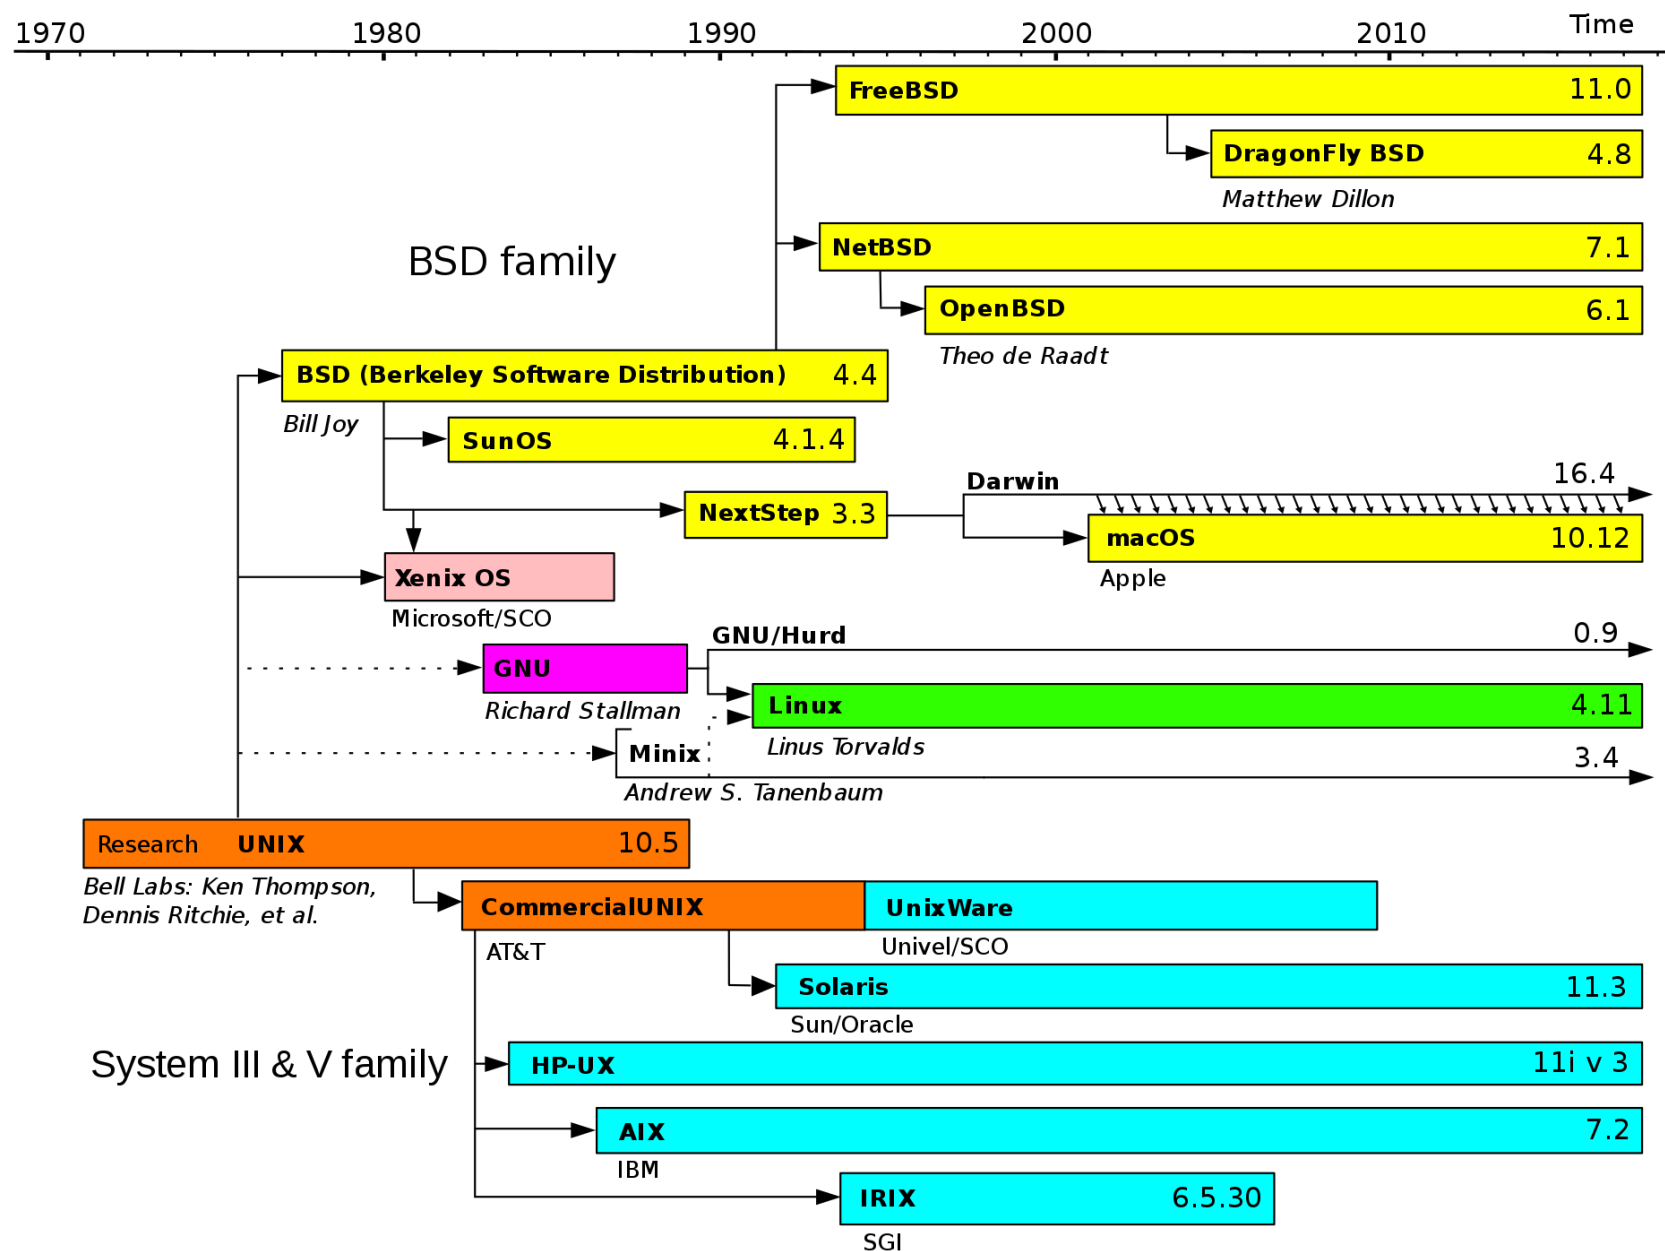
- OS is highly complicated software running on most machines
 - Windows: 50M lines of source code
 - Linux: 25M lines of source code
- It contains many important system concepts
 - complexity hiding, performance tuning, resource allocation...
- Studying OS internals makes you a more **capable** programmer hacker
 - know how it works, and how it works better



Popular Operating Systems



UNIX Family Tree




History of System Software

■ Linus Torvalds

- Released the first Linux prototypes in late 1991 (21 years old)
- Devote his whole life to Linux

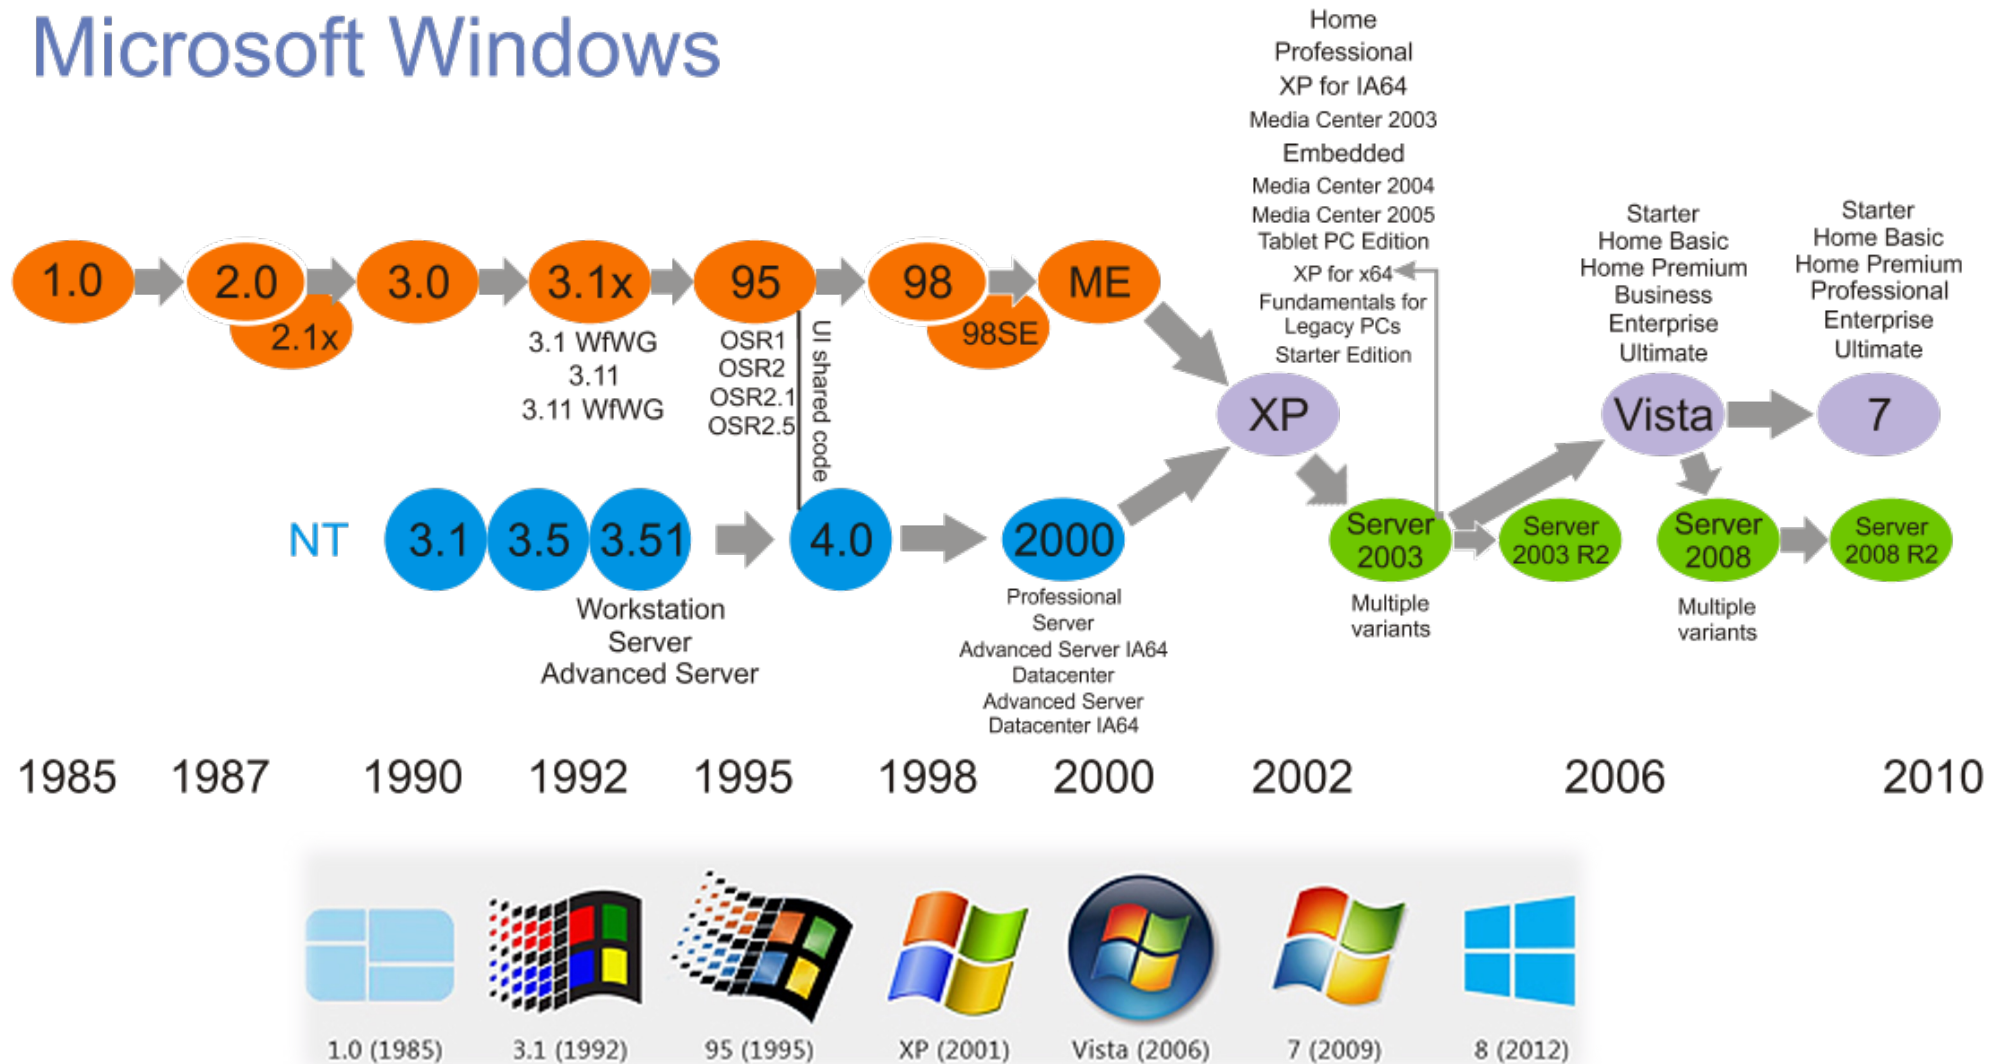
■ Chris Lattner

- Started LLVM project in 2000 at UIUC with his PhD advisor Vikram Adve
- Published numerous papers on program analysis
- Devote his whole life to LLVM and swift (Apple)



Windows Family Tree

Microsoft Windows



Why are we studying OS?

- After all, you probably will not develop an OS

Why are we studying OS?

- ~~After all, you probably will not develop an OS~~
 - Build first mainstream OS from China

Why are we studying OS?

- ~~After all, you probably will not develop an OS~~
 - Build first mainstream OS from China
- OS concepts benefit whole life
 - OS concepts are re-usable when implementing other software
 - Lessons learned from OS study can be applied to complex software systems, such as map-reduce, DNS
- Foundation of ALL software
 - Better user-space software, including apps
 - ▶ Invoke proper kernel API
 - ▶ What can and cannot be done
 - Better performance
 - ▶ Caching
 - ▶ Scheduling
 - ▶ Memory management

Why are we studying OS?

■ Non-textbook answers

- For know-everything-feeling
- For Hacking
 - ▶ The more you know OS, the better hacker you are
 - ▶ Because the thing you are trying to hack into, probably is running an OS



Why are we studying OS?

■ Non-textbook answers

● For Profit

- ▶ Interview = coding + system design
- ▶ Great System == Great Product == Great Company



Google MapReduce

Google File System

Google Cloud



Fuchsia
by Google



ANDROID



Wear OS by Google



Google Chrome OS

Why are we studying this?

- After all, you probably will not develop an OS
 - Unless you land a super interesting job :)
- Important to understand what you use
 - Understanding how the OS works helps you develop (better) apps, understand what you can and cannot do, understand performance, understand why some OS may be better/worse than some other in some situations
- Pervasive abstractions
 - OS concepts are fundamental and re-usable when implementing apps that are not operating systems
- Complex software systems
 - Many of you will participate in complex software systems
 - OSES are among the most interesting such systems and lessons from OSES (and their evolutions) can be applied in many other contexts

Studying OS Today

- Thanks to the open-source movement we have access to a lot of OS code
- Before OSes were even more mysterious
 - We can now look at “old” commercial OSes, which often reveals that they were pretty cool (or pretty scary)
- In fact, it’s become possible for any student to create an OS after reading other OS code
 - Or to contribute to an existing OS
- And thanks to virtualization technology, one can play with and run OSes easily
 - Without compromising one’s computer

This Set of Lecture Notes

- This set of lecture notes is a 10,000 ft overview of the OS
- Many details will be explained throughout the semester
- Some terms are used, which you may not be familiar with, and that will all be explained later
- Some simplifying assumptions are made
 - If you know better, then bear with us until a further lecture

What is an OS?

- What do you think the answer is?

(there are many possible answers)

What is an OS?

- What do you think the answer is?

When buying a phone


<

When using a phone



What is an OS?

- One answer: software layer between the applications and the hardware because the hardware would be too difficult for users to use



- Or: it's “all the code you didn't have to write” when you wrote your application
 - Not quite right as there are tons of non-OS libraries that you didn't write as well


What is an OS?

- It's a resource **abstractor** and a resource **allocator**
 - The OS defines a set of logical resources that correspond to hardware resources, and a set of well-defined operations on logical resources
 - ▶ e.g., physical resources: CPU, Disks, RAM
 - ▶ e.g., logical resources: processes, files, arrays
 - The OS decides who (which running program) gets what resource (share) and when
- Popular yet very wrong definition: It's the one program (in fact, its “kernel” component) that runs at all times
 - It's a misleading view of the kernel (which is not a running program at all)

How big is an OS?

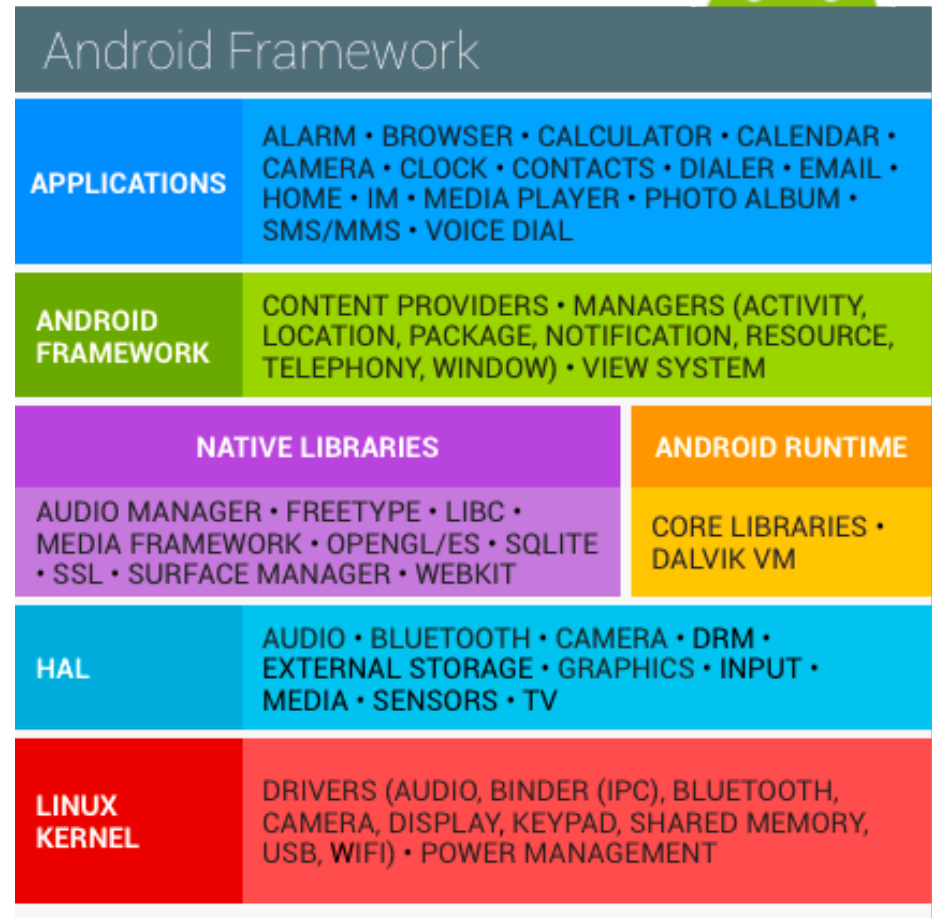

- The question “What is part of the OS and what isn’t?” is a difficult one
 - What about the windowing system? “system” programs?
 - The 1998 lawsuit against Microsoft putting “too much” in what they called the Operating System
- But here are a few SLOC (Source Line of Code) numbers
 - Windows NT (1993): 6 Million
 - Windows XP: ~50 Million
 - Windows Vista: ~XP + 10
 - Mac OS X 10.4: ~86 Million
 - Ubuntu distribution: > 230 Million
 - ▶ But tons of things are not part of the OS
 - ▶ Kernel 2.6.29: 11 Million
- OSes are BIG

Linux Kernel Lines of Code



■ <https://www.linuxcounter.net/statistics/kernel>

OS is complex



How does one start an OS?

- When a computer boots, it needs to run a first program: the **bootstrap program**
 - Stored in Read Only Memory (ROM)
 - Called the “firmware” or **bootloader**
 - We use OpenSBI (Supervisor Binary Interface) for our labs
- The bootstrap program initializes the computer
 - Register content, device controller contents, etc.
- It then locates and loads the OS kernel into memory
- The kernel starts the first process (called “init” on Linux, “launchd” on Mac OS X).. let's see it...
- And then, nothing happens until an **event** occurs
 - more on events in a few slides


Multi-Programming

- **Multi-Programming**: Modern OSes allow multiple “jobs” (running programs) to reside in memory simultaneously
 - The OS picks and begins to execute one of the jobs in memory
 - When the job has to wait for “something”, then the OS picks another job to run
 - This is called a **context-switch**, and improves productivity
- We are used to this now, but it wasn’t always so
 - Single-user mode
 - ▶ Terrible productivity (while you “think”, nobody else is using the machine)
 - Batch processing (jobs in a queue)
 - ▶ Low productivity (CPU idle during I/O operations)

Time-Sharing


- **Time-Sharing**: Multi-programming with rapid context-switching
- Jobs cannot run for “too long”
- Allows for interactivity
 - Response time very short
 - Each job has the illusion that it is alone on the system
- In modern OSes, jobs are called **processes**
 - **A process is a running program**
- There are many processes, some of which are (partly) in memory concurrently
 - Let's run the “ps aux” command on my laptop

The Running OS



- The code of the operating system resides in memory at a specified address, as loaded by the bootstrap program
- At times, some of this code can be executed by a process
 - Branch to some OS code segment
 - Return to the program's code later
- Each process is loaded in a subset of the memory
 - Code + data
- Memory protection among processes is ensured by the OS
 - A process cannot step on another process' toes

Running the OS Code?



- The kernel is NOT a running job
- It's code (i.e., a data and a text segment) that resides in memory and is ready to be executed at any moment
 - When some event occurs
- It can be executed on behalf of a job whenever requested
- It can do special/dangerous things
 - having to do with hardware

A Note on Kernel Size

- In the previous figure you see that the kernel uses some space in the physical memory
- As a kernel designer you want to be careful to not use too much memory!
 - Hence the fight about whether new features are truly necessary in the kernel
 - Hence the need to write lean/mean code
 - **lean** → nothing more; **mean** → single-minded
- Furthermore, there is no memory protection within the kernel
 - The kernel's the one saying to a process "segmentation fault"
 - Nobody's watching over the kernel
- So one must be extremely careful when developing kernels
- Hence the reason why kernel "hacking" is highly respected

How OS achieves ...

- It's a resource abstractor and a resource allocator
 - The OS defines a set of logical resources that correspond to hardware resources, and a set of well-defined operations on logical resources
 - ▶ e.g., physical resources: CPU, Disks, RAM
 - ▶ e.g., logical resources: processes, files, arrays
 - The OS decides who (which running program) gets what resource (share) and when

- How to achieve?

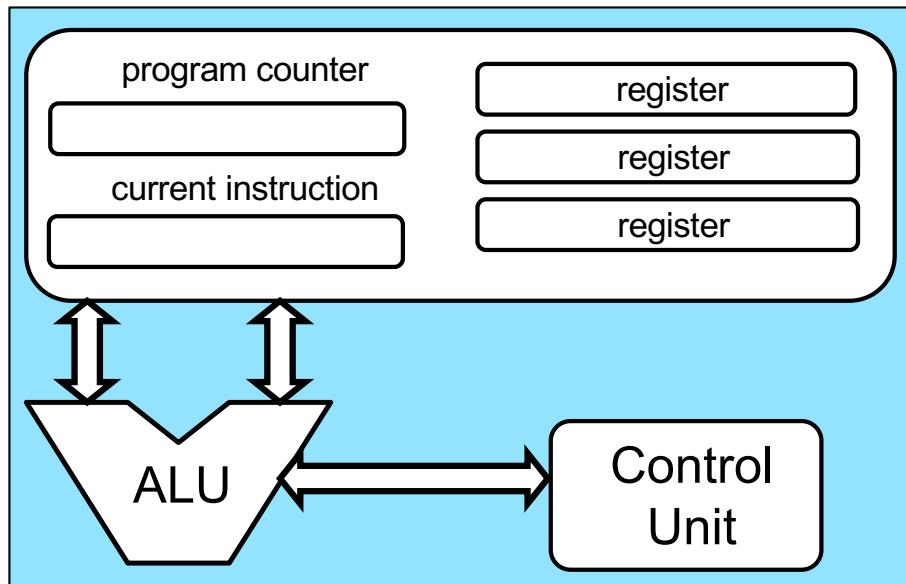
Protected Instructions

- A subset of instructions of every CPU is restricted in usage: only the OS can execute them
 - Known as **protected** (or **privileged**) instructions
- For instance, only the OS can:
 - Directly access I/O devices (printer, disk, etc.)
 - ▶ Fairness, security
 - Manipulate memory management state
 - ▶ Fairness, security
 - Manipulate protected control registers
 - ▶ Kernel mode, interrupt level (more on all this later)
 - Execute the halt instruction that shuts down the processor
- The CPU needs to know whether it can execute a protected instruction or not...
 - How to achieve?

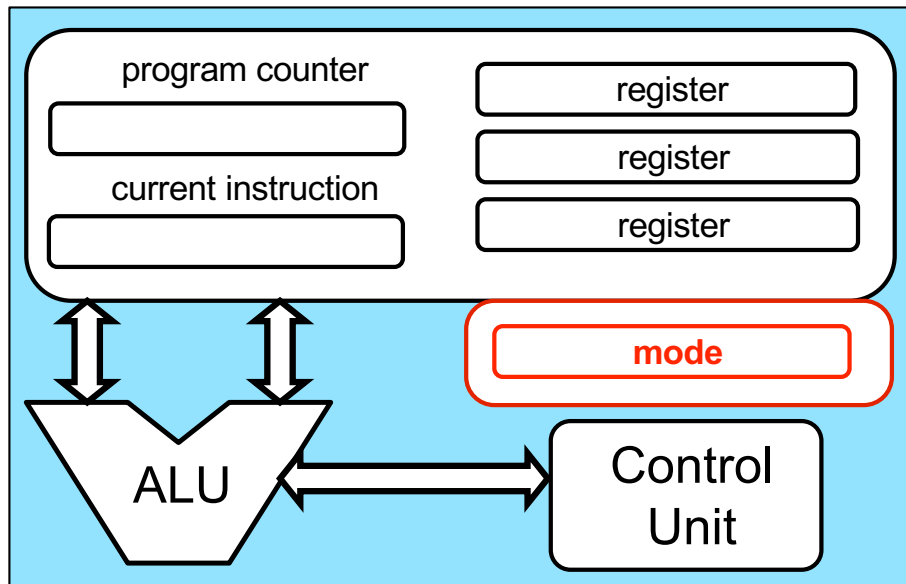
User vs. Kernel Mode

- All modern processors support (at least) two modes of execution:
 - **Unprivileged mode**: In this mode protected instructions cannot be executed
 - **Privileged mode**: In this mode all instructions can be executed
- The mode is indicated by a status bit in a protected control register
 - The CPU checks this bit before executing a protected instruction
 - User code (unprotected insn) executes in user mode
 - OS code (all insn) executes in kernel mode
- Setting the mode bit is, of course, a protected instruction

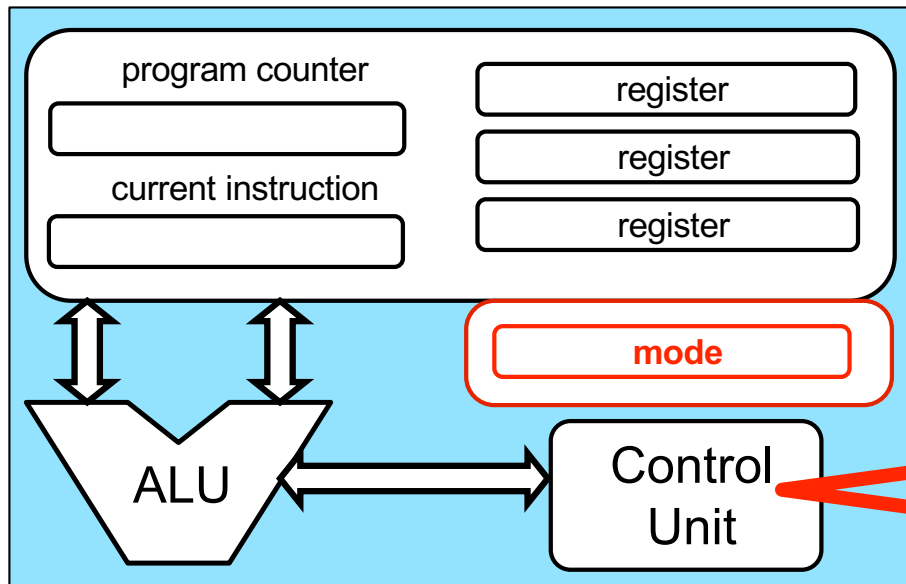
User vs. Kernel Mode



User vs. Kernel Mode



User vs. Kernel Mode



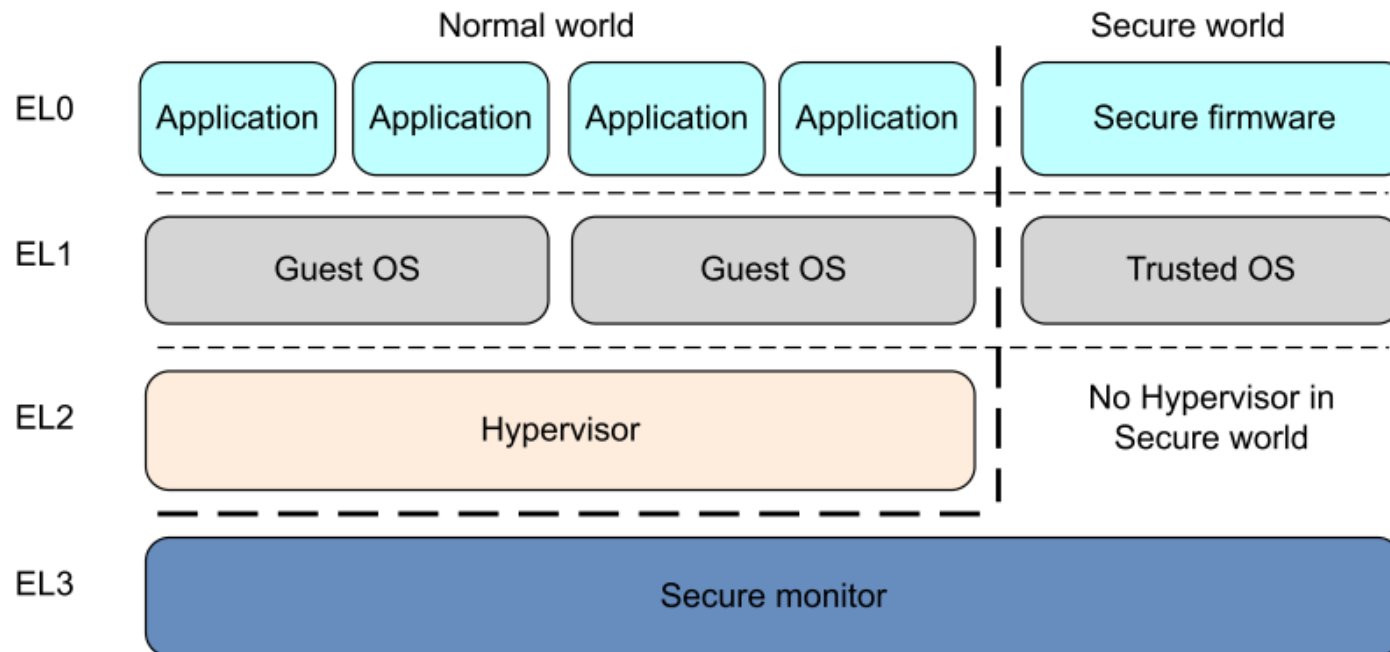
- Decode instruction
- Determine if instruction is privileged or not
 - Based on the instruction code (e.g., the binary code for all privileged instructions could start with '00')
- If instruction is privileged and mode == user, then abort!
 - Raise a "trap"

User vs. Kernel Mode

- There can be multiple modes
 - e.g., multiple levels in CPU
- MS-DOS had only one mode, because it was written for the Intel 8088 (1979), which had no mode bit
 - A user program could wipe out the whole system due to a bug (or a malicious user)
 - Multiple user programs could write to the same device concurrently, leading to incoherent behavior

ARM64 Modes

- There can be multiple modes
 - e.g., multiple levels in ARM64 (2011)



RISCV Modes

Level	Encoding	Name	Abbreviation
0	00	User/Application	U
1	01	Supervisor	S
2	10	<i>Reserved</i>	
3	11	Machine	M

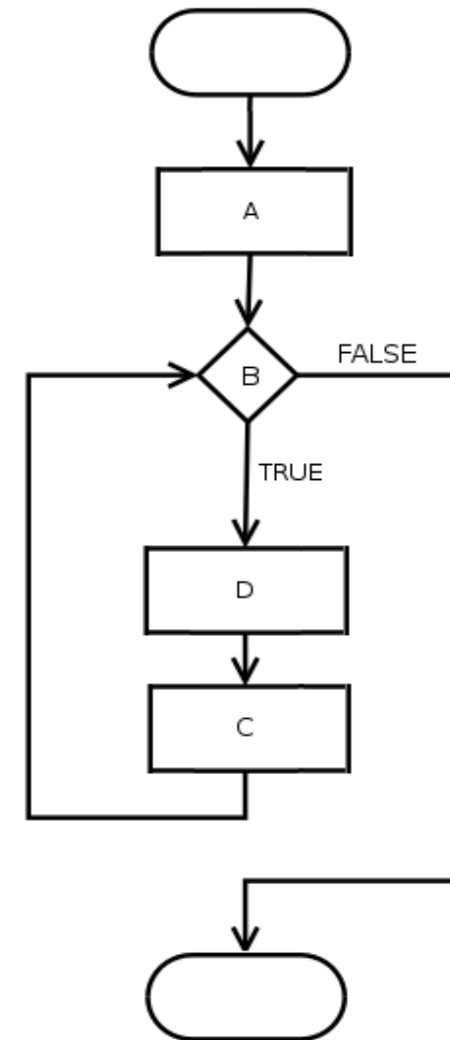
OS Events

- An event is an “unusual” change in control flow
 - A usual change is some “branch” instruction within a user program for instance
- An event stops execution, changes mode, and changes context
 - i.e., it starts running kernel code
- The kernel defines a **handler** for each event type
 - i.e., a piece of code executed in kernel mode
- Once the system is booted, all entries to the kernel occur as the result of an event
 - The OS can be seen as a huge event handler

Control flow

- In computer science, control flow (or flow of control) is the order in which individual statements, instructions or function calls of an program are executed.

```
for(A;B;C)  
D;
```



10K-foot View of Kernel Code

```
void processEvent(event) {  
    switch (event.type) {  
        case NETWORK_COMMUNICATION:  
            NetworkManager.handleEvent(event);  
            break;  
  
        case SEGMENTATION_FAULT:  
        case INVALID_MODE:  
            ProcessManager.handleEvent(event);  
            break;  
        ...  
    }  
    return;  
}
```

OS Events

- There are two kinds of events: **interrupts** and **exceptions (traps)**
 - The two terms are often confused (even in the textbook)
 - The term fault often refers to unexpected events
- Interrupts are caused by external events
 - Hardware-generated
 - e.g., some device controller says “something happened”
- Exceptions are caused by executing instructions
 - Software-generated interrupts
 - e.g., the CPU tried to execute a privileged instruction but it's not in kernel mode (Traps)
 - e.g., a division by zero (Aborts)

OS Events

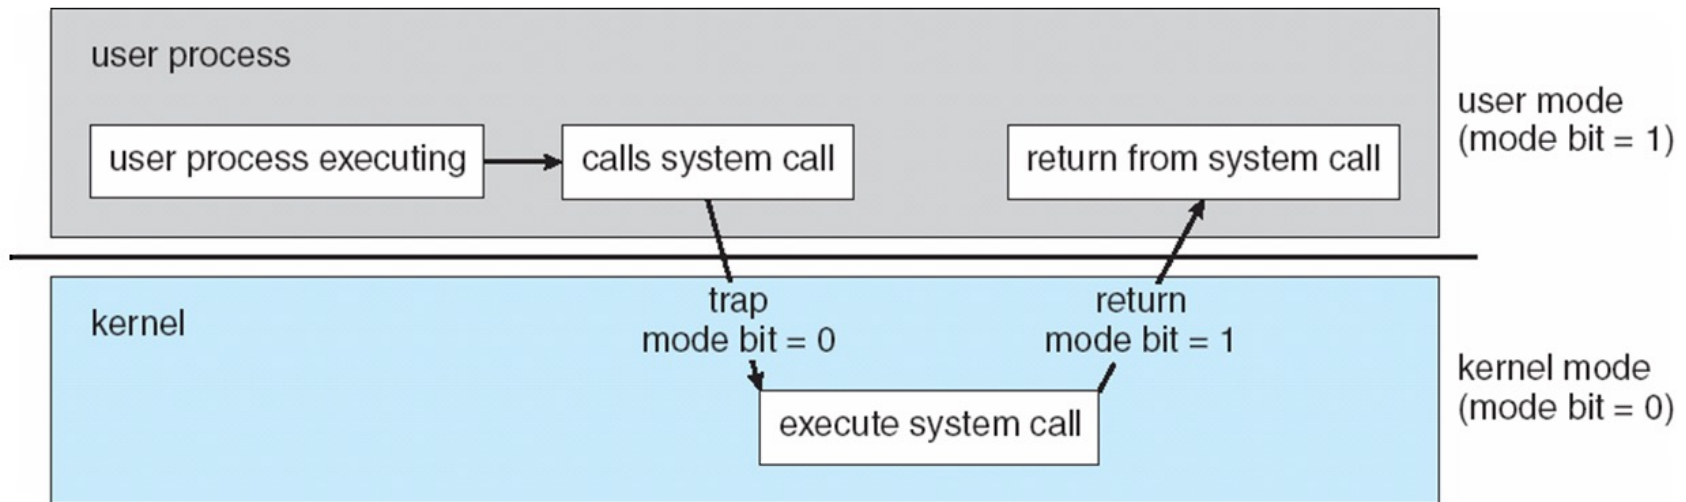
- When the CPU is interrupted, it stops what it is doing and immediately transfers execution to a fixed location in the kernel code
 - the “processEvent()” method in my mock-up kernel code a couple slides ago
- Could result in:
 - Some work being done by the kernel
 - A user process being terminated (e.g., segmentation fault)
- What about “faults” in the kernel?
 - Say dereferencing of a NULL pointer, or a divide by zero
 - This is a fatal fault
 - UNIX Panic, Windows blue screen of death
 - ▶ Kernel is halted, state dumped to a core file, machine is locked up

System Calls

- When a user program needs to do something privileged, it calls a **system call**
 - e.g., to create a process, write to disk, read from the network card
- **A system call is a special kind of trap**
- Every Instruction Set Architecture (ISA) provides a system call instruction that
 - Causes a trap, which maps to a kernel handler
 - Passes a parameter determining which system call to use (a number)
 - Saves caller state (PC, regs, mode) so it can be restored later
- On the x86-64 architecture the instruction is called `syscall`

```
mov    $0x1, $eax,  
syscall           // call system call #1
```

System Calls



10K-foot View of Kernel Code

```
void processEvent(event) {  
    switch (event.type) {  
    case NETWORK_COMMUNICATION:  
        NetworkManager.handleEvent(event);  
        break;  
    case SEGMENTATION_FAULT:  
    case INVALID_MODE:  
        ProcessManager.handleEvent(event);  
        break;  
    case SYSTEM_CALL:  
        SystemCallManager.execute(event);  
        break;  
        ...  
    }  
    return;  
}
```

Timers

- The OS must keep control of the CPU
 - OS must have a concept of “time”
 - Programs cannot gain an unfair share of the computer
- One way in which the OS (or kernel) retrieves control is when an interrupt occurs
- To make sure that an interrupt will occur reasonably soon, we can use a **timer**
- The timer interrupts the computer regularly
 - For example, every 1ms-1s
 - The OS always makes sure the timer is set before turning over control to user code
- Modifying the timer is done via privileged instructions

10K-foot View of Kernel Code

```
void processEvent(event) {  
    Timer.set(1000); // Will generate an event in 1000 time units  
    switch (event.type) {  
    case NETWORK_COMMUNICATION:  
        NetworkManager.handleEvent(event);  
        break;  
  
    case SEGMENTATION_FAULT:  
    case INVALID_MODE:  
        ProcessManager.handleEvent(event);  
        break;
```

10K-foot View of Kernel Code

```
case SYSTEM_CALL:  
    SystemCallManager.execute(event);  
    break;
```

```
case TIMER:  
    Timer.handleEvent(event);  
    break;
```

```
...
```

```
}
```

```
return;
```

```
}
```

Main OS Services

- Process Management
- Memory Management
- Storage Management
- I/O Management
- Protection and Security

Process Management

- A **process** is a program in execution
 - Program: passive entity
 - Process: active entity

- The OS is responsible for :
 - Creating and deleting processes
 - Suspending and resuming processes
 - Providing mechanisms for process synchronization
 - Providing mechanisms for process communication
 - Providing mechanisms for deadlock handling

Memory Management

- Memory management determines what is in memory when
 - The kernel is ALWAYS in memory
- The OS is responsible for:
 - Keeping track of which parts of memory are currently being used and by which process
 - Deciding which processes (or parts thereof) and data to move into and out of memory
 - Allocating and deallocating memory space as needed
- The OS is not responsible for memory caching, cache coherency, etc.
 - These are managed by the hardware

Storage Management

- The OS provides a uniform, logical view of information storage
 - It abstracts physical properties to logical storage unit (e.g., as a “file”)
- The OS operates File-System management
 - Creating and deleting files and directories
 - Manipulating files and directories
 - Mapping files onto secondary storage
 - Backup files onto stable (non-volatile) storage media
 - Free-space management
 - Storage allocation
 - Disk scheduling

I/O Management

- The OS hides peculiarities of hardware devices from the user
- The OS is responsible for
 - Memory management of I/O including buffering (storing data temporarily while it is being transferred), spooling (the overlapping of output of one job with input of other jobs), etc.
 - General device-driver interface
 - ▶ So that multiple devices can be used with the same kernel as long as they offer some standard interface
 - Drivers for specific hardware devices

Protection and Security

- **Protection**: mechanisms for controlling access of processes to resources defined by the OS
- **Security**: defense of the system against internal and external attacks, for example, due to bugs
 - including denial-of-service, worms, viruses, identity theft, theft of service
- The OS provides:
 - Memory protection, device protection
 - User IDs associated to processes and files
 - Group IDs for sets of users
 - Definition of privilege levels

Privileged Instructions

- In class discussion: which of these instructions should be privileged, and why?
 - Set value of the system timer
 - Read the clock
 - Clear memory
 - Issue a system call instruction
 - Turn off interrupts
 - Modify entries in device-status table
 - Access I/O device

Privileged Instructions

- In class discussion: which of these instructions should be privileged, and why?
 - Set value of the system timer
 - Read the clock
 - Clear memory
 - Issue a system call instruction
 - Turn off interrupts
 - Modify entries in device-status table
 - Access I/O device

Sections 1.10 and 1.11

- Operating System Concepts (10th Edition)
- The textbook has sections on “Computing environments” and “Free and Open-Source Operating Systems”
 - Make sure you read them!
 - They talk about a number of topics that are part of the general culture that you should have, in case you don’t have it already

Conclusion

- This set of slides gave a grand tour of what an OS is and what it does
- We have purposely left many elements not fully explained... they will be elucidated throughout the semester
- Reading assignment: Chapter 1