

# Threads




**Operating Systems**  
**Wenbo Shen**

# Revisit - Process Concept

---


- Process =
  - **code** (also called the **text**)
    - initially stored on disk in an executable file
  - a **data section**
    - global variables (.bss and .data in x86 assembly)
  - **program counter**
    - points to the next instruction to execute (i.e., an address in the code)
  - content of the processor's **registers**
  - a runtime **stack**
  - a **heap**
    - for dynamically allocated memory (malloc, new, etc.)

# Revisit - Memory Layout of a C Program



# Why thread?

- How can we make a process run faster
  - Multiple execution units with a process




# Thread Definition

---

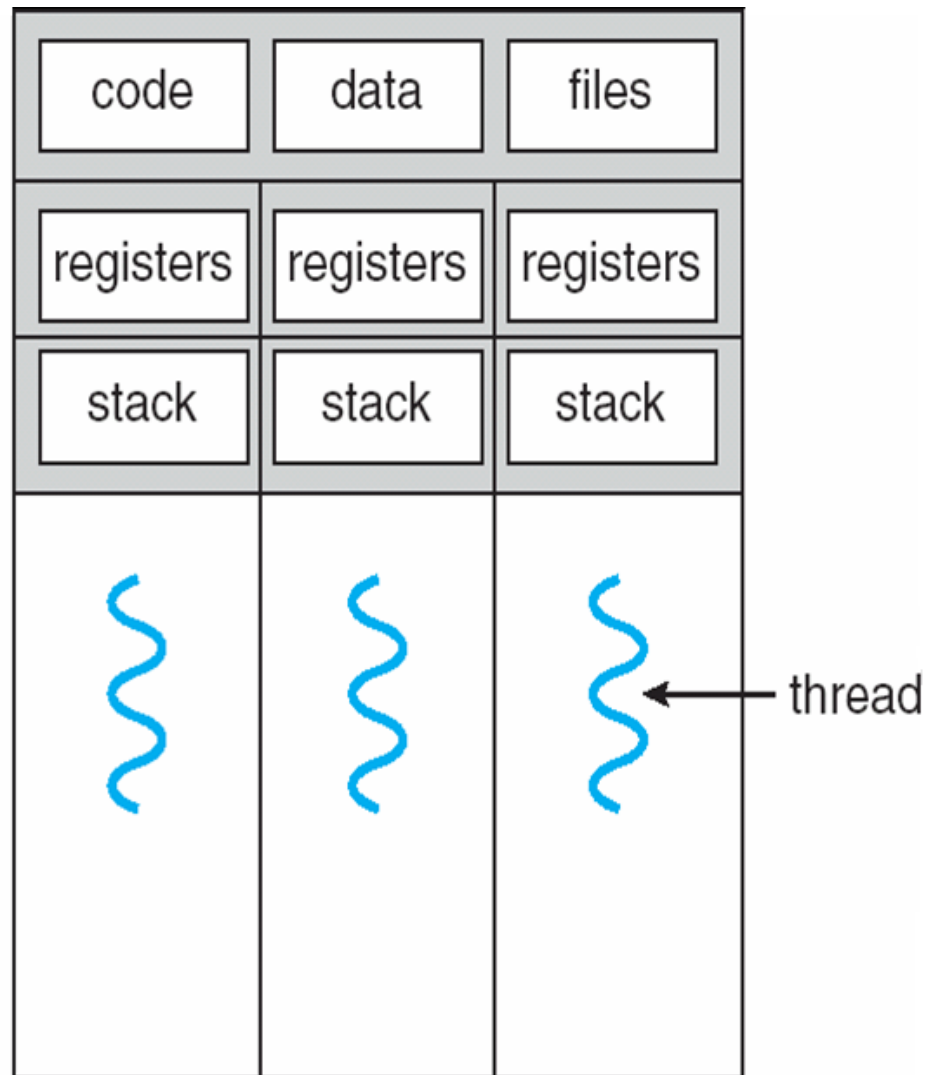
- A thread is a basic unit of execution within a process
- Each thread has its own
  - thread ID
  - program counter
  - register set
  - Stack
- It shares the following with other threads within the same process
  - code section
  - data section
  - the heap (dynamically allocated memory)
  - open files and signals
- **Concurrency:** A multi-threaded process can do multiple things at once

# The Typical Figure

---

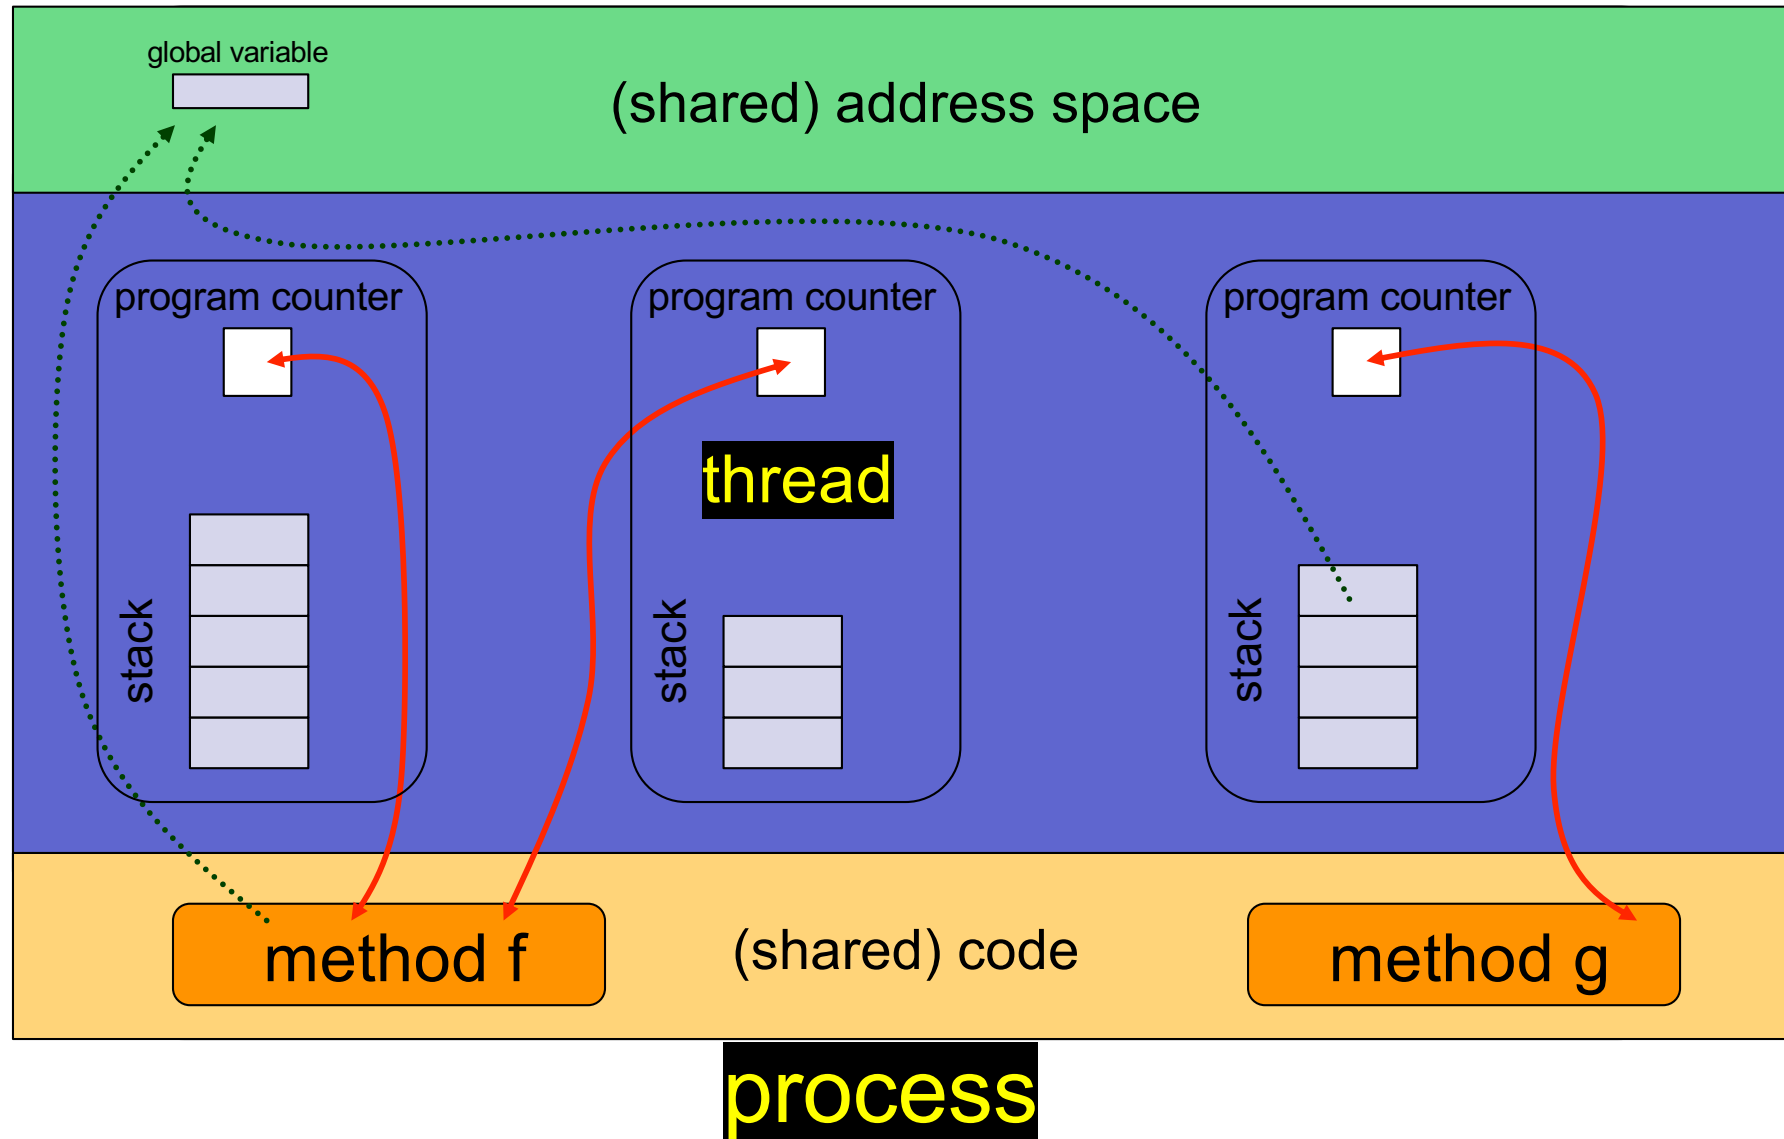


single-threaded process



multithreaded process

# A More Detailed Figure



# Multi-Threaded Program

- Source-code view

- a blue thread
- a red thread
- a green thread

[illegible]

```

Terminal — vim — 101x71

A
# print_queue()
#
void print_queue(xbt_fifo_t q) {
    xbt_fifo_tmap_t tmap;
    job_descriptor_t descriptor;

    fprintf(stderr, "---- Job queue size: %i\n", xbt_fifo_size(q));
    xbt_fifo_foreach(q, &descriptor, &descriptor_t {
        fprintf(stderr, "Job descriptor: %i, duration: %i, actual_duration: %i\n",
            descriptor->requested_duration, descriptor->actual_duration);
    });
    printf(stderr, "\n");
}

# start_job()
#
void start_job(job_descriptor_t jd, scheduler_bookkeeping_t *bk)
{
    /* Get the time = NSG_get_clock(); */
    /* Notify the job simulator of a new job */
    NSG_task_create(&job_started, &b, (void*)jd);
    task NSG_task_put(task, NSG_host_set(&f), PORT_START_JOB) = NSG_OK; <
    xbt_assert(0, "Error while sending a job start notification to the job simulator");
}

# sub_iter()
#
/* Notify the subiter */
{
    s_task t task;
    int i = jd->id;
    task NSG_task_create(&job_started, &b, (void*)jd);
    task NSG_task_put(task, jd->subiter, jd->channel_started) = NSG_OK; <
    xbt_assert(0, "Error while sending a job start notification to the subiter");
}

# sub_iter_size()
#
/* Notify the subiter of my queue size */
/* (this should really be done anywhere) */
{
    s_task t task;
    int queue_size = xbt_fifo_size(&queue);
    task NSG_task_create(&sub_iter_size, &b, (void*)queue_size);
    if (NSG_task_put(task, jd->subiter, jd->channel_queue_size) = NSG_OK) <
    xbt_assert(0, "Error while sending a queue size notification to the subiter");
}

return 0;
}

# scheduler_init()
#
void scheduler_init(job_scheduling_algorithm_t alg, scheduler_bookkeeping_t *bk)
{
    /* Initialize the job descriptor queues and the number of free nodes */
    bk->queue = xbt_fifo_new();
    bk->running = xbt_fifo_new();
}

# scheduler_finalize()
#
void scheduler_finalize(scheduler_bookkeeping_t *bk)
{
    case POFB;
    scheduler_init(&fbs(bk);
}

```

```

Terminal -- vim -- 98x69
                                #num_target_schedulers,num_pruned_schedulers,
                                num_target_schedulers, target_schedulers);
}

    break;
}

if (num_target_schedulers > 1)
    pending_job->flooded_across = 1;

free(pruned_schedulers);
return;
}

/*
 * find_Rc_target_schedulers()
 */
void find_Rc_target_schedulers(int x, scheduler_info_t preferred,
                               scheduler_info_t *schedulers, int num_schedulers,
                               int *num_target_schedulers, scheduler_info_t **target_schedulers)
{
    int i;
    int n;
    done;

    num_target_schedulers = 0;

    while (QMIN(x,num_schedulers) > 0) {
        for (i = 0; /* first, pick the preferred scheduler */
             i < num_schedulers; i++) {
            if (schedulers[i] == preferred) {
                n++;
                break;
            }
        }
        if (i == num_schedulers)
            continue;
        n++;
        r = random_integer_biased(0,num_schedulers-1);
        n = random_integer(0,num_schedulers-1);
        done = 1;
        for (i = j; i < n; i++)
            if (target_schedulers[i] == schedulers[r])
                done = 0;
        if (done)
            break;
    }
    num_target_schedulers++;
    target_schedulers = REALLLOC(*target_schedulers,
                                (num_target_schedulers) * sizeof(scheduler_info_t));
    target_schedulers[num_target_schedulers-1] = schedulers[r];
}

return;
}

/*
 * find_Sc_target_schedulers()
 */
void find_Sc_target_schedulers(int x, scheduler_info_t preferred,
                               scheduler_info_t *schedulers, int num_schedulers,
                               int *num_target_schedulers, scheduler_info_t **target_schedulers)
{
    scheduler_info_t *sorted;

    sorted = malloc(sizeof(scheduler_info_t) * num_schedulers);
    if (sorted == NULL)
        return;
    for (i = 0; i < num_schedulers; i++)
        sorted[i] = schedulers[i];
    qsort(sorted, num_schedulers, sizeof(scheduler_info_t),
          (int (*)(const void *, const void *)) scheduler_info_cmp);
    num_target_schedulers = 0;
    for (i = 0; i < num_schedulers; i++)
        if (sorted[i] == preferred)
            num_target_schedulers++;
    if (num_target_schedulers == 0)
        return;
    target_schedulers = REALLLOC(*target_schedulers,
                                num_target_schedulers * sizeof(scheduler_info_t));
    for (i = 0; i < num_target_schedulers; i++)
        target_schedulers[i] = sorted[i];
}
}

```

```

vim - vim - 96x42
#include "xltutil.h"
#include "stdlib.h"

XBT_LOG_NEW_DEFAULT_CATEGORY(RECEIVER, "Logging for the receiver process");

/** Receiver function
 * @arg #1: port
 * @arg #2: dynar name
 */
int receiver(int argc, char *argv[])
{
    int port = 1;
    xbt_dict_t fifo;
    dynar_t dynar;

    /* Process the first argument */
    if (sizeof(argv[1]) > 0) {
        xbt_dict_set1(fifo, "invalid port 'x" argv[1]');
    }

    /* Get the dynar */
    if ((dynar = (dynar_t) xbt_dict_get_or_null(receiver_dynar_dict, (const char *) argv[2])) == NULL) {
        xbt_dict_set1(fifo, "cannot find dynar 'x" argv[2]');
    }

    /* Loop */
    while (1) {
        int i;
        m_t task = NULL;

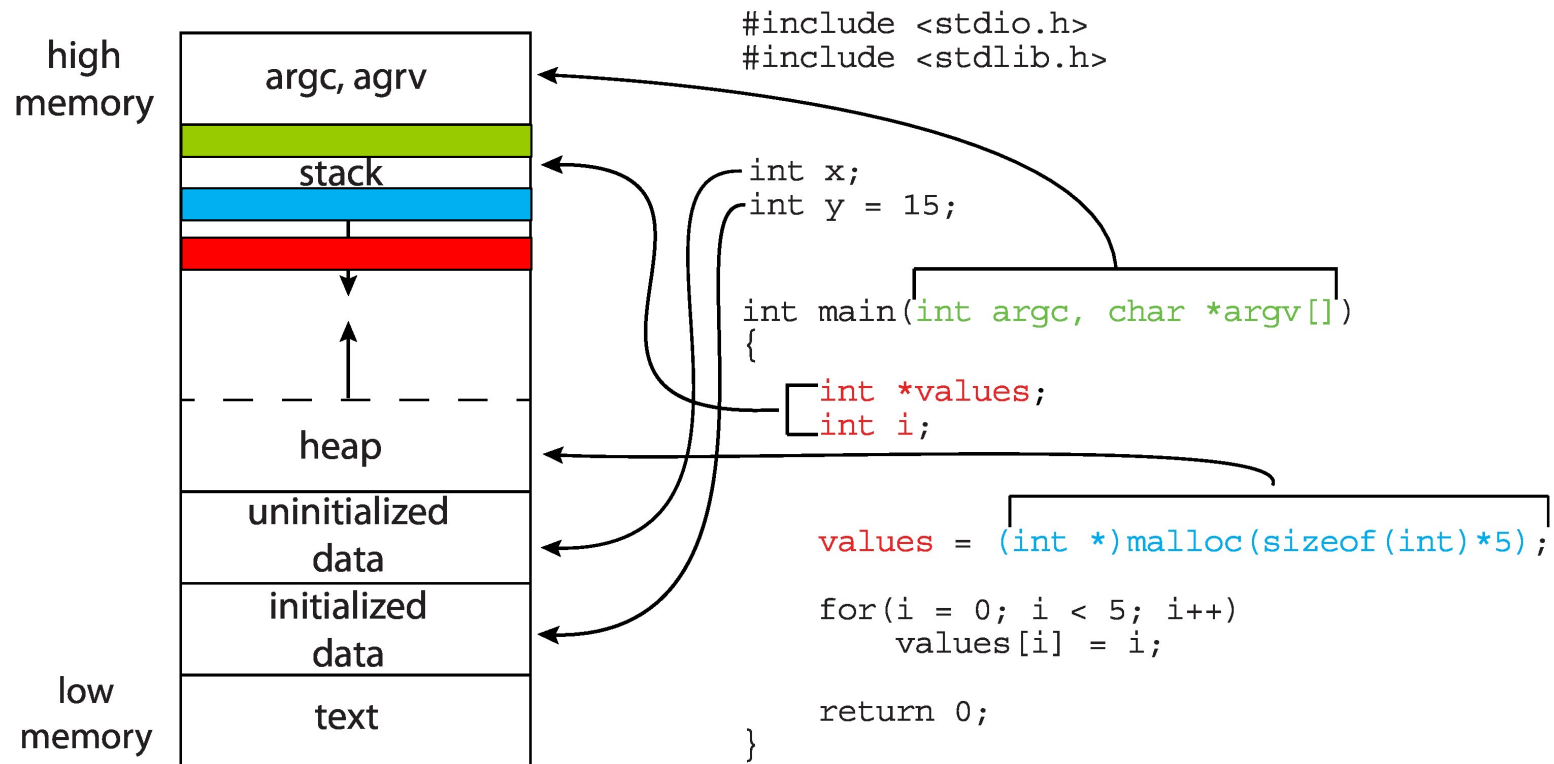
        if (sizeof(task) > 0) {
            xbt_dict_set1(fifo, "Error while receiving a task in a receiver process");
        }

        /* Get the task at a random location in the dynar */
        i = random_integer(0, MAX((0, dynar_length(dynar) - 1));
        dynar_index_t task_index = (dynar_index_t) (dynar_length(dynar) * (double) i / MAX((0, dynar_length(dynar) - 1));
        return 0;
    }
}

```



# Revisit - Memory Layout of a C Program



# Advantages of Threads?

---

## ■ Economy:

- Creating a thread is cheap
  - Much cheaper than creating a process
    - Code, data and heap are already in memory
- Context-switching between threads is cheap
  - Much cheaper than between processes
    - No cache flush

## ■ Resource Sharing:

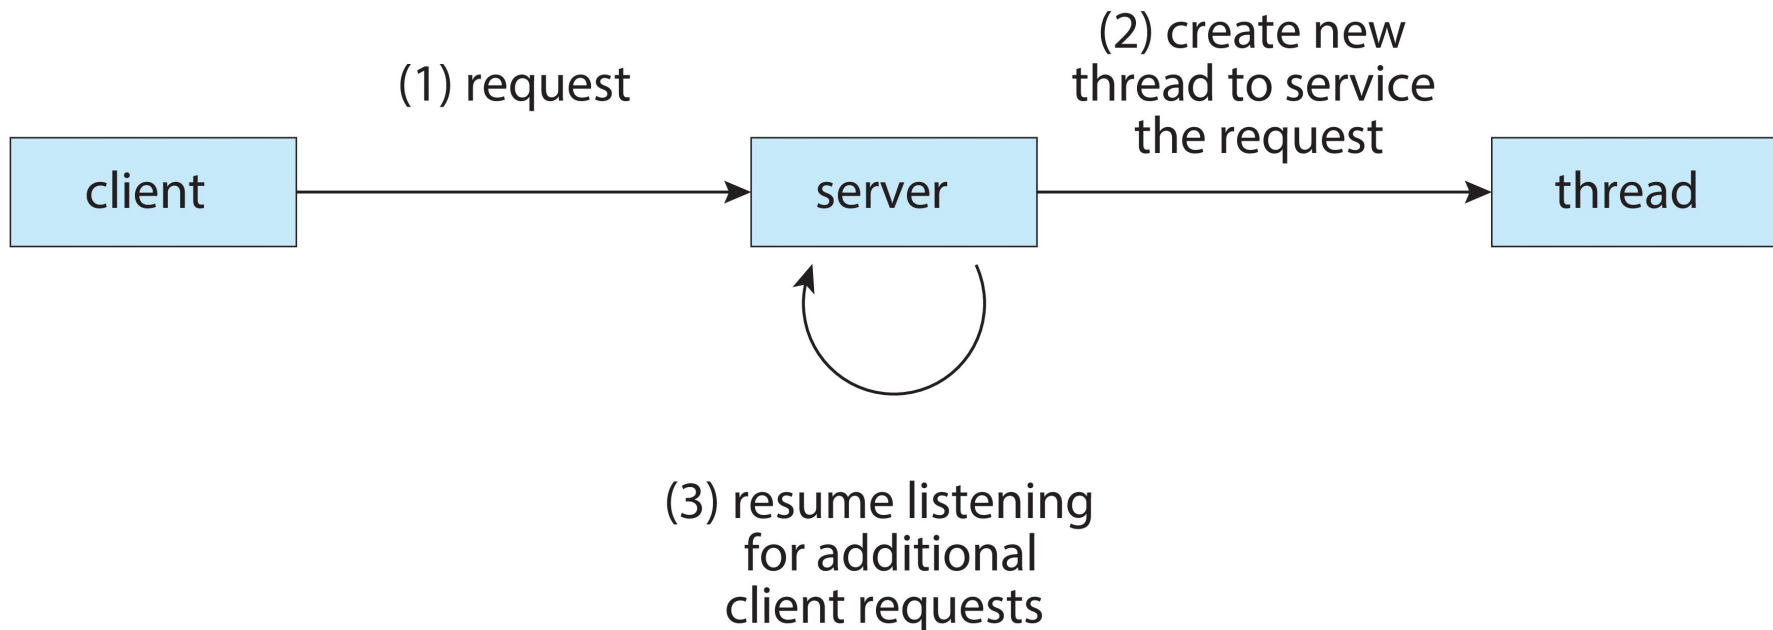
- Threads naturally share memory
  - With processes you have to use possibly complicated IPC (e.g., Shared Memory Segments)
  - **IPC is not needed**
- Having concurrent activities in the same address space is very powerful
  - But fraught with danger

# Advantages of Threads?

---

## ■ Responsiveness

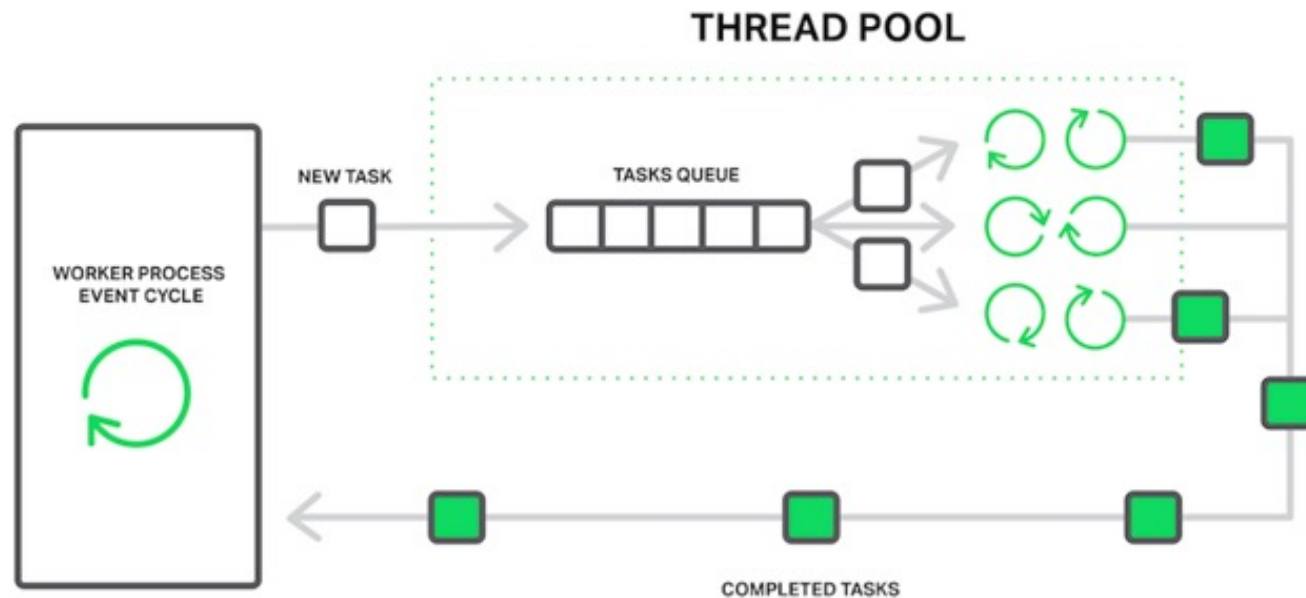
- A program that has concurrent activities is more responsive
  - While one thread blocks waiting for some event, another can do something
  - e.g. Spawn a thread to answer a client request in a client-server implementation
- This is true of processes as well, but with threads we have better sharing and economy



# Advantages of Threads?

---

- Thread Pools in NGINX Boost Performance 9x
  - nginx : worker process -> thread pool



# Advantages of Threads?

---

## ■ Responsiveness

- A program that has concurrent activities is more responsive
  - While one thread blocks waiting for some event, another can do something
  - e.g. Spawn a thread to answer a client request in a client-server implementation
- This is true of processes as well, but with threads we have better sharing and economy

## ■ Scalability

- Running multiple “threads” at once uses the machine more effectively
  - e.g., **on a multi-core machine**
- This is true of processes as well, but with threads we have better sharing and economy

# Drawbacks of Threads

---

- Weak isolation between threads: If **one thread fails** (e.g., a segfault), then **the process fails**
  - And therefore the whole program
- This leads to process-based concurrency
  - e.g., The Google Chrome Web browser
  - See <http://www.google.com/googlebooks/chrome/>
- Sort of a throwback to the pre-thread era
  - Threads have been available for 20+ years
  - Very trendy recently due to multi-core architectures

# Drawbacks of Threads

---

- Threads may be more memory-constrained than processes
  - Due to OS limitation of the address space size of a single process
  - Not a problem any more on 64-bit architecture
- Threads do not benefit from memory protection
  - Concurrent programming with threads is hard
    - But so is it with Processes and Shared Memory Segments

# Threads on My Machine?

---

- Let's run `ps aux` and look at several applications
  - `ps -eLf`
  - Firefox
  - Terminal
  - ...



# Multi-Threading Challenges

---

- Typical challenges of multi-threaded programming
  - Deal with data dependency and synchronization
  - Dividing activities among threads
  - Balancing load among threads
  - Split data among threads
  - Testing and debugging

# User Threads vs. Kernel Threads


---

- Threads can be supported solely in User Space
  - Threads are managed by some user-level thread library (e.g., Java Green Threads)
- Threads can also be supported in Kernel Space
  - The kernel has data structure and functionality to deal with threads
  - Most modern OSes support kernel threads
    - ▶ In fact, Linux doesn't really make a difference between processes and threads (same data structure)

# Many-to-One Model


---

- Advantage: multi-threading is efficient and low-overhead
  - No syscalls to the kernel
- **Major Drawback #1:** cannot take advantage of a multi-core architecture!
- **Major Drawback #2:** if one threads blocks, then all the others do!
- Examples (User-level Threads):
  - Java Green Threads
  - GNU Portable Threads



# One-to-One Model

---




- Removes both drawbacks of the Many-to-One Model
- Creating a new threads requires work by the kernel
  - Not as fast as in the Many-to-One Model
- Example:
  - Linux
  - Windows
  - Solaris 9 and later

# Many-to-Many Model


---

- A compromise
- If a user thread blocks, the kernel can create a new kernel threads to avoid blocking all user threads
- A new user thread doesn't necessarily require the creation of a new kernel thread
- True concurrency can be achieved on a multi-core machine
- Examples:
  - Solaris 9 and earlier
  - Win NT/2000 with the ThreadFiber package



# Two-Level Model

---



- The user can say: “Bind this thread to its own kernel thread”
- Example:
  - IRIX, HP-UX, Tru64 UNIX
  - Solaris 8 and earlier

# Thread Libraries

---

- Thread libraries provide users with ways to create threads in their own programs
  - In C/C++: pthreads and Win32 threads
    - Implemented by the kernel
  - In C/C++: OpenMP
    - A layer above pthreads for convenient multithreading in “easy” cases
  - In Java: Java Threads
    - Implemented by the JVM, which relies on threads implemented by the kernel

# Pthreads

---

- May be provided either as user-level or kernel-level
- A POSIX standard (IEEE 1003.1c) API for thread creation and synchronization
- ***Specification***, not ***implementation***
- API specifies behavior of the thread library, implementation is up to development of the library
- Common in UNIX operating systems (Linux & Mac OS X)



# Pthreads Example

---

```
#include <pthread.h>
#include <stdio.h>
#include <stdlib.h>
int sum; /* this data is shared by the thread(s) */

/* The thread will execute in this function */
void *runner(void *param) {
    int i, upper = atoi(param);
    sum = 0;

    for (i = 1; i <= upper; i++)
        sum += 1;
    pthread_exit(0);
}

int main(int argc, char *argv[]) {
    pthread_t tid; /* the thread identifier */
    pthread_attr_t attr; /* set of thread attributes */

    /* set the default attributes of the thread */
    pthread_attr_init(&attr);

    /* create the thread */
    pthread_create(&tid, &attr, runner, argv[1]);

    /* wait for the thread to exit */
    pthread_join(tid, NULL);

    printf("sum = %d\n", sum);
}
```

# Pthreads Code for Joining 10 Threads

---

```
#define NUM_THREADS 10

/* an array of threads to be joined upon */
pthread_t workers[NUM_THREADS];

for (int i = 0; i < NUM_THREADS; i++)
    pthread_join(workers[i], NULL);
```

# Windows Multithreaded C Program

---

```
#include <windows.h>
#include <stdio.h>
DWORD Sum; /* data is shared by the thread(s) */

/* The thread will execute in this function */
DWORD WINAPI Summation(LPVOID Param){
    DWORD Upper = *(DWORD *)Param;
    for (DWORD i = 1; i ≤ Upper; i++)
        Sum += i;
    return 0;
}

int main(int argc, char *argv[])
{
    DWORD ThreadId;
    HANDLE ThreadHandle;
    int Param;

    Param = atoi(argv[1]);

    /* create the thread */
    ThreadHandle = CreateThread(NULL, 0,
        Summation, /* thread function */
        &Param, /* parameter to thread function */
        0, &ThreadId); /* returns the thread identifier */

    /* now wait for the thread to finish */
    WaitForSingleObject(ThreadHandle, INFINITE);

    /* close the thread handle */ CloseHandle(ThreadHandle);
    printf("sum = %d\n", Sum);
}
```

# OpenMP

---

- Set of compiler directives and an API for C, C++, FORTRAN
- Provides support for parallel programming in shared-memory environments
- Identifies **parallel regions** – blocks of code that can run in parallel

**#pragma omp parallel**

Create as many threads as there are cores

```
#include <omp.h>
#include <stdio.h>

int main(int argc, char *argv[])
{
    /* sequential code */

    #pragma omp parallel
    {
        printf("I am a parallel region.");
    }

    /* sequential code */
    return 0;
}

#pragma omp parallel for
for (i = 0; i < N; i++)
{
    c[i] = a[i] + b[i];
}
```

# Java Threads

---


- All memory-management headaches go away with Java Threads
  - In nice Java fashion
- Several programming languages have long provided constructs/abstractions for writing concurrent programs
  - Modula, Ada, etc.
- Java threads may be created by:
  - Extending Thread class
  - Implementing the Runnable interface

```
class MyThread extends Thread {  
    public void run() {  
        . . .  
    }  
}  
MyThread t = new MyThread();  
  
public interface Runnable {  
    public abstract void run();  
}
```

# Thread Scheduling

---

- The JVM keeps track of threads, enacts the thread state transition diagram
- Question: who decides which runnable thread to run?
- Old versions of the JVM used **Green Threads**
  - User-level threads implemented by the JVM
  - Invisible to the O/S



# Beyond Green Threads


---

- Green threads have all the disadvantages of user-level threads (see earlier)
  - Most importantly: Cannot exploit multi-core, multi-processor architectures
- The JVM now provides **native threads**
  - Green threads are typically not available anymore
  - you can try to use “java -green” and see what your system says

# Java Threads / Kernel Threads

---

- In modern JVMs, application threads are *mapped* to kernel threads





# Threading Issues

---

- Semantics of **fork()** and **exec()** system calls
- Signal handling
  - Synchronous and asynchronous
- Thread cancellation of target thread
  - Asynchronous or deferred
- Thread-local storage
- Scheduler Activations

# Semantics of fork() and exec()

---

- What happens when a thread calls fork()?
- Two possibilities:
  - A new process is created that has only one thread (the copy of the thread that called fork()), or
  - A new process is created with all threads of the original process (a copy of all the threads, including the one that called fork())
- Some OSes provide both options
  - In Linux the first option above is used
- If one calls exec() after fork(), all threads are “wiped out” anyway

# Signals

---

- We've talked about signals for processes
  - Signal handlers are either default or user-specified
  - `signal()` and `kill()` are the system calls
- In a multi-threaded program, what happens?
- Multiple options
  - Deliver the signal to the thread to which the signal applies
  - Deliver the signal to every thread in the process
  - Deliver the signal to certain threads in the process
  - Assign a specific thread to receive all signals

# Signals

---

- Most UNIX versions: a thread can say which signals it accepts and which signals it doesn't accept
- On Linux: dealing with threads and signals is tricky but well understood with many tutorials on the matter and man pages
  - `man pthread_sigmask`
  - `man sigemptyset`
  - `man sigaction`

# Safe Thread Cancellation

---

- One potentially useful feature would be for a thread to simply terminate another thread
- Two possible approaches:
  - **Asynchronous** cancellation
    - One thread terminates another immediately
  - **Deferred** cancellation
    - A thread periodically checks whether it should terminate

# Safe Thread Cancellation

---

- Two possible approaches:
  - **Asynchronous** cancellation
    - One thread terminates another immediately
  - **Deferred** cancellation
    - A thread periodically checks whether it should terminate
  
- Pthread code to create and cancel a thread:

```
pthread_t tid;

/* create the thread */
pthread_create(&tid, 0, worker, NULL);

. . .

/* cancel the thread */
pthread_cancel(tid);

/* wait for the thread to terminate */
pthread_join(tid, NULL);
```

# Safe Thread Cancellation

---

- Invoking thread cancellation requests cancellation, but actual cancellation depends on thread state

Mode	State	Type
Off	Disabled	–
Deferred	Enabled	Deferred
Asynchronous	Enabled	Asynchronous

- If thread has cancellation disabled (**off**), cancellation remains pending until thread enables it
- Default type is deferred
  - Cancellation only occurs when thread reaches **cancellation point**
    - I.e. `pthread_testcancel()`
    - Then **cleanup handler** is invoked
- On Linux systems, thread cancellation is handled through signals

# Safe Thread Cancellation

---

- The problem with asynchronous cancellation:
  - may lead to an inconsistent state or to a synchronization problem if the thread was in the middle of “something important”
  - Absolutely terrible bugs lurking in the shadows
- The problem with deferred cancellation: the code is cumbersome due to multiple cancellation points
  - should I die? should I die? should I die?
- In Java, the `Thread.stop()` method is deprecated, and so cancellation has to be deferred



# Operating System Examples

---

- Windows Threads
- Linux Threads

# Windows Threads

---

- Windows API – primary API for Windows applications
- Implements the one-to-one mapping, kernel-level
- Each thread contains
  - A thread id
  - Register set representing state of processor
  - Separate user and kernel stacks for when thread runs in user mode or kernel mode
  - Private data storage area used by run-time libraries and dynamic link libraries (DLLs)
- The register set, stacks, and private storage area are known as the **context** of the thread

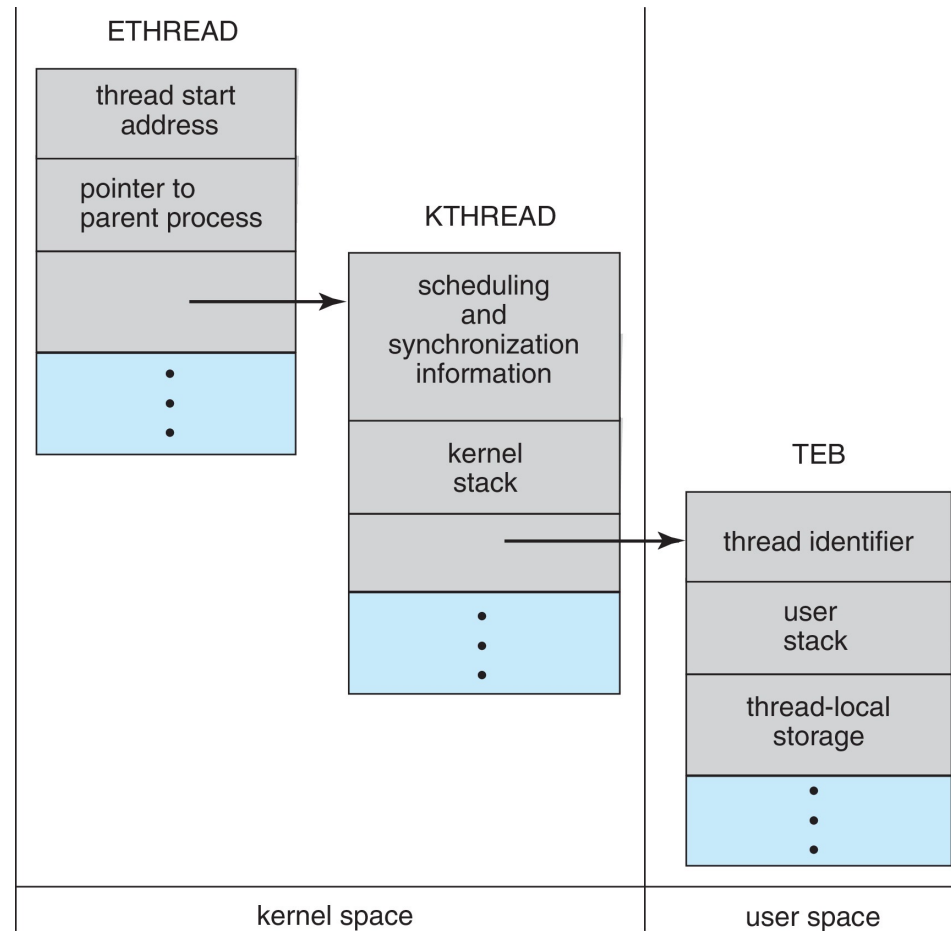
# Windows Threads (Cont.)

---

- The primary data structures of a thread include:
  - ETHREAD (executive thread block) – includes pointer to process to which thread belongs and to KTHREAD, in kernel space
  - KTHREAD (kernel thread block) – scheduling and synchronization info, kernel-mode stack, pointer to TEB, in kernel space
  - TEB (thread environment block) – thread id, user-mode stack, thread-local storage, in user space

# Windows Threads Data Structures

---



# Linux Threads

---

- In Linux, a thread is also called a light-weight process (LWP)
- The clone() syscall is used to create a thread or a process
  - Shares execution context with its parent
  - pthread library uses clone() to implement threads. Refer to `./nptl/sysdeps/pthread/createthread.c`

flag	meaning
CLONE_FS	File-system information is shared.
CLONE_VM	The same memory space is shared.
CLONE_SIGHAND	Signal handlers are shared.
CLONE_FILES	The set of open files is shared.

# Linux Threads

---

- Linux does not distinguish between PCB and TCB
  - Kernel data structure: task\_struct


```
591
592 struct task_struct {
593 #ifdef CONFIG_THREAD_INFO_IN_TASK
594     /*
595      * For reasons of header soup (see current_thread_info()), this
596      * must be the first element of task_struct.
597      */
598     struct thread_info          thread_info;
599 #endif
600     /* -1 unrunnable, 0 runnable, >0 stopped: */
601     volatile long               state;
602
603     /*
604      * This begins the randomizable portion of task_struct. Only
605      * scheduling-critical items should be added above here.
606      */
607     randomized_struct_fields_start
608
609     void                        *stack;
610     atomic_t                    usage;
611     /* Per task flags (PF_*), defined further below: */
612     unsigned int                flags;
613     unsigned int                ptrace;
614
```

# Linux Threads

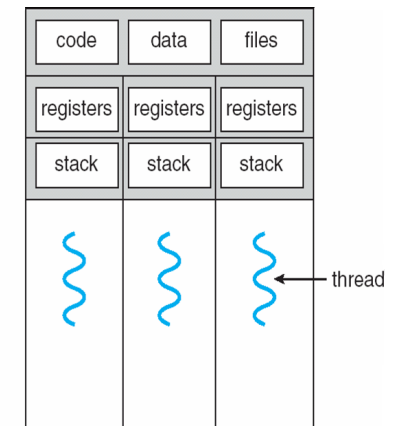
## ■ Single-threaded process vs multi-threaded process

```
wenbo@wenbo-desktop:~/KERNEL/linux.git$ ps -eLf
```

UID	PID	PPID	LWP	C	NLWP	STIME	TTY	TIME	CMD
root	1	0	1	0	1	3月11 ?		00:00:19	/sbin/init splash
root	2	0	2	0	1	3月11 ?		00:00:00	[kthreadd]
root	4	2	4	0	1	3月11 ?		00:00:00	[kworker/0:0H]
root	6	2	6	0	1	3月11 ?		00:00:00	[mm_percpu_wq]
root	7	2	7	0	1	3月11 ?		00:00:00	[ksoftirqd/0]
root	8	2	8	0	1	3月11 ?		00:00:31	[rcu_sched]
root	9	2	9	0	1	3月11 ?		00:00:00	[rcu_bh]
root	10	2	10	0	1	3月11 ?		00:00:00	[migration/0]
root	11	2	11	0	1	3月11 ?		00:00:00	[watchdog/0]
root	704	1	704	0	1	3月11 ?		00:00:00	/usr/sbin/cron -f
root	718	1	718	0	16	3月11 ?		00:00:00	/usr/lib/snapd/snapd
root	718	1	882	0	16	3月11 ?		00:00:00	/usr/lib/snapd/snapd
root	718	1	883	0	16	3月11 ?		00:00:00	/usr/lib/snapd/snapd
root	718	1	884	0	16	3月11 ?		00:00:00	/usr/lib/snapd/snapd
root	718	1	885	0	16	3月11 ?		00:00:00	/usr/lib/snapd/snapd
root	718	1	917	0	16	3月11 ?		00:00:00	/usr/lib/snapd/snapd
root	718	1	921	0	16	3月11 ?		00:00:01	/usr/lib/snapd/snapd
root	718	1	922	0	16	3月11 ?		00:00:00	/usr/lib/snapd/snapd
root	718	1	923	0	16	3月11 ?		00:00:01	/usr/lib/snapd/snapd
root	718	1	924	0	16	3月11 ?		00:00:01	/usr/lib/snapd/snapd



single-threaded process



multithreaded process

# Linux Threads

- Single-threaded process vs multi-threaded process

```
wenbo@wenbo-desktop:~/KERNEL/linux.git$ ps -eLf
```

UID	PID	PPID	LWP	C	NLWP	STIME	TTY	TIME	CMD
root	1	0	1	0	1	3月11 ?		00:00:19	/sbin/init splash
root	2	0	2	0	1	3月11 ?		00:00:00	[kthreadd]
root	4	2	4	0	1	3月11 ?		00:00:00	[kworker/0:0H]
root	6	2	6	0	1	3月11 ?		00:00:00	[mm_percpu_wq]
root	7	2	7	0	1	3月11 ?		00:00:00	[ksoftirqd/0]
root	8	2	8	0	1	3月11 ?		00:00:31	[rcu_sched]
root	9	2	9	0	1	3月11 ?		00:00:00	[rcu_bh]
root	10	2	10	0	1	3月11 ?		00:00:00	[migration/0]
root	11	2	11	0	1	3月11 ?		00:00:00	[watchdog/0]
root	704	1	704	0	1	3月11 ?		00:00:00	/usr/sbin/cron -f
root	718	1	718	0	16	3月11 ?		00:00:00	/usr/lib/snapd/snapd
root	718	1	882	0	16	3月11 ?		00:00:00	/usr/lib/snapd/snapd
root	718	1	883	0	16	3月11 ?		00:00:00	/usr/lib/snapd/snapd
root	718	1	884	0	16	3月11 ?		00:00:00	/usr/lib/snapd/snapd
root	718	1	885	0	16	3月11 ?		00:00:00	/usr/lib/snapd/snapd
root	718	1	917	0	16	3月11 ?		00:00:00	/usr/lib/snapd/snapd
root	718	1	921	0	16	3月11 ?		00:00:01	/usr/lib/snapd/snapd
root	718	1	922	0	16	3月11 ?		00:00:00	/usr/lib/snapd/snapd
root	718	1	923	0	16	3月11 ?		00:00:01	/usr/lib/snapd/snapd
root	718	1	924	0	16	3月11 ?		00:00:01	/usr/lib/snapd/snapd

```
787      /* PID/PID hash table linkage. */
788      struct pid *thread_pid;
789      struct hlist_node pid_links[PIDTYPE];
790      struct list_head thread_group;
791      struct list_head thread_node;
792
793      struct completion *vfork_done;
794
795      /* CLONE_CHILD_SETTID: */
796      int __user *set_child_tid;
797      ---
```




# Linux Threads

---

- Linux uses the same `task_struct` for ~~process and thread~~
- A process is
  - either a **single** thread + an address space
    - PID is thread ID
  - or **multiple** threads + an address space
    - PID is the leading thread ID

# Threads within Process



# Threads with Process – What is shared

```
--
29 static void traversal_thread_group(struct task_struct * tsk){
30     struct task_struct * curr_thread = NULL;
31     unsigned long tg_offset = offsetof(struct task_struct, thread_group);
32
33     curr_thread = (struct task_struct *) (((unsigned long)tsk->thread_group.next) - tg_offset);
34     while (curr_thread != tsk){
35         printk("\t\tTHREAD TSK=%llx\tPID=%d\tSTACK=%llx \tCOMM=%s\tMM=%llx\tACTIVE_MM=%llx\n",
36             (u64)curr_thread, curr_thread->pid, (u64)curr_thread->stack,
37             curr_thread->comm, (u64)curr_thread->mm, (u64)curr_thread->active_mm);
38         curr_thread = (struct task_struct *) (((unsigned long)curr_thread->thread_group.next) - tg_offset);
39     }
40 }
41
42 static void traversal_process(void) {
43     struct task_struct * tsk = NULL;
44
45     traversal_thread_group(&init_task);
46     for_each_process(tsk){
47         printk("PROCESS\tTHREAD TSK=%llx\tPID=%d\tSTACK=%llx \tCOMM=%s\tMM=%llx\tACTIVE_MM=%llx\n",
48             (u64)tsk, tsk->pid, (u64)tsk->stack, tsk->comm,
49             (u64)tsk->mm, (u64)tsk->active_mm);
50         traversal_thread_group(tsk);
51     }
52 }
```

# Threads with Process – What is shared

```
--
29 static void traversal_thread_group(struct task_struct * tsk){
30     struct task_struct * curr_thread = NULL;
31     unsigned long tg_offset = offsetof(struct task_struct, thread_group);
32
33     curr_thread = (struct task_struct *) (((unsigned long)tsk->thread_group.next) - tg_offset);
34     while (curr_thread != tsk){
35         printk("\t\tTHREAD TSK=%llx\tPID=%d\tSTACK=%llx \tCOMM=%s\tMM=%llx\tACTIVE_MM=%llx\n",
36             (u64)curr_thread, curr_thread->pid, (u64)curr_thread->stack,
37             curr_thread->comm, (u64)curr_thread->mm, (u64)curr_thread->active_mm);
38         curr_thread = (struct task_struct *) (((unsigned long)curr_thread->thread_group.next) - tg_offset);
39     }
40 }
41
42 static void traversal_process(void) {
43     struct task_struct * tsk = NULL;
44
45     traversal_thread_group(&init_task);
46     for_each_process(tsk){
47         printk("PROCESS\t\tTHREAD TSK=%llx\tPID=%d\tSTACK=%llx \tCOMM=%s\tMM=%llx\tACTIVE_MM=%llx\n",
48             (u64)tsk, tsk->pid, (u64)tsk->stack, tsk->comm,
49             (u64)tsk->mm, (u64)tsk->active_mm);
50         traversal_thread_group(tsk);
51     }
52 }
```

PROCESS	THREAD	TSK=ffff8c4c4bf3c5c0	PID=718	STACK=ffff985c82268000	COMM=snapd	MM=ffff8c4c46400840	ACTIVE_MM=ffff8c4c46400840
	THREAD	TSK=ffff8c4c46d52e80	PID=882	STACK=ffff985c82390000	COMM=snapd	MM=ffff8c4c46400840	ACTIVE_MM=ffff8c4c46400840
	THREAD	TSK=ffff8c4c46d545c0	PID=883	STACK=ffff985c822e8000	COMM=snapd	MM=ffff8c4c46400840	ACTIVE_MM=ffff8c4c46400840
	THREAD	TSK=ffff8c4c491b45c0	PID=884	STACK=ffff985c8218c000	COMM=snapd	MM=ffff8c4c46400840	ACTIVE_MM=ffff8c4c46400840
	THREAD	TSK=ffff8c4c4cbeb1740	PID=885	STACK=ffff985c821ec000	COMM=snapd	MM=ffff8c4c46400840	ACTIVE_MM=ffff8c4c46400840
	THREAD	TSK=ffff8c4c4cae1ae80	PID=917	STACK=ffff985c823c8000	COMM=snapd	MM=ffff8c4c46400840	ACTIVE_MM=ffff8c4c46400840
	THREAD	TSK=ffff8c4c4b562e80	PID=921	STACK=ffff985c82418000	COMM=snapd	MM=ffff8c4c46400840	ACTIVE_MM=ffff8c4c46400840
	THREAD	TSK=ffff8c4c48340000	PID=922	STACK=ffff985c823b0000	COMM=snapd	MM=ffff8c4c46400840	ACTIVE_MM=ffff8c4c46400840
	THREAD	TSK=ffff8c4c472bae80	PID=923	STACK=ffff985c821f4000	COMM=snapd	MM=ffff8c4c46400840	ACTIVE_MM=ffff8c4c46400840
	THREAD	TSK=ffff8c4c4b5945c0	PID=924	STACK=ffff985c81fa8000	COMM=snapd	MM=ffff8c4c46400840	ACTIVE_MM=ffff8c4c46400840
	THREAD	TSK=ffff8c4c46775d00	PID=925	STACK=ffff985c822a8000	COMM=snapd	MM=ffff8c4c46400840	ACTIVE_MM=ffff8c4c46400840
	THREAD	TSK=ffff8c4c4b692e80	PID=973	STACK=ffff985c82438000	COMM=snapd	MM=ffff8c4c46400840	ACTIVE_MM=ffff8c4c46400840
	THREAD	TSK=ffff8c4c4b78ae80	PID=974	STACK=ffff985c823c0000	COMM=snapd	MM=ffff8c4c46400840	ACTIVE_MM=ffff8c4c46400840
	THREAD	TSK=ffff8c4c46e1dd00	PID=975	STACK=ffff985c824b8000	COMM=snapd	MM=ffff8c4c46400840	ACTIVE_MM=ffff8c4c46400840

# Threads within Process – What is shared

PROCESS	THREAD	TSK=ffff8c4c4bf3c5c0	PID=718	STACK=ffff985c82268000	COMM=snapd	MM=ffff8c4c46400840	ACTIVE_MM=ffff8c4c46400840
	THREAD	TSK=ffff8c4c46d52e80	PID=882	STACK=ffff985c82390000	COMM=snapd	MM=ffff8c4c46400840	ACTIVE_MM=ffff8c4c46400840
	THREAD	TSK=ffff8c4c46d545c0	PID=883	STACK=ffff985c822e8000	COMM=snapd	MM=ffff8c4c46400840	ACTIVE_MM=ffff8c4c46400840
	THREAD	TSK=ffff8c4c491b45c0	PID=884	STACK=ffff985c8218c000	COMM=snapd	MM=ffff8c4c46400840	ACTIVE_MM=ffff8c4c46400840
	THREAD	TSK=ffff8c4c4beeb1740	PID=885	STACK=ffff985c821ec000	COMM=snapd	MM=ffff8c4c46400840	ACTIVE_MM=ffff8c4c46400840
	THREAD	TSK=ffff8c4c4ae1ae80	PID=917	STACK=ffff985c823c8000	COMM=snapd	MM=ffff8c4c46400840	ACTIVE_MM=ffff8c4c46400840
	THREAD	TSK=ffff8c4c4b562e80	PID=921	STACK=ffff985c82418000	COMM=snapd	MM=ffff8c4c46400840	ACTIVE_MM=ffff8c4c46400840
	THREAD	TSK=ffff8c4c48340000	PID=922	STACK=ffff985c823b0000	COMM=snapd	MM=ffff8c4c46400840	ACTIVE_MM=ffff8c4c46400840
	THREAD	TSK=ffff8c4c472bae80	PID=923	STACK=ffff985c821f4000	COMM=snapd	MM=ffff8c4c46400840	ACTIVE_MM=ffff8c4c46400840
	THREAD	TSK=ffff8c4c4b5945c0	PID=924	STACK=ffff985c81fa8000	COMM=snapd	MM=ffff8c4c46400840	ACTIVE_MM=ffff8c4c46400840
	THREAD	TSK=ffff8c4c46775d00	PID=925	STACK=ffff985c822a8000	COMM=snapd	MM=ffff8c4c46400840	ACTIVE_MM=ffff8c4c46400840
	THREAD	TSK=ffff8c4c4b692e80	PID=973	STACK=ffff985c82438000	COMM=snapd	MM=ffff8c4c46400840	ACTIVE_MM=ffff8c4c46400840
	THREAD	TSK=ffff8c4c4b78ae80	PID=974	STACK=ffff985c823c0000	COMM=snapd	MM=ffff8c4c46400840	ACTIVE_MM=ffff8c4c46400840
	THREAD	TSK=ffff8c4c46e1dd00	PID=975	STACK=ffff985c824b8000	COMM=snapd	MM=ffff8c4c46400840	ACTIVE_MM=ffff8c4c46400840
		task_struct	pid	stack	comm		mm_struct

Not Shared

Shared

# User thread to kernel thread mapping

---

- One task in Linux
  - Same task\_struct (PCB) means same thread
    - Also viewed as 1:1 mapping
    - One user thread maps to one kernel thread
    - But actually, they are the same thread
  - Can be executed in user space
    - User code, user space stack
  - Can be executed in kernel space
    - Kernel code, kernel space stack

# User thread to kernel thread mapping

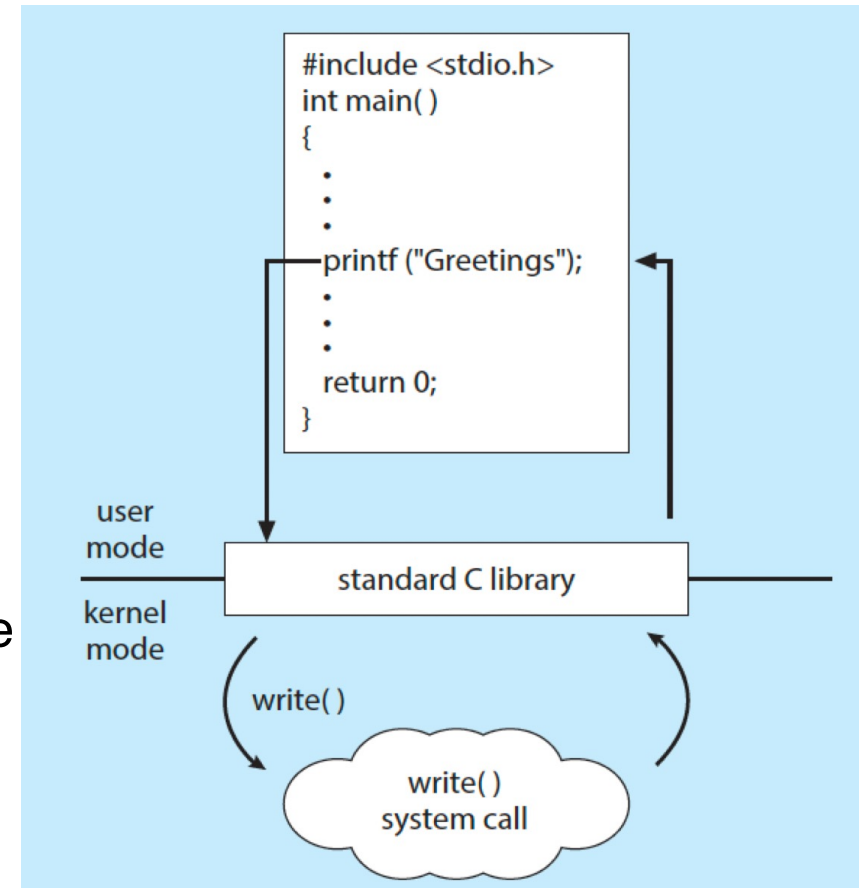
---

- One task
  - Can be a single-threaded process
    - One task\_struct – PCB
    - PID is thread ID
  - Can be a multi-threaded process
    - Multiple task\_struct
    - What is the PID?
  - Can be executed in user space
    - User code, user space stack
  - Can be executed in kernel space
    - Such as calls a system call
    - Execution flow traps to kernel
    - Execute kernel code, use kernel space stack



# User thread to kernel thread mapping

- One task
  - Can be a single-threaded process
    - One task\_struct – PCB
  - Can be a multi-threaded process
    - Multiple task\_struct
  - Execution
    - Can be executed in user space
      - User code, user space stack
    - Can be executed in kernel space
      - Such as calls a system call
      - Execution flow traps to kernel
      - Execute kernel code, use kernel space stack



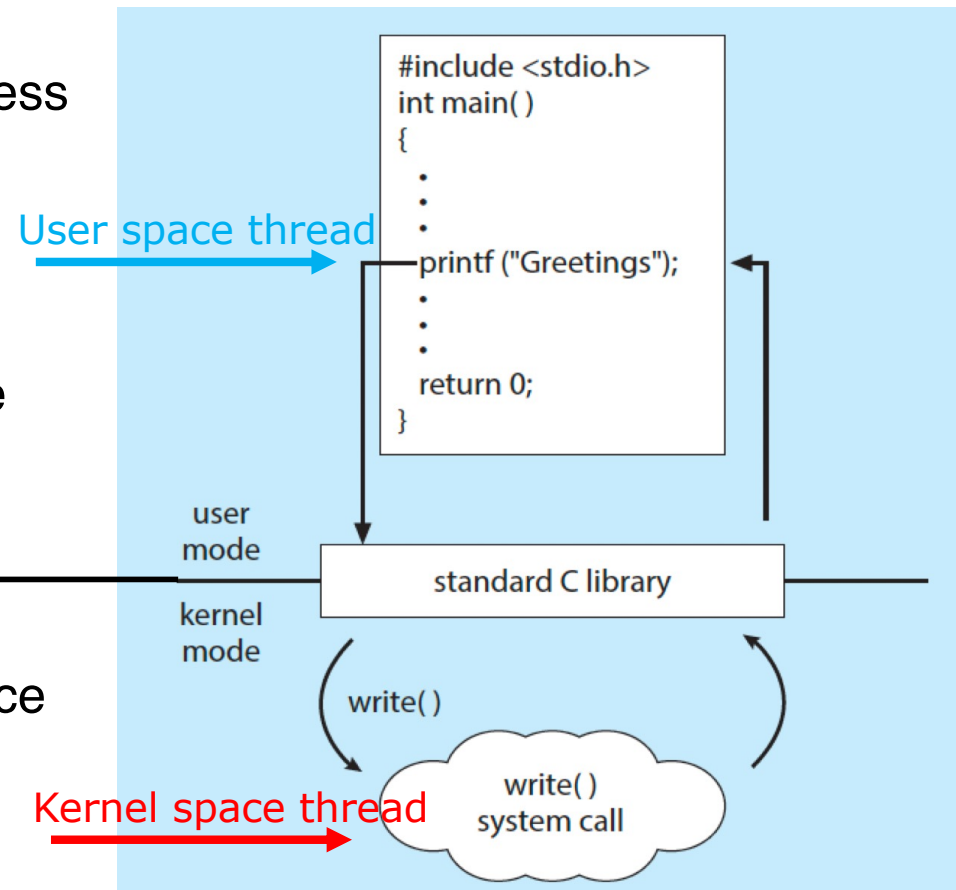


# User thread to kernel thread mapping

## ■ One task

- Can be a single-threaded process
  - One task\_struct – PCB
- Can be a multi-threaded process
  - Multiple task\_struct
- Can be executed in user space
  - User code, user space stack

- Can be executed in kernel space
  - Such as calls a system call
  - Execution flow traps to kernel
  - Execute kernel code, use kernel space stack



# Takeaway

---

- Thread is the basic execution unit
  - Has its own registers, pc, stack
- Thread vs Process
  - What is shared and what is not
- Pros and cons of thread