

# Paging




Operating Systems  
Wenbo Shen

# Paging: basic idea

---

## Contiguous → Noncontiguous

- Physical address space of a process can be **noncontiguous**; process is allocated physical memory whenever the latter is available
- Fixed and variable partitions are physical **contiguous** allocation
- Avoids **external fragmentation**
- Avoids problem of varying sized memory chunks



# Paging

---

## Basic methods

- Divide **physical** address into fixed-sized blocks called **frames**
  - Size is power of 2, usually 4KB
- Divide **logical** address into blocks of same size called **pages**
- Keep track of all free frames
- To run a program of size **N** pages, need to find **N** free frames and load program
- Set up a mapping to translate logical to physical addresses
  - This mapping is called page table

# Paging: Internal Fragmentation

---


- Paging has **no external fragmentation**, but **internal fragmentation**
  - e.g., page size: 2,048, program size: 72,766 (35 pages + 1,086 bytes)
    - internal fragmentation:  $2,048 - 1,086 = 962$
  - **worst case** internal fragmentation: **1 frame – 1 byte**
  - **average** internal fragmentation: 1 / 2 frame size
    - In practice, average is much smaller
    - Program contains several pages, only last page has fragmentation
- Small frame sizes more desirable than large frame size?
  - Small size means **less** internal fragmentation, but **more** page table entries
  - Large size means **more** ..., but **less** ...
  - page size is usually 4KB, can be 64KB, block can be 2MB or 1GB (64-bit)

# Page Table

---

- Page table
  - Stores the logical page to physical frame mapping
- Frame table
  - OS is managing physical memory, it should be aware of the allocation details of physical memory (frame)
  - Which frame is free, and how many frames have been allocated ...
  - One entry for each physical frame
  - The allocated frame belongs to which process

# Paging example



Before allocation

After allocation

# Paging: Address Translation


---

- A logical address is divided into:
  - **page number (p)**
    - used as an index into a page table
    - page table entry contains the corresponding **physical** frame number
  - **page offset (d)**
    - offset within the page/frame
    - combined with frame number to get the physical address

page number	page offset
$p$	$d$
$m - n \text{ bits}$	$n \text{ bits}$

$m$  bit logical address space,  $n$  bit page size

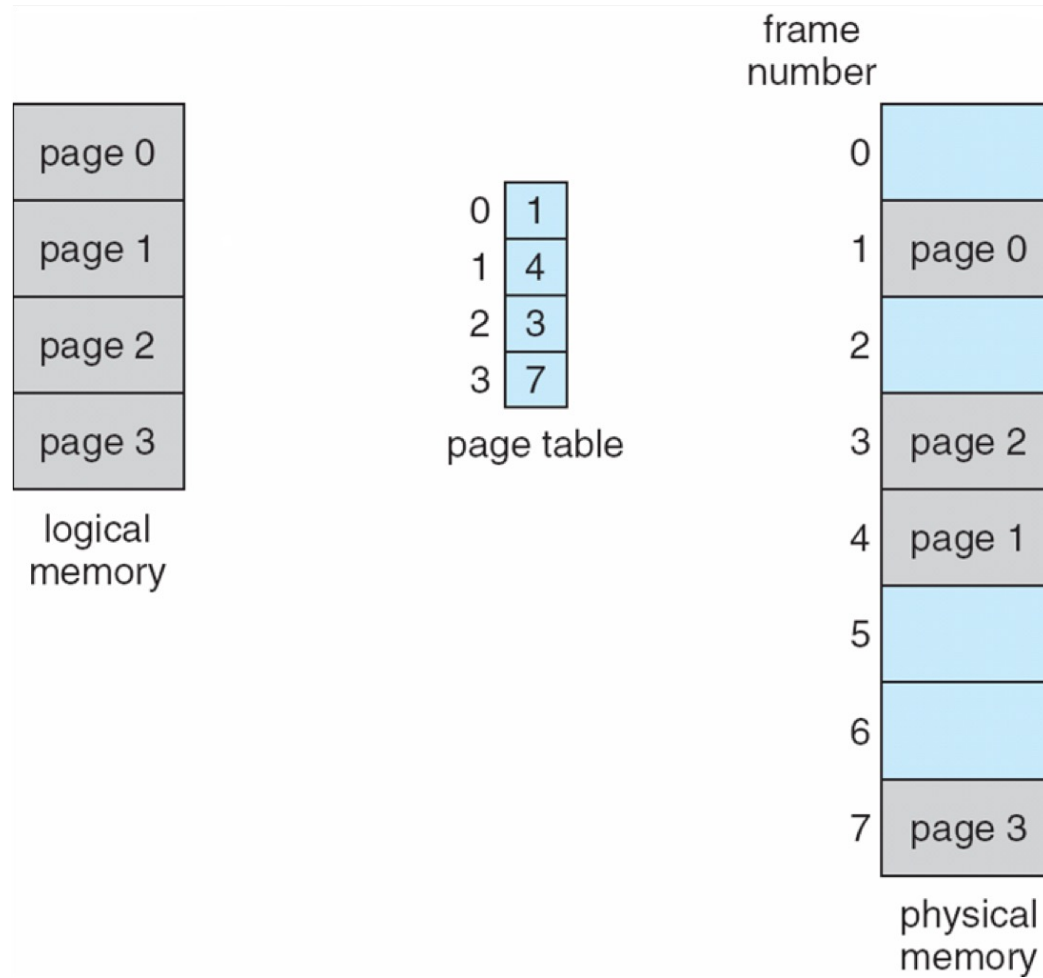
# Paging Hardware






# Paging Example

---



# Paging Example II



$m = 4$  and  $n = 2$  32-byte memory and 4-byte pages

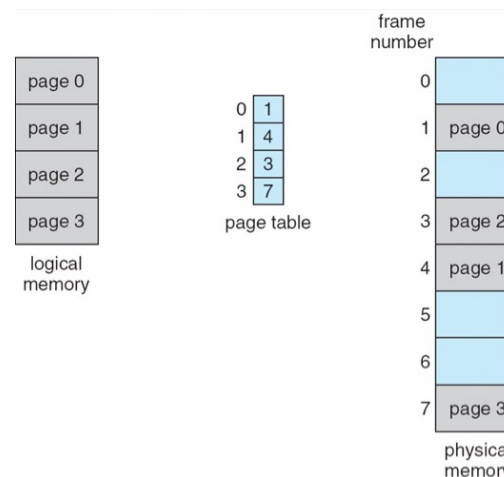
# Page Table Hardware Support: Simplest Case

---

- Page table is in a set of dedicated registers
  - Advantages: very efficient - access to register is fast
  - Disadvantages:
    - register number is limited → the table size is very small
    - the context switch need to save and restore these registers

# Page Table Hardware Support: Alternative Way

- One big page table maps logical address to physical address
  - the page table should be **kept in main memory**
  - **page-table base register (PTBR)** points to the page table
    - **does PTBR contain physical or logical address?**
  - **page-table length register (PTLR)** indicates the size of the page table
- Every data/instruction access requires **two memory accesses**
  - One for the page table and one for the data / instruction
  - How to reduce memory accesses caused by page table?
    - CPU can cache the translation to avoid one memory access (**TLB**)



# TLB

---

- TLB (**translation look-aside buffer**) caches the address translation
  - if page number is in the TLB, no need to access the page table
  - if page number is not in the TLB, need to replace one TLB entry
  - TLB usually use a fast-lookup hardware cache called **associative memory**
  - TLB is usually small, 64 to 1024 entries
- Use with page table
  - TLB contains a few page table entries
  - Check whether page number is in TLB
    - If -> frame number is available and used
    - If not -> **TLB miss** access page table and then fetch into TLB

# TLB

---

- TLB and context switch
  - Each process has its own page table
    - switching process needs to switch page table
  - **TLB must be consistent with page table**
    - Option I: Flush TLB at every context switch, or,
    - Option II: Tag TLB entries with **address-space identifier (ASID)** that uniquely identifies a process
  - Some TLB entries can be **shared** by processes, and fixed in the TLB
    - e.g., TLB entries for the kernel
- TLB and operating system
  - MIPS: OS should deal with TLB miss exception
  - X86: TLB miss is handled by hardware

# Associative Memory

---

- Associative memory: memory that supports **parallel search**
- Associative memory is not addressed by "addresses", but **contents**
  - if page# is in associative memory's key, return frame# (value) directly

Page #	Frame #
1	7
2	12
3	15
4	31

# More on TLB Entries

---

- Instruction micro TLB
- Data micro TLB
- Main TLB
  - A 4-way, set-associative, 1024 entry cache which stores VA to PA mappings for 4KB, 16KB, and 64KB page sizes.
  - A 2-way, set-associative, 128 entry cache which stores VA to PA mappings for 1MB, 2MB, 16MB, 32MB, 512MB, and 1GB block sizes.



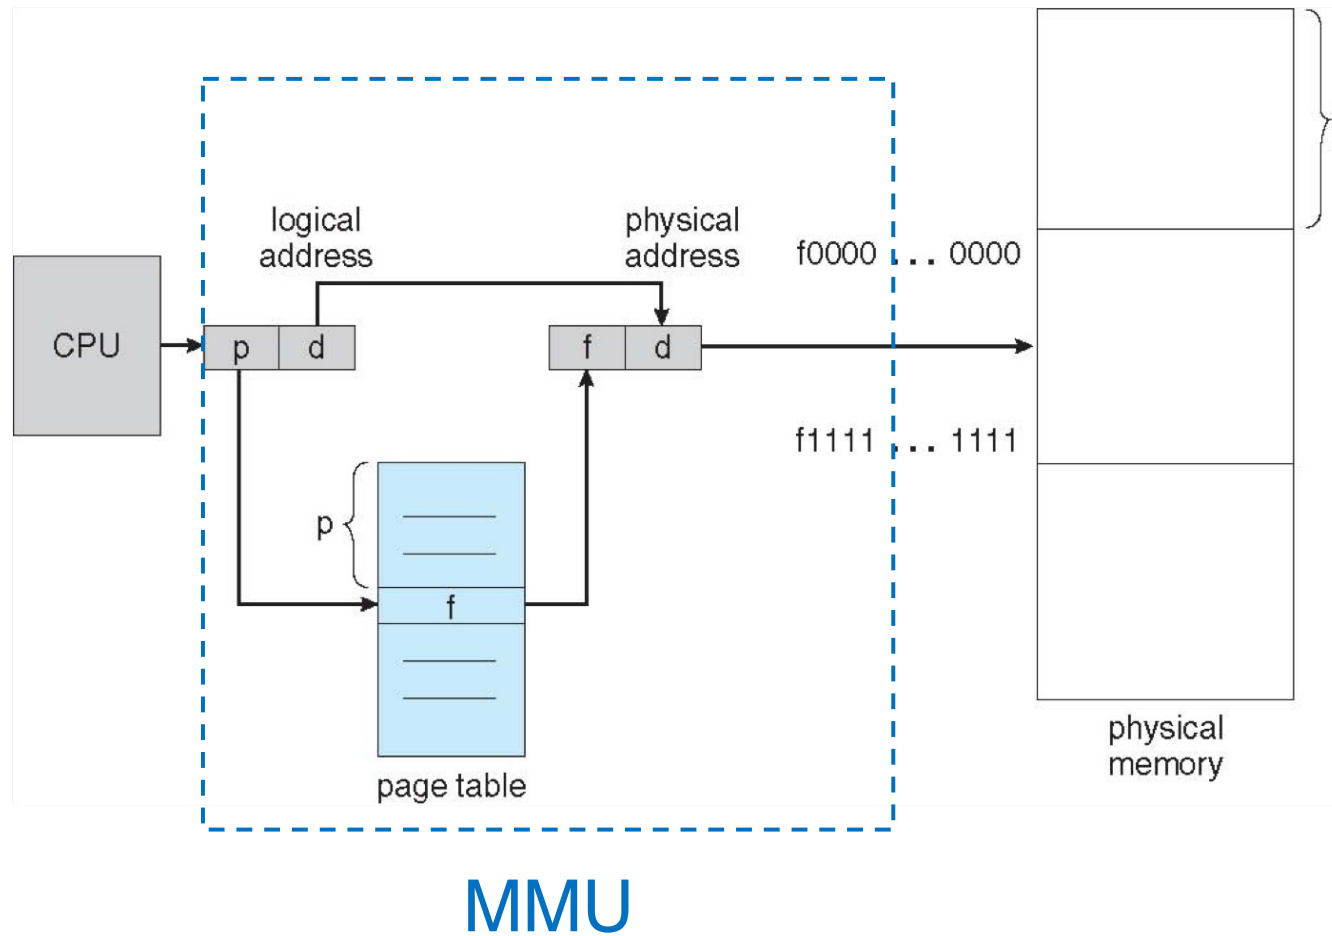
# More on TLB Match Process

---

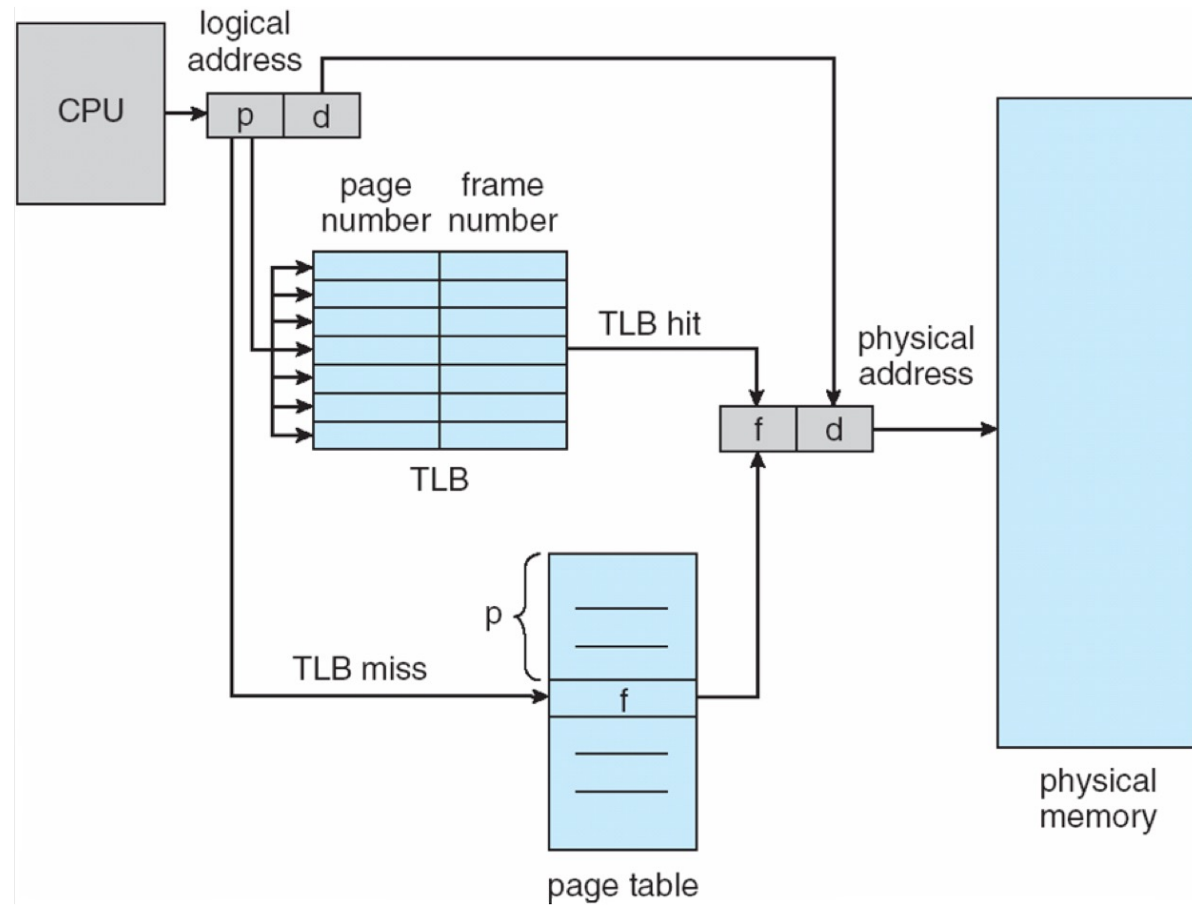
A TLB entry match occurs when the following conditions are met:

- Its VA, moderated by the page size such as the VA bits[47:12], matches that of the requested address
- The memory space matches the memory space state of the requests
- The ASID matches the current ASID held in the CONTEXTIDR, TTBR0, or TTBR1 register or the entry is marked global
- The VMID matches the current VMID held in the VTTBR register

# Paging Hardware



# Paging Hardware With TLB



# Effective Access Time

---

- **Hit ratio** - percentage of times that a page number is found in the TLB
- An 80% hit ratio means that we find the desired page number in the TLB 80% of the time.
- Suppose that 100 nanoseconds to access memory.
  - If we find the desired page in TLB then a mapped-memory access take 100 ns
  - Otherwise we need **two memory access** so it is 200 ns: page table + memory access
- Effective Access Time (EAT)
  - $EAT = 0.80 \times 100 + 0.20 \times 200 = 120$  nanoseconds
  - implying 20% slowdown in access time
- Consider a more realistic hit ratio of 99%,
  - $EAT = 0.99 \times 100 + 0.01 \times 200 = 101$ ns
  - implying only 1% slowdown in access time

# Memory Protection

---

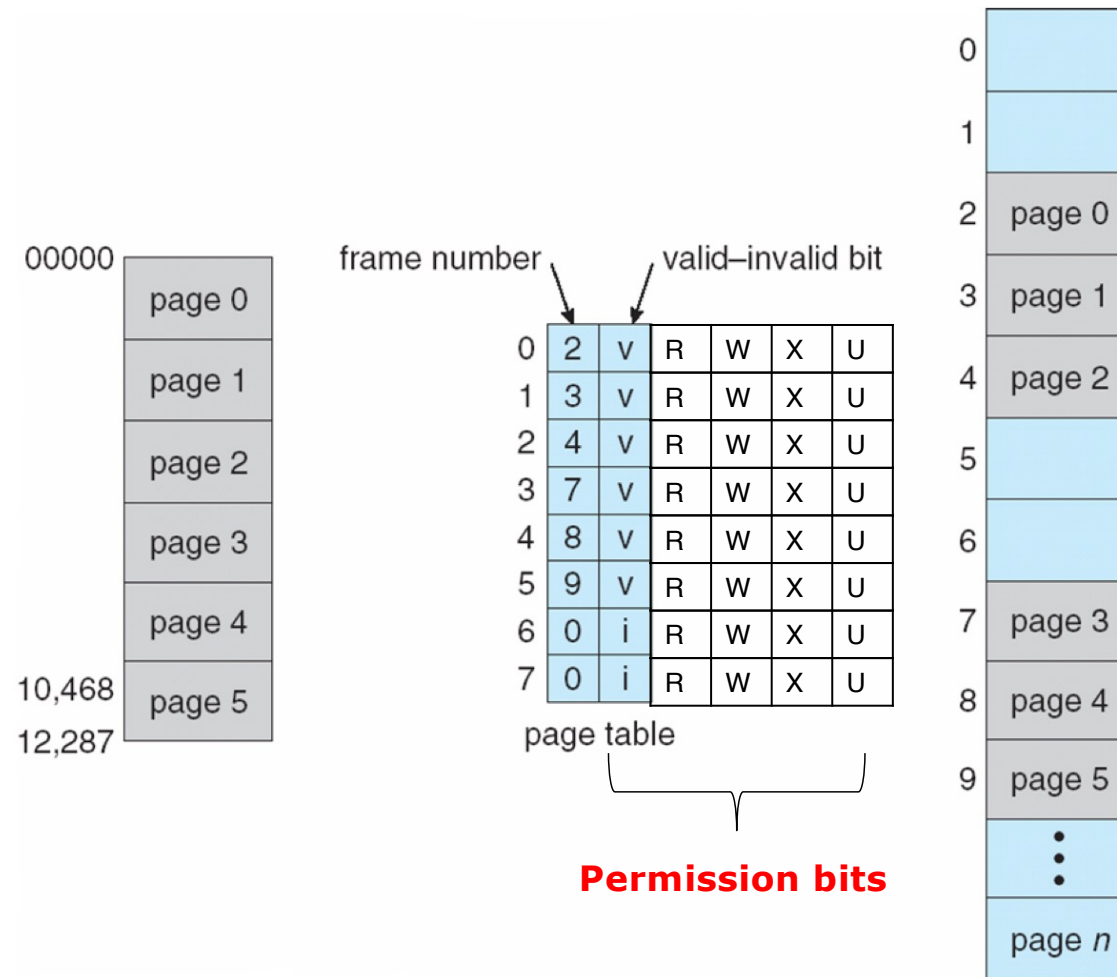
- Accomplished by protection bits with each frame
- Each page table entry has a **present** (aka. valid) bit
  - present: the page has a valid physical frame, thus can be accessed
- Each page table entry contains some protection bits
  - **kernel/user, read/write, execution?, kernel-execution?**
  - why do we need them?
- Any violations of memory protection result in a trap to the kernel

# Memory Protection (more)

---

- XN: protecting code
  - Segregate areas of memory for use by either storage of processor instructions (code) or for storage of data
  - Intel: XD(execute disable), AMD: EVP (enhanced virus protection), ARM: XN (execute never)
- PXN: Privileged Execute Never
  - A Permission fault is generated if the processor is executing at **EL1(kernel)** and attempts to execute an instruction fetched from the corresponding memory region when this PXN bit is 1 (usually user space memory)
  - Intel: SMEP

# Memory Protection



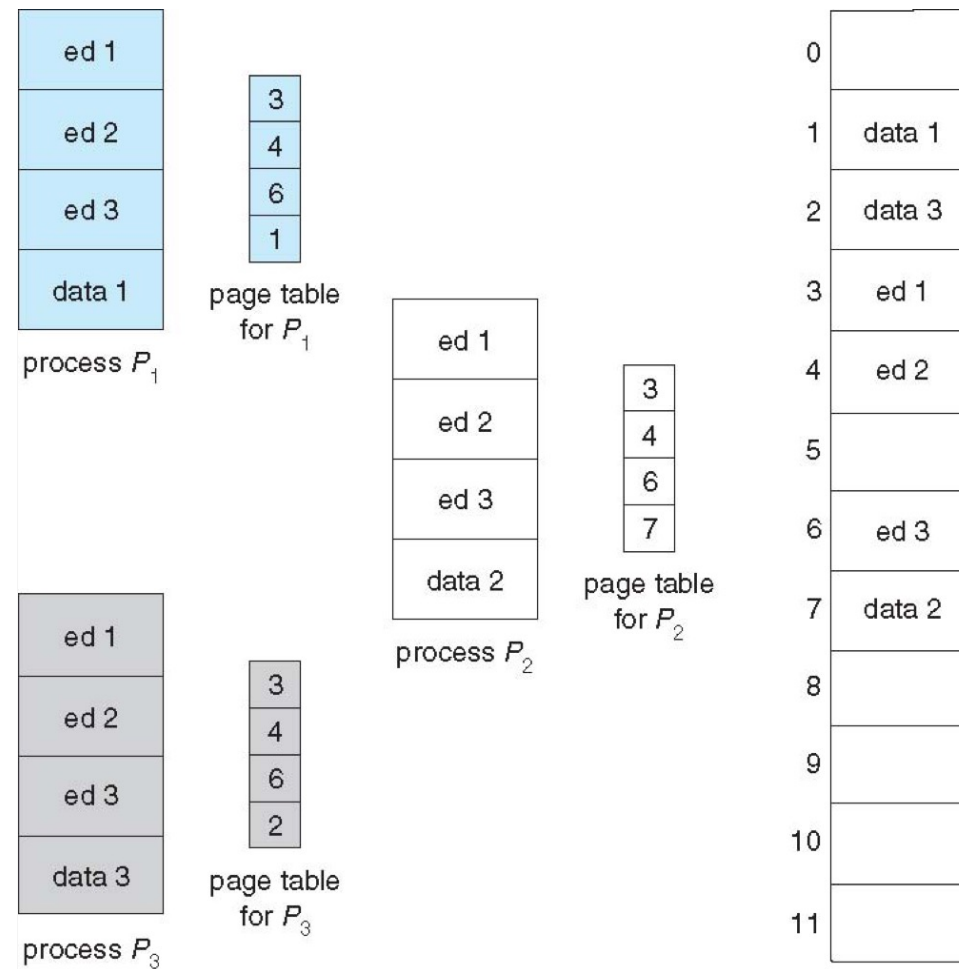
# Page Sharing

---

- Paging allows to share memory between processes
  - e.g., one copy of **code** shared by **all processes of the same program**
    - text editors, compilers, browser..
  - shared memory can be used for **inter-process communication**
  - shared libraries
- Reentrant code: non-self-modifying code: never changes between execution
- Each process can, of course, have its private code and data



# Page Sharing

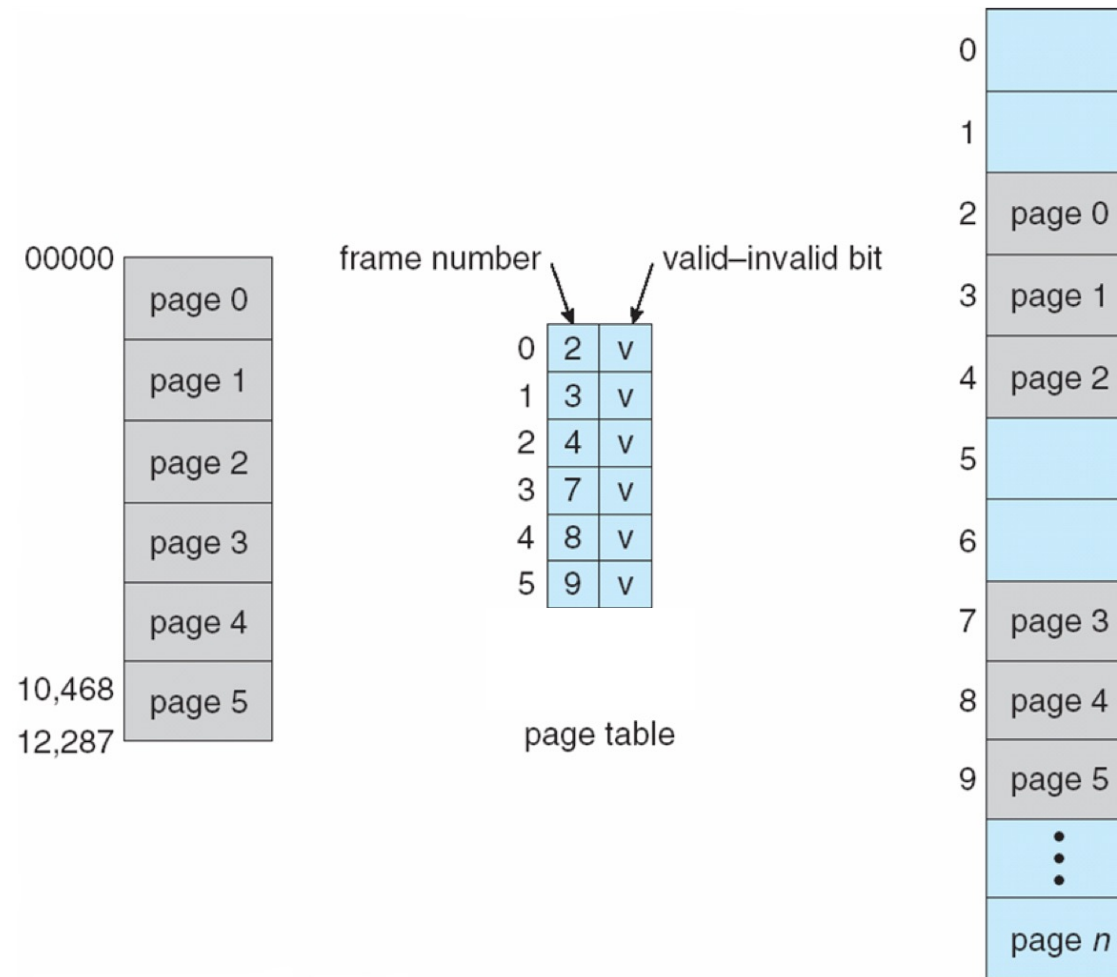


# Structure of Page Table

---

- One-level page table can consume lots of memory for page table
  - e.g., 32-bit logical address space and 4KB page size
    - page table would have 1 million entries ( $2^{32} / 2^{12}$ )
    - if each entry is 4 bytes  $\rightarrow$  4 MB of memory for page table alone
  - each process requires its own page table
  - page table must be **physically contiguous**

# One-Level Page Table



# ELF binary basics

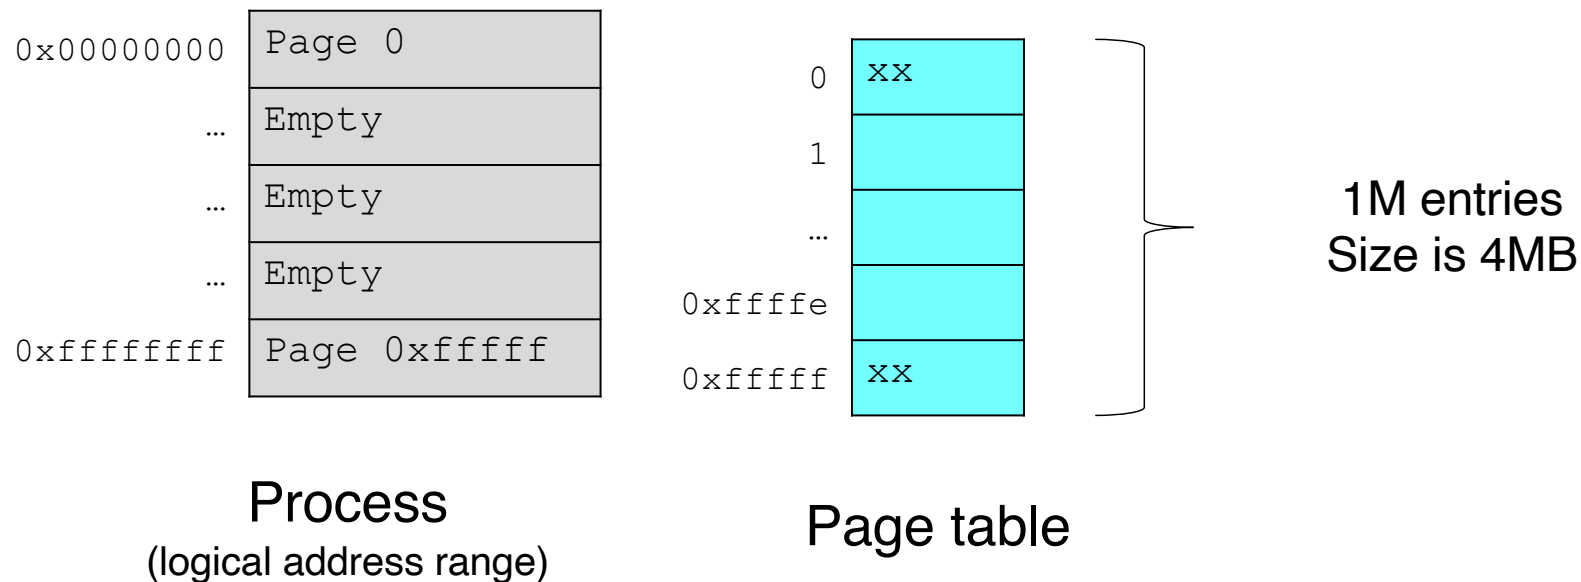
---

- Logical addresses have holes

```
wenbo@wenbo-ThinkPad:~$ cat /proc/self/maps
559a61b8d000-559a61b95000 r-xp 00000000 08:05 3145753 /bin/cat
559a61d94000-559a61d95000 r--p 00007000 08:05 3145753 /bin/cat
559a61d95000-559a61d96000 rw-p 00008000 08:05 3145753 /bin/cat
559a63860000-559a63881000 rw-p 00000000 00:00 0 [heap]
7f8fc2690000-7f8fc305f000 r--p 00000000 08:05 4725339 /usr/lib/locale/locale-archive
7f8fc305f000-7f8fc3246000 r-xp 00000000 08:05 2626368 /lib/x86_64-linux-gnu/libc-2.27.so
7f8fc3246000-7f8fc3446000 ---p 001e7000 08:05 2626368 /lib/x86_64-linux-gnu/libc-2.27.so
7f8fc3446000-7f8fc344a000 r--p 001e7000 08:05 2626368 /lib/x86_64-linux-gnu/libc-2.27.so
7f8fc344a000-7f8fc344c000 rw-p 001eb000 08:05 2626368 /lib/x86_64-linux-gnu/libc-2.27.so
7f8fc344c000-7f8fc3450000 rw-p 00000000 00:00 0
7f8fc3450000-7f8fc3477000 r-xp 00000000 08:05 2626340 /lib/x86_64-linux-gnu/ld-2.27.so
7f8fc3635000-7f8fc3659000 rw-p 00000000 00:00 0
7f8fc3677000-7f8fc3678000 r--p 00027000 08:05 2626340 /lib/x86_64-linux-gnu/ld-2.27.so
7f8fc3678000-7f8fc3679000 rw-p 00028000 08:05 2626340 /lib/x86_64-linux-gnu/ld-2.27.so
7f8fc3679000-7f8fc367a000 rw-p 00000000 00:00 0
7fffd416e0000-7fffd41701000 rw-p 00000000 00:00 0 [stack]
7fffd4178b000-7fffd4178e000 r--p 00000000 00:00 0 [vvar]
7fffd4178e000-7fffd41790000 r-xp 00000000 00:00 0 [vdso]
fffffffff600000-fffffffff601000 r-xp 00000000 00:00 0 [vsyscall]
```

# One-Level Page Table

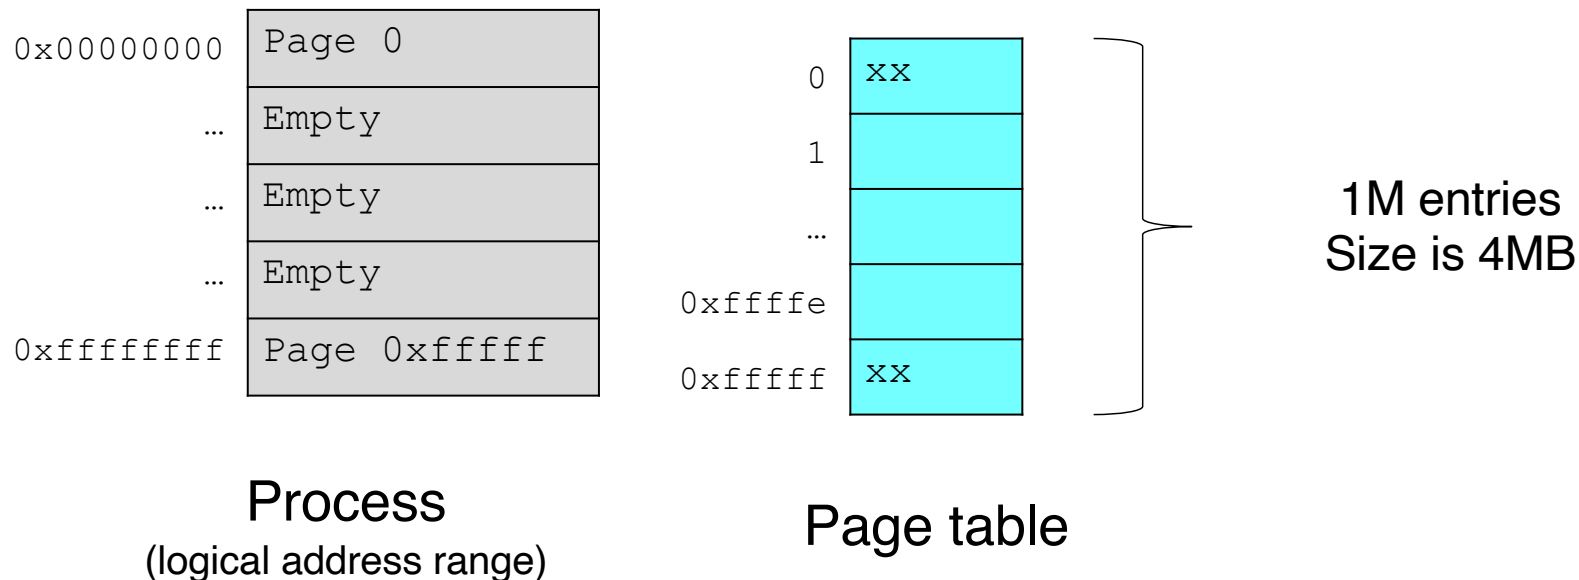
- Worst case for 32-bit
  - Access first page and last page
  - Page table needs to have 0x100000 entries
    - Size is 4MB, physically continuous
      - 1K=1024=0x400
      - 4K=4096=0x1000
      - 1M=1024\*1024=0x100000



# One-Level Page Table

---


- Can we break down page table into pages?
  - Page size is 4KB, how many entries can fit in one page?



# One-Level Page Table

---

- Can we break down page table into pages?
  - Page size is 4KB, how many entries can fit in one page?
    - 1K entries / page
    - Can break down to 1K pages -> Can be indexed by 1K entries -> Put into another page



Page table


# Hierarchical Page Tables

---

- Break up the logical address space into **multiple-level** of page tables
  - e.g., two-level page table
  - first-level page table contains the **frame#** for second-level page tables
    - “page” the page table
- Why hierarchical page table can save memory for page table?




# Two-Level Page Table



# Two-Level Paging

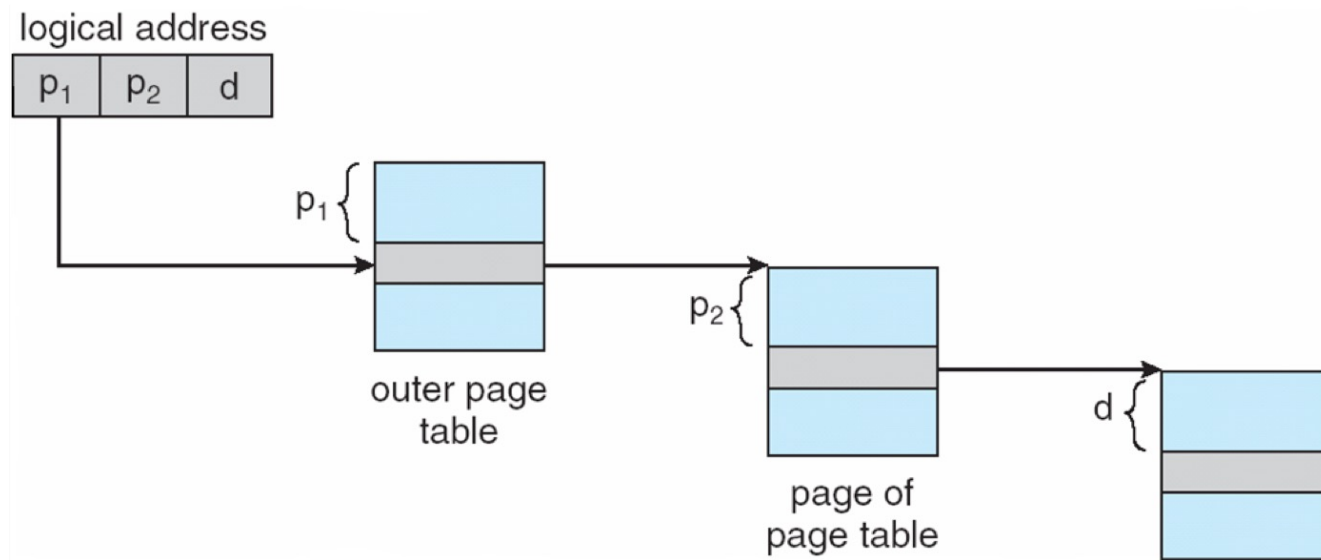
---

- A logical address is divided into:
  - a **page directory number** (first level page table)
  - a **page table number** (2nd level page table)
  - a **page offset**
- Example: 2-level paging in 32-bit Intel CPUs
  - 32-bit address space, 4KB page size
  - 10-bit page directory number, 10-bit page table number
  - each page table entry is 4 bytes, one frame contains 1024 entries ( $2^{10}$ )




# Address-Translation Scheme

---



# Page Table in Linux



\*) 32 bits aligned to a 4-KByte boundary


# 64-bit Logical Address Space

---

- 64-bit logical address space requires more levels of paging
  - two-level paging is not sufficient for 64-bit logical address space
    - if page size is 4 KB ( $2^{12}$ ), outer page table has  $2^{42}$  entries, inner page tables have  $2^{10}$  4-byte entries
  - one solution is to add more levels of page tables
    - e.g., three levels of paging: 1st level page table is  $2^{32}$  bytes in size
    - and possibly 4 memory accesses to get to one physical memory location
- **usually not support full 64-bit virtual address space**
  - AMD-64 supports 48-bit
  - ARM64 supports 39-bit, 48-bit

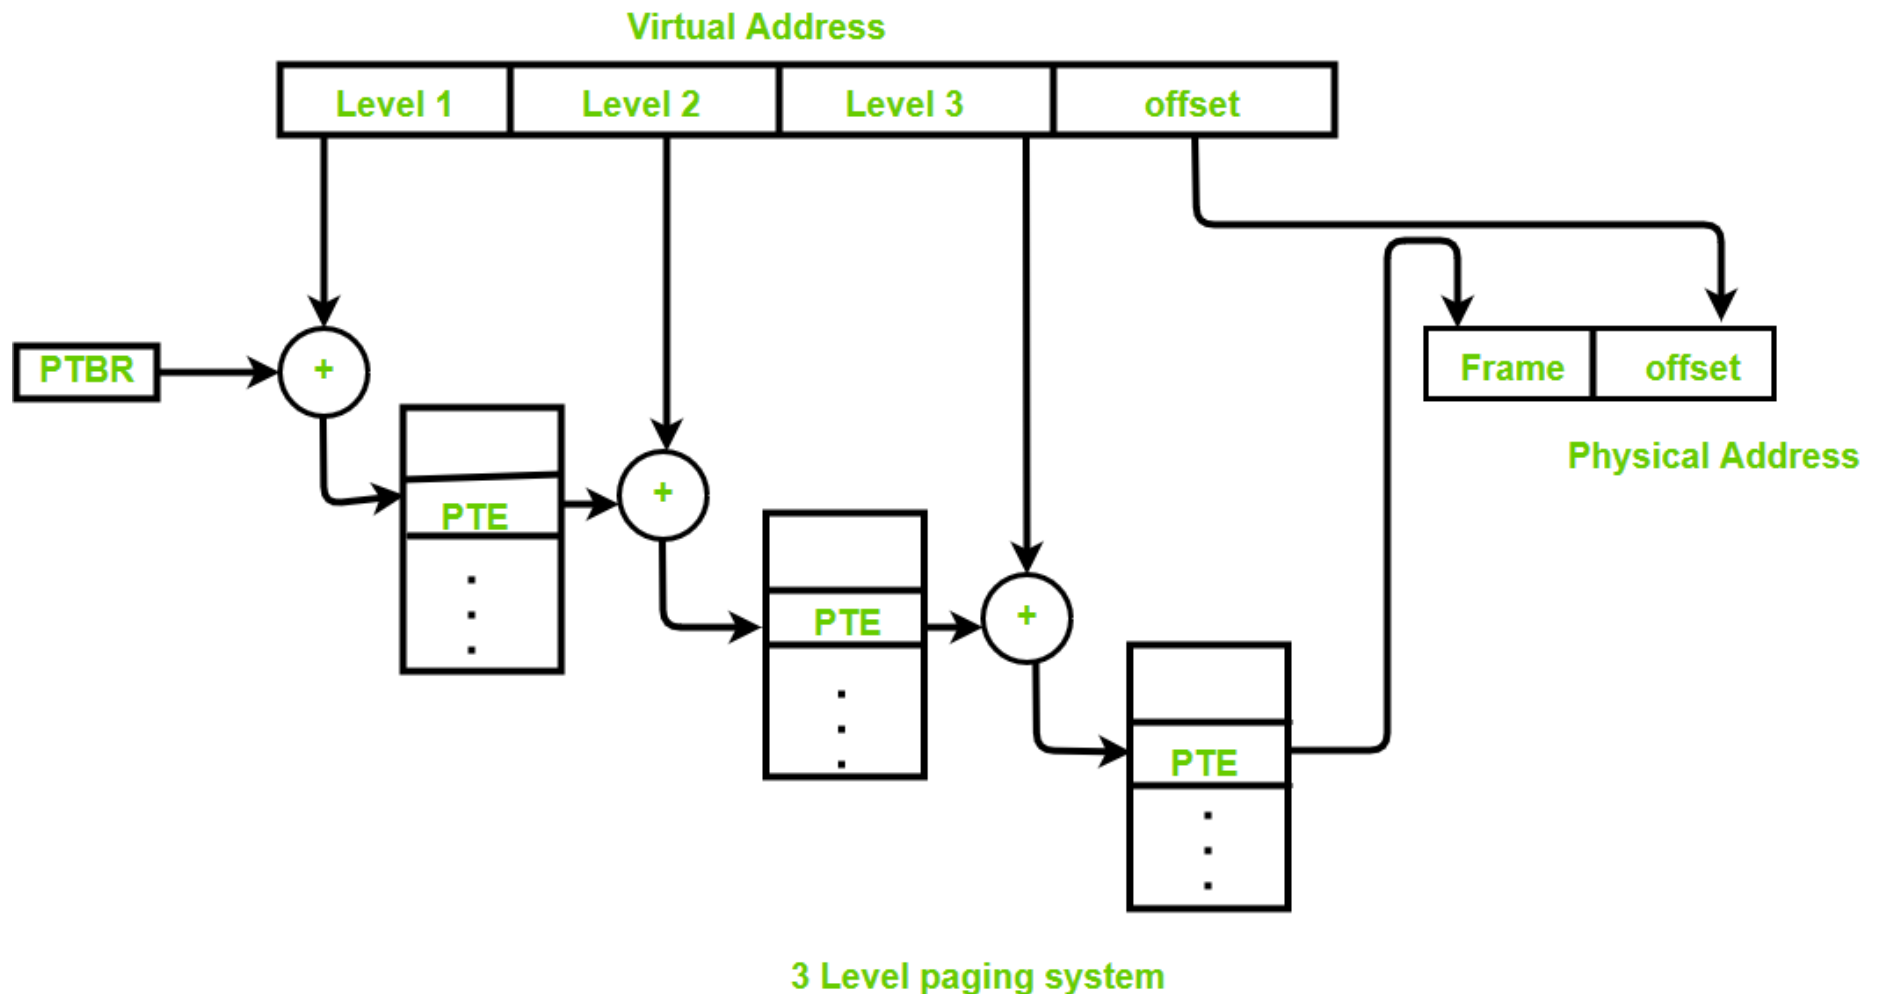
# 64-bit Logical Address Space

- ARM64: 39 bits = 9+9+9+12



# 64-bit Logical Address Space

- ARM64: 39 bits = 9+9+9+12



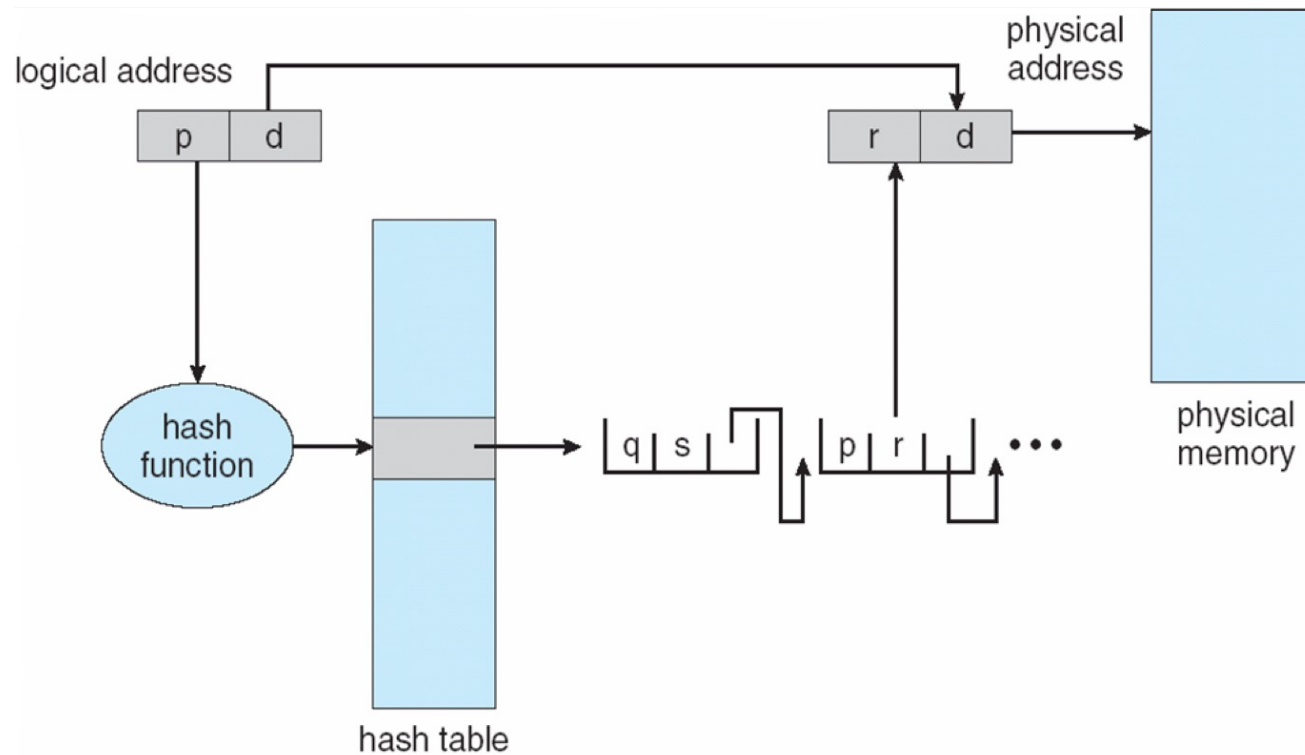
# Hashed Page Tables

---

- In hashed page table, virtual page# is hashed into a frame#
  - the page table contains a chain of elements hashing to the same location
  - each element contains: **page#**, **frame#**, and a **pointer to the next element (resolving conflict)**
    - virtual page numbers are compared in this chain searching for a match
    - if a match is found, the corresponding frame# is returned
- Hashed page table can be used in address spaces > 32 bits
- **Clustered page tables**
  - Each entry refers to several pages



# Hashed Page Table



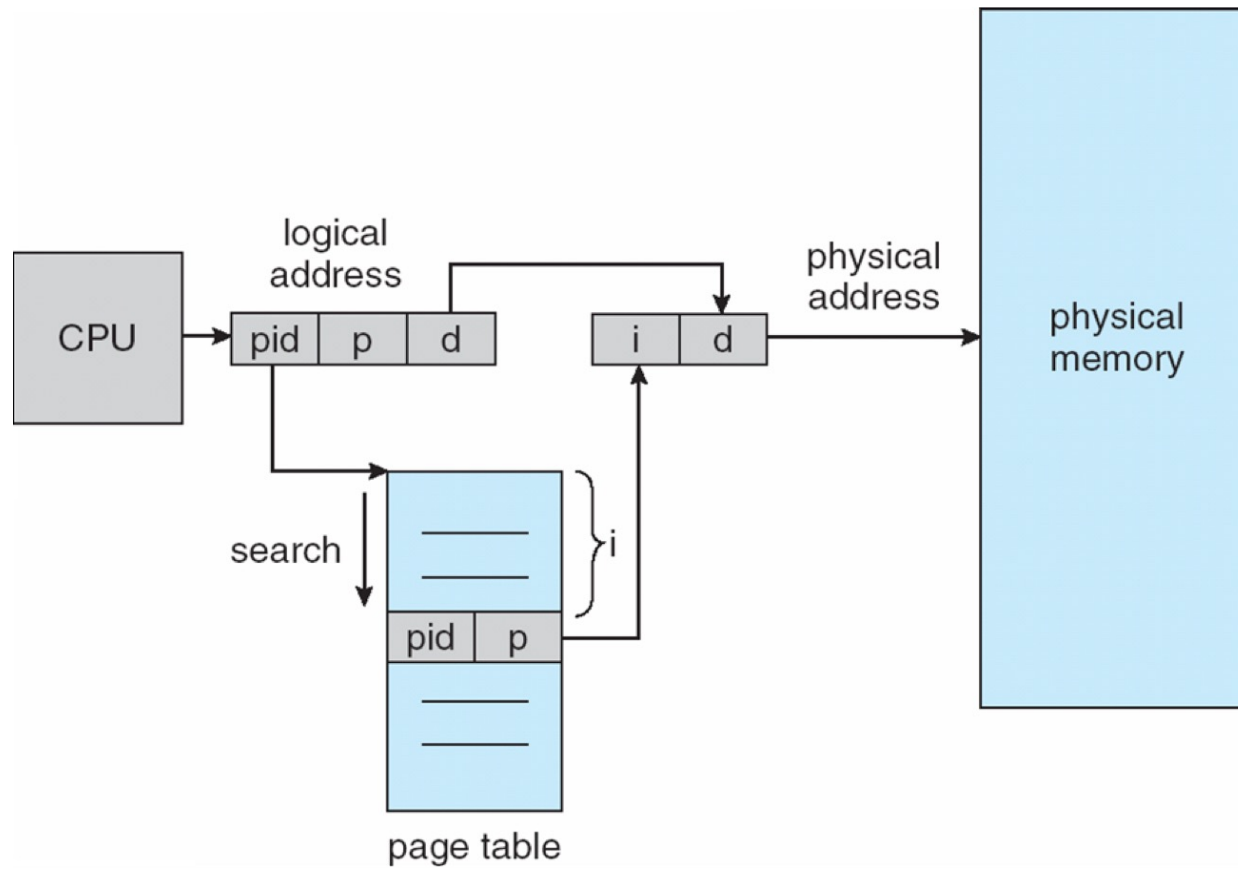
# Inverted Page Table

---

- Inverted page table tracks allocation of physical frame to a process
  - One entry for each physical frame → fixed amount of memory for page table
    - Physical memory is usually much smaller than virtual memory
      - Why?
  - Each entry has the **process id** and the **page#** (virtual address)
- Sounds like a brilliant idea?
  - to translate a virtual address, it is necessary to search the (**whole**) page table
    - can use TLB to accelerate access, TLB miss could be very expensive
  - how to implement shared memory?
    - a physical frame can only be mapped into one process!
    - Because one physical memory page cannot have multiple virtual page entry!

# Inverted Page Table

---




# Swapping

---

- **Swapping** extends physical memory with backing disks
  - A process can be swapped temporarily out of memory to a backing store
    - Backing store is usually a (fast) disk
  - The process will be brought back into memory for continued execution
    - Does the process need to be swapped back in to same physical address?
- Swapping is usually only initiated under memory pressure
- Context switch time can become very high due to swapping
  - If the next process to be run is not in memory, need to swap it in
  - Disk I/O has high latency

# Swapping

---



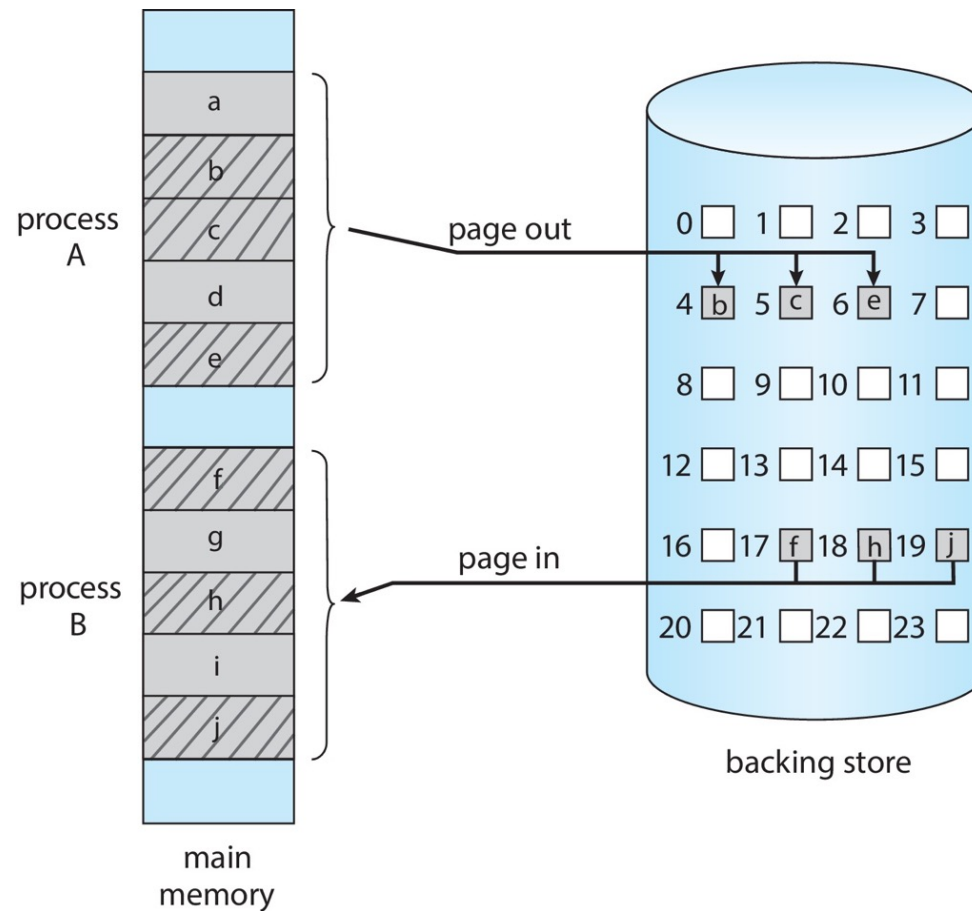
# Context Switch Time including Swapping

---

- If next processes to be put on CPU is not in memory, need to swap out a process and swap in target process
- Context switch time can then be **very high**
- 100MB process swapping to hard disk with transfer rate of 50MB/sec
  - Swap out time of 2,000 ms
  - Plus swap in of same sized process
  - Total context switch swapping component time of 4000ms (4 seconds)
- Can reduce if reduce size of memory swapped - by knowing how much memory really being used

# Swapping with Paging

- Swap pages instead of entire process



# Swapping on Mobile Systems

---

- Not typically supported
  - Flash memory based
    - Small amount of space
    - Limited number of write cycles
    - Poor throughput between flash memory and CPU on mobile platform
- Instead use other methods to free memory if low
  - iOS **asks** apps to voluntarily relinquish allocated memory
    - Read-only data thrown out and reloaded from flash if needed
    - Failure to free can result in termination
  - Android terminates apps if low free memory, but first writes **application state** to flash for fast restart



## Example: The Intel 32 and 64-bit Architectures

---

- Dominant industry chips
- Pentium CPUs are 32-bit and called IA-32 architecture
- Current Intel CPUs are 64-bit and called IA-64 architecture
- Many variations in the chips, cover the main ideas here

# Example: The Intel IA-32 Architecture


---

- Supports both segmentation and segmentation with paging
  - Each segment can be 4 GB
  - Up to 16 K segments per process
  - Divided into two partitions
  - First partition of up to 8 K segments are private to process (kept in **local descriptor table (LDT)**)
  - Second partition of up to 8K segments shared among all processes (kept in **global descriptor table (GDT)**)

# Example: The Intel IA-32 Architecture

---

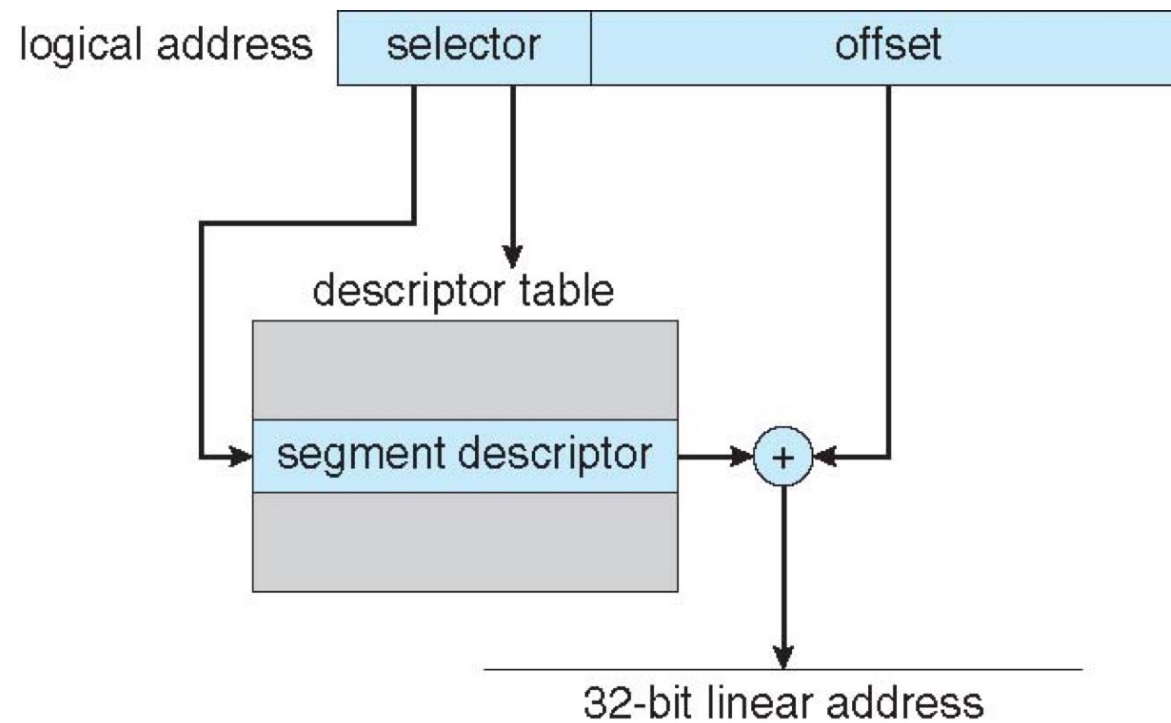
- CPU generates logical address
  - Selector given to segmentation unit
    - Which produces linear addresses



- *s* → segment number, *g* → local/global, *p* → protection
- Segment address is stored in the segment registers
  - CS: segment address of the code segment of the currently executing instruction
  - DS: many data segments, can be one data segment address
- GDTR, LDTR → base address of the descriptor table
- descriptor: base, limit and other bits

# Intel IA-32 Segmentation

---




# Example: The Intel IA-32 Architecture

---

- Linear address given to paging unit
  - Which generates physical address in main memory
  - Paging units form equivalent of MMU
  - Pages sizes can be 4 KB or 4 MB

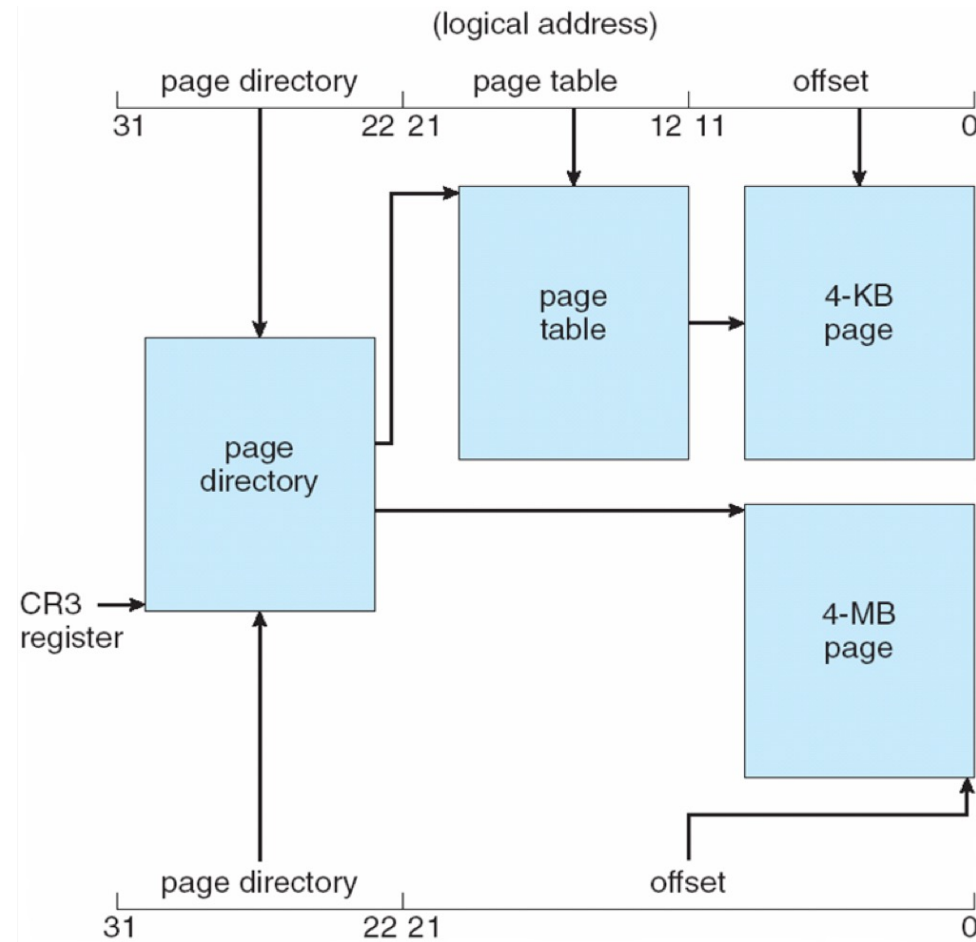
# Logical to Physical Address Translation in IA-32

---




page number		page offset
$p_1$	$p_2$	$d$
10	10	12

# Intel IA-32 Paging Architecture



# Intel IA-32 Page Address Extensions

- 32-bit address limits led Intel to create **page address extension (PAE)**, allowing 32-bit apps access to more than 4GB of memory space
  - Paging went to a 3-level scheme
  - Top two bits refer to a **page directory pointer table**
  - Page-directory and page-table entries moved to 64 bits in size
  - Net effect is increasing address space to 36 bits - 64GB of **physical** memory
    - Virtual address space is still 4GB
    - Page table entries is 64 bits, can index 36 bits physical address (64GB)



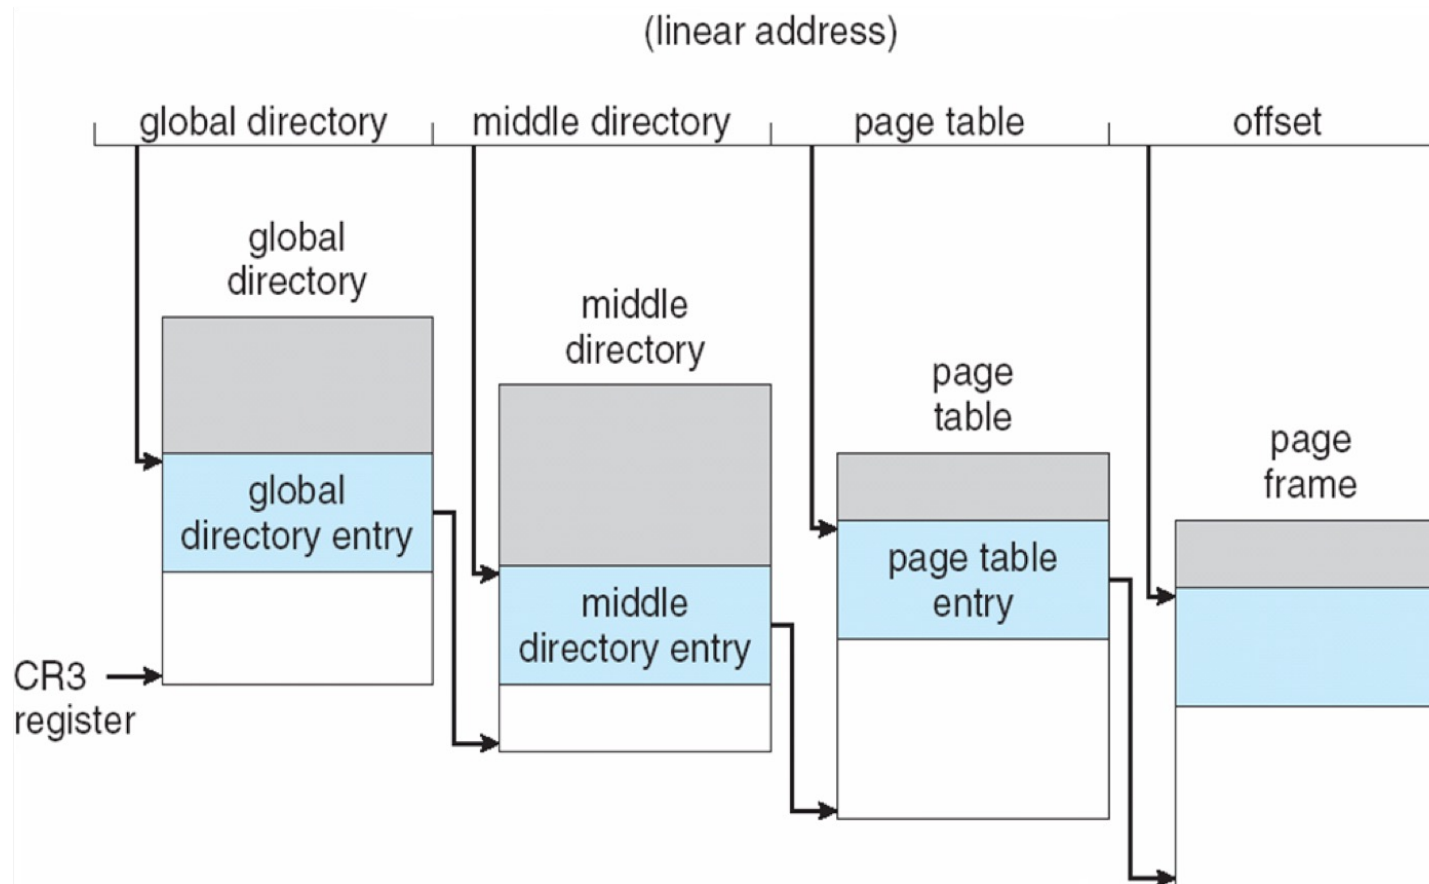


# Linux Support for Intel Pentium

---

- Linux uses only 6 segments
  - kernel code, kernel data, user code, user data
  - task-state segment (TSS), default LDT segment
- Linux only uses two of four possible modes
  - kernel: ring 0, user space: ring 3
- Uses a generic four-level paging for 32-bit and 64-bit systems
  - Two-level paging, middle and upper directories are omitted
  - Three-level page, upper directories are omitted


# Three-level Paging in Linux



# Intel x86-64

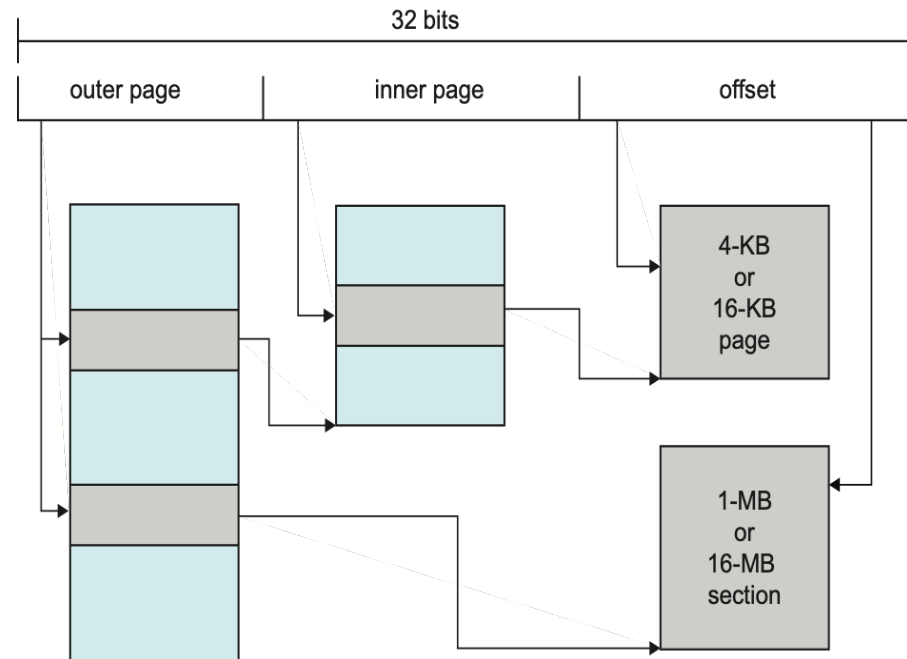
---

- Current generation Intel x86 architecture
- 64 bits is ginormous (> 16 exabytes)
- In practice only implement 48 bit addressing
  - Page sizes of 4 KB, 2 MB, 1 GB
  - **Four levels** of paging hierarchy
- Can also use PAE so virtual addresses are 48 bits and physical addresses are 52 bits
  - Usually no needs



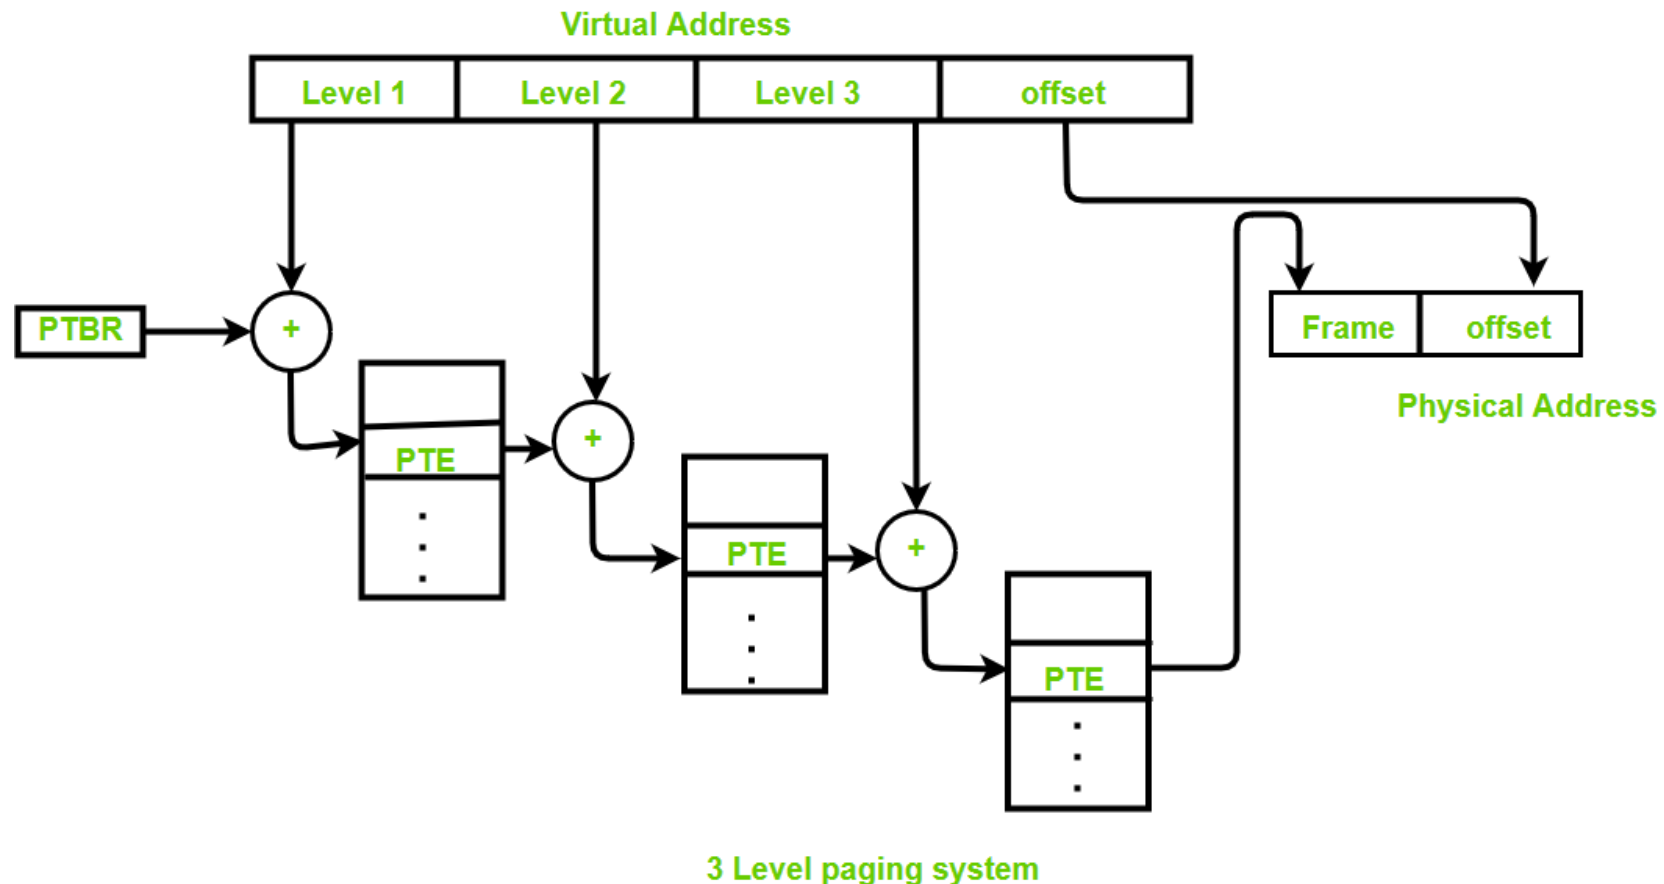
# Example: ARM32 Architecture

- Energy efficient, **32-bit CPU**
- 4 KB and 16 KB pages
- 1 MB and 16 MB pages (termed sections)
- One-level paging for sections, two-level for smaller pages
- Two levels of TLBs
  - Outer level has **two micro TLBs** (one data, one instruction)
  - Inner is single **main TLB**
  - First inner is checked, on miss outers are checked, and on miss page table walk performed by CPU



# Example: ARM64 Architecture

- Dominant mobile platform chip (Apple iOS and Google Android devices for example)
- Support
  - 39 bits addressing = 9+9+9+12, three-level page table
  - 48 bits addressing = 9+9+9+9+12, four-level page table



# Takeaway

---

- Contiguous allocation
  - Fixed, variable
  - Segmentation
- Fragmentation
  - Internal, external
- MMU
- Paging
  - Page table
    - Hierarchical, hashed page table, inverted
    - Two-level, three-level, four-level
    - For 32 bits and 64 bits architectures

# Page table quiz

---

- 1) In 32-bit architecture, for 1-level page table, how large is the whole page table?
- 2) In 32-bit architecture, for 2-level page table, how large is the whole page table?
  - 1) How large for the 1<sup>st</sup> level PGT?
  - 2) How large for the 2<sup>nd</sup> level PGT?
- 3) Why can 2-level PGT save memory?
- 4) 2-level page table walk example
  - 1) Page table base register holds 0x0061 9000
  - 2) Virtual address is 0xf201 5202
  - 3) Page table base register holds 0x1051 4000
  - 4) Virtual address is 0x2190 7010

# Page table quiz


---

- How about page size is 64KB
  - What is the virtual address format for 32-bit?
  - What is the virtual address format for 64-bit?
    - 39-bit VA
    - 48-bit VA



# Virtual address format

- 48-bit VA with 4KB page



# Virtual address format

- 48-bit VA with 64KB page

