

# Virtual Memory



Operating Systems  
Wenbo Shen

# Review

---

- Logical vs physical address
- Memory allocation
  - Contiguous allocation: first-, best-, worst-fit
  - Fragmentation: external vs internal
- paging: page number + page offset
  - Hierarchical page table, hashed page table, inverted page table
  - 1-level vs 2-level, why save memory, page table walk
- Swapping
- MMU
  - TLB

# Outline

---

- Demanding paging
- Copy-on-write
- Page replacement algorithm
  - FIFO, optimal, LRU, ...
- Allocation of frames
- Thrashing
- Examples

# Background

---

- Code needs to be in memory to execute, but entire program **rarely** needed or used at the same time
  - **unused code**: error handling code, unusual routines
  - **unused data**: large data structures
- Consider ability to execute **partially-loaded program**
  - program no longer constrained by limits of physical memory
  - programs could be larger than physical memory


# Background

---

- Virtual memory: separation of **logical memory** from **physical memory**
  - only part of the program needs to be in memory for execution
    - logical address space can be much larger than physical address space
    - more programs can run concurrently
    - less I/O needed to load or swap processes (part of it)
  - allows memory (e.g., shared library) to be shared by several processes
    - better IPC performance
  - allows for more efficient process **forking** (copy-on-write)
- Virtual memory can be implemented via:
  - **Paging**

# Virtual Memory Larger Than Physical Memory


- Virtual memory is larger than physical memory
  - Virtual memory is just address ranges, all storage is backed by physical memory



# Virtual-address Space


---

- Usually design virtual address space for stack to start at **Max logical address and grow "down" while heap grows "up"**
  - Maximizes address space use
  - Unused address space between the two is hole
    - No physical memory needed until heap or stack grows to a given new page
- Enables **sparse** address spaces with holes left for growth, dynamically linked libraries, etc
- System libraries shared via mapping into virtual address space
- Pages can be shared during fork(), speeding process creation: COW



# Shared Library Using Virtual Memory

---





# Shared Library Using Virtual Memory

```
wenbo@parallels: ~  
wenbo@parallels: ~ 107x30  
7ffc75a5f000-7ffc75a80000 rw-p 00000000 00:00 0 [stack]  
7ffc75aa7000-7ffc75aaa000 r--p 00000000 00:00 0 [vvar]  
7ffc75aaa000-7ffc75aac000 r-xp 00000000 00:00 0 [vdso]  
fffffffffff600000-fffffffffff601000 r-xp 00000000 00:00 0 [vsyscall]  
wenbo@parallels:~$ which cat  
/bin/cat  
wenbo@parallels:~$ file /bin/cat  
/bin/cat: ELF 64-bit LSB shared object, x86-64, version 1 (SYSV), dynamically linked, interpreter /lib64/l,  
for GNU/Linux 3.2.0, BuildID[sha1]=747e524bc20d33ce25ed4aea108e3025e5c3b78f, stripped  
wenbo@parallels:~$ cat /proc/self/maps  
55b793b79000-55b793b81000 r-xp 00000000 08:01 1048601 /bin/cat  
55b793d80000-55b793d81000 r--p 00007000 08:01 1048601 /bin/cat  
55b793d81000-55b793d82000 rw-p 00008000 08:01 1048601 /bin/cat  
55b794d33000-55b794d54000 rw-p 00000000 00:00 0 [heap]  
7f1974b90000-7f197555f000 r--p 00000000 08:01 662494 /usr/lib/locale/locale-archive  
7f197555f000-7f1975746000 r-xp 00000000 08:01 267596 /lib/x86_64-linux-gnu/libc-2.27.so  
7f1975746000-7f1975946000 ---p 001e7000 08:01 267596 /lib/x86_64-linux-gnu/libc-2.27.so  
7f1975946000-7f197594a000 r--p 001e7000 08:01 267596 /lib/x86_64-linux-gnu/libc-2.27.so  
7f197594a000-7f197594c000 rw-p 001eb000 08:01 267596 /lib/x86_64-linux-gnu/libc-2.27.so  
7f197594c000-7f1975950000 rw-p 00000000 00:00 0  
7f1975950000-7f1975977000 r-xp 00000000 08:01 267568 /lib/x86_64-linux-gnu/ld-2.27.so  
7f1975b3c000-7f1975b60000 rw-p 00000000 00:00 0  
7f1975b77000-7f1975b78000 r--p 00027000 08:01 267568 /lib/x86_64-linux-gnu/ld-2.27.so  
7f1975b78000-7f1975b79000 rw-p 00028000 08:01 267568 /lib/x86_64-linux-gnu/ld-2.27.so  
7f1975b79000-7f1975b7a000 rw-p 00000000 00:00 0  
7ffc73010000-7ffc73031000 rw-p 00000000 00:00 0 [stack]  
7ffc73148000-7ffc7314b000 r--p 00000000 00:00 0 [vvar]  
7ffc7314b000-7ffc7314d000 r-xp 00000000 00:00 0 [vdso]  
fffffffffff600000-fffffffffff601000 r-xp 00000000 00:00 0 [vsyscall]  
wenbo@parallels:~$
```

# Demand Paging

---

- **Demand paging** brings a **page** into memory only when it is **demanded**
  - demand means access (read/write)
  - if page is invalid (error)  $\Rightarrow$  abort the operation
  - if page is valid but not in memory  $\Rightarrow$  bring it to memory
    - Memory here means **physical** memory
    - This is called **page fault**
    - via swapping for swapped pages
    - via mapping for new page
    - no unnecessary I/O, less memory needed, slower response, more apps

# Demand Paging


---

- Demand paging vs page fault
  - What is the relationship?
- What causes page fault ?
  - User space program accesses an address
- Which hardware issues page fault?
  - MMU
- Who handles page fault ?
  - Operating system

# What causes page fault?


- User space program accesses an address
  - Will kernel code cause page fault?

1. Program calls `brk()` to grow its heap




2. `brk()` enlarges heap VMA.


New pages are **not** mapped onto physical memory.



3. Program tries to access new memory.  
Processor page faults.



4. Kernel assigns page frame to process,  
creates PTE, resumes execution. Program is  
unaware anything happened.



# MMU issues page fault

---

- How does MMU know the physical frame is not mapped?
- Each page table entry has a valid-invalid (present) bit
  - $\underline{v} \Rightarrow$  frame mapped,  $\underline{i} \Rightarrow$  frame not mapped
  - initially, valid-invalid bit is set to  $\underline{i}$  on all entries
  - during address translation, if the entry is invalid, it will trigger a **page fault**
- Example of a page table snapshot:

Frame #	v/i bit
	<b>v</b>
	<b>v</b>
	<b>v</b>
	<b>v</b>
	<b>i</b>
....	
	<b>i</b>
	<b>i</b>

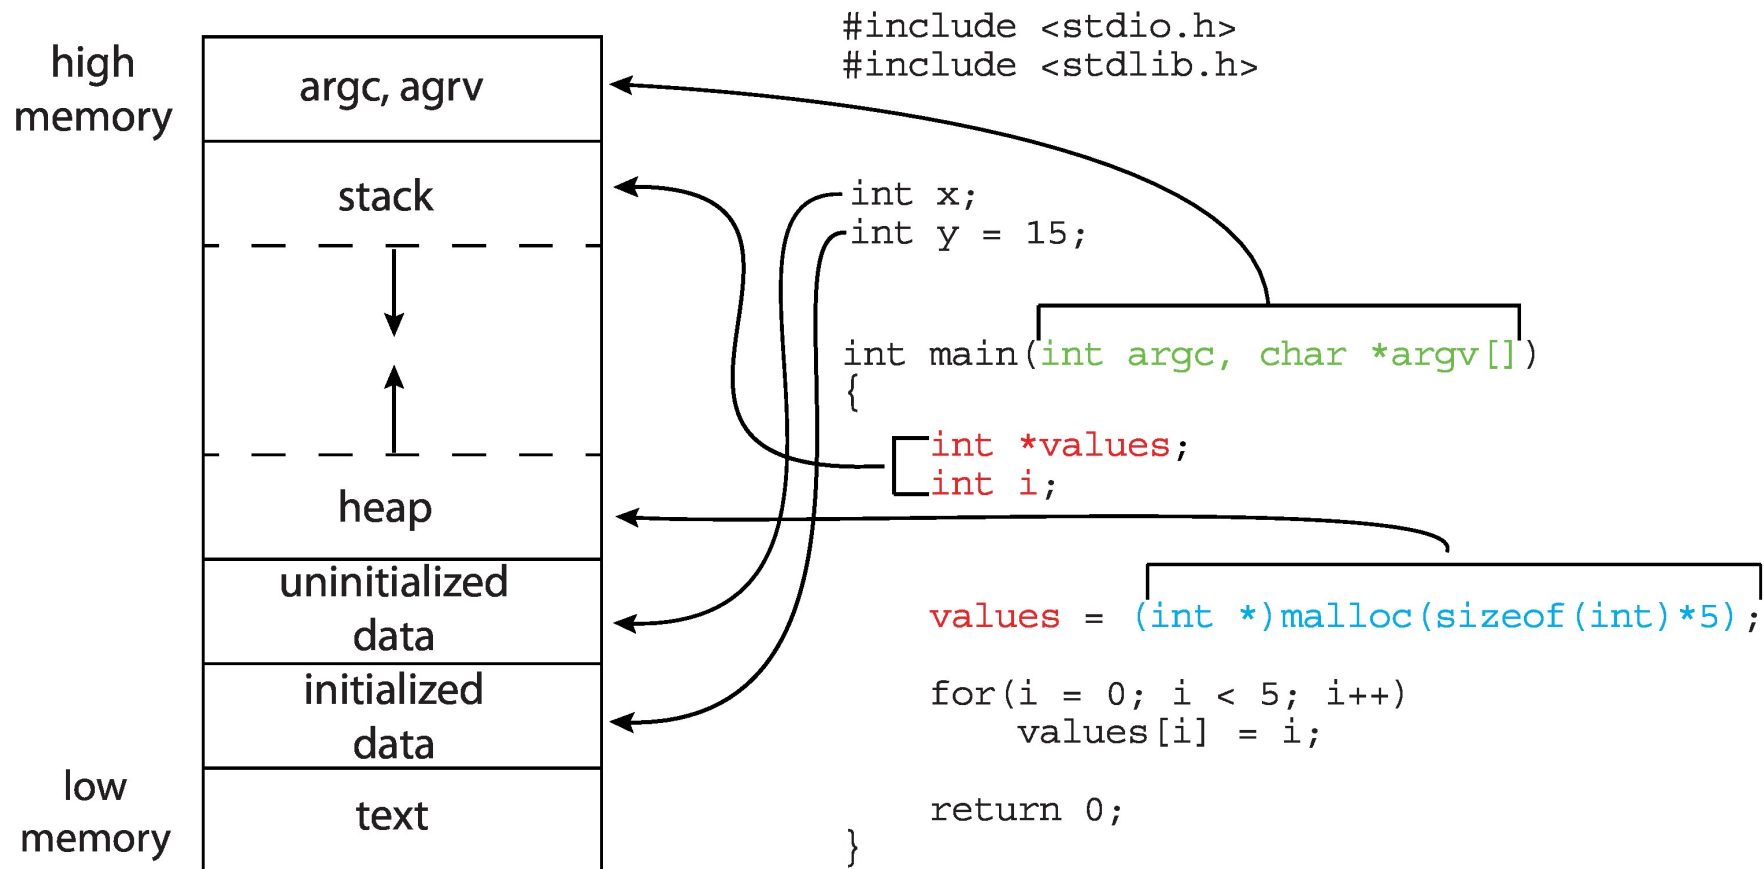
page table

# Who handles page fault?

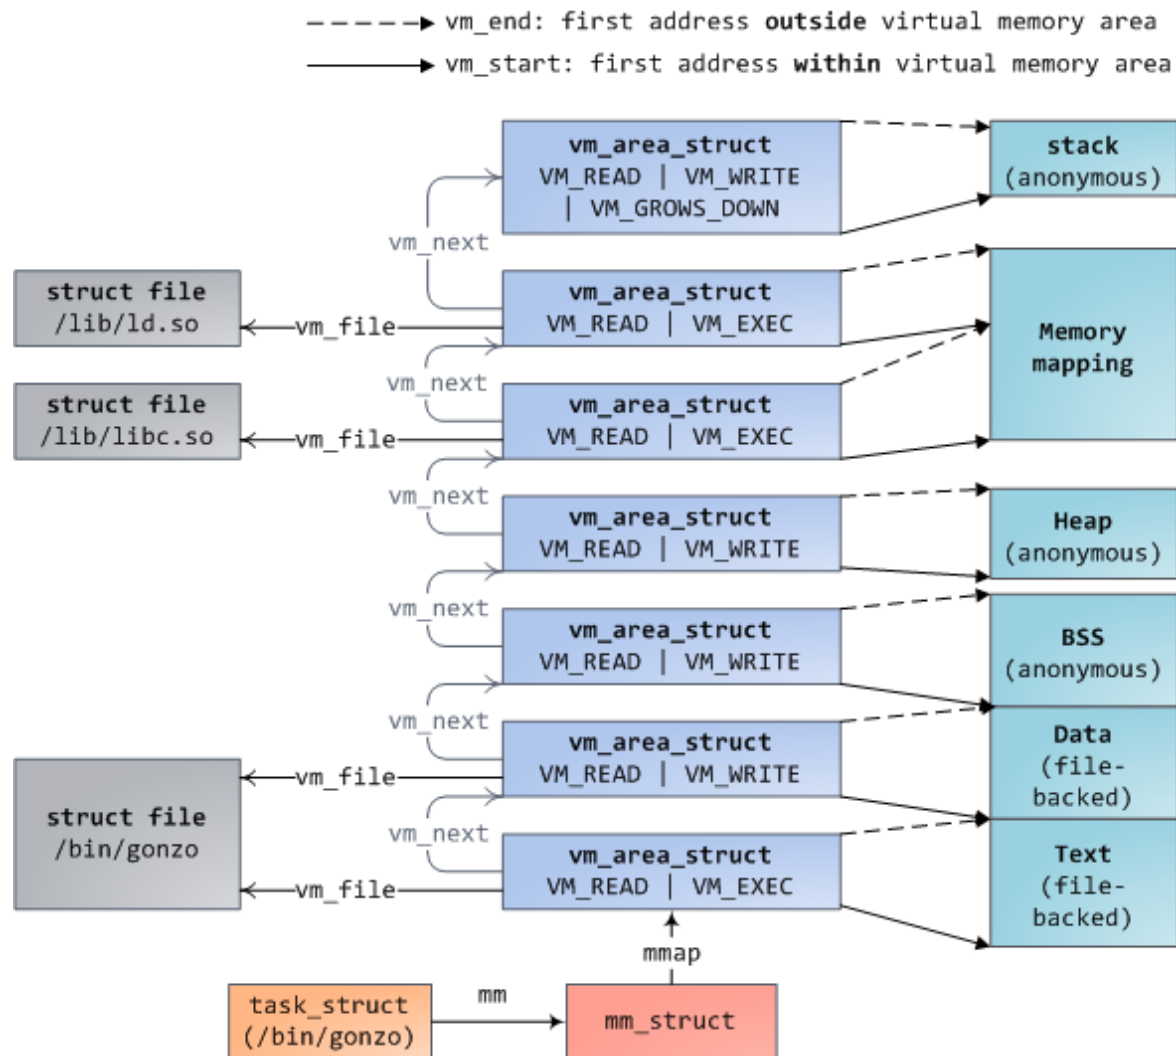
---

- Linux implementation
  - First check vma to decide fault type
    - Address in vm\_area
    - Address out of vm\_area
      - Error, abort
  - Then map physical frame

# Memory Layout of a C Program



# Page Faults



```


305 struct vm_area_struct {
306     /* The first cache line has the info for VMA tree walking. */
307
308     unsigned long vm_start;          /* Our start address within vm_mm. */
309     unsigned long vm_end;           /* The first byte after our end address
310                                     within vm mm. */

```



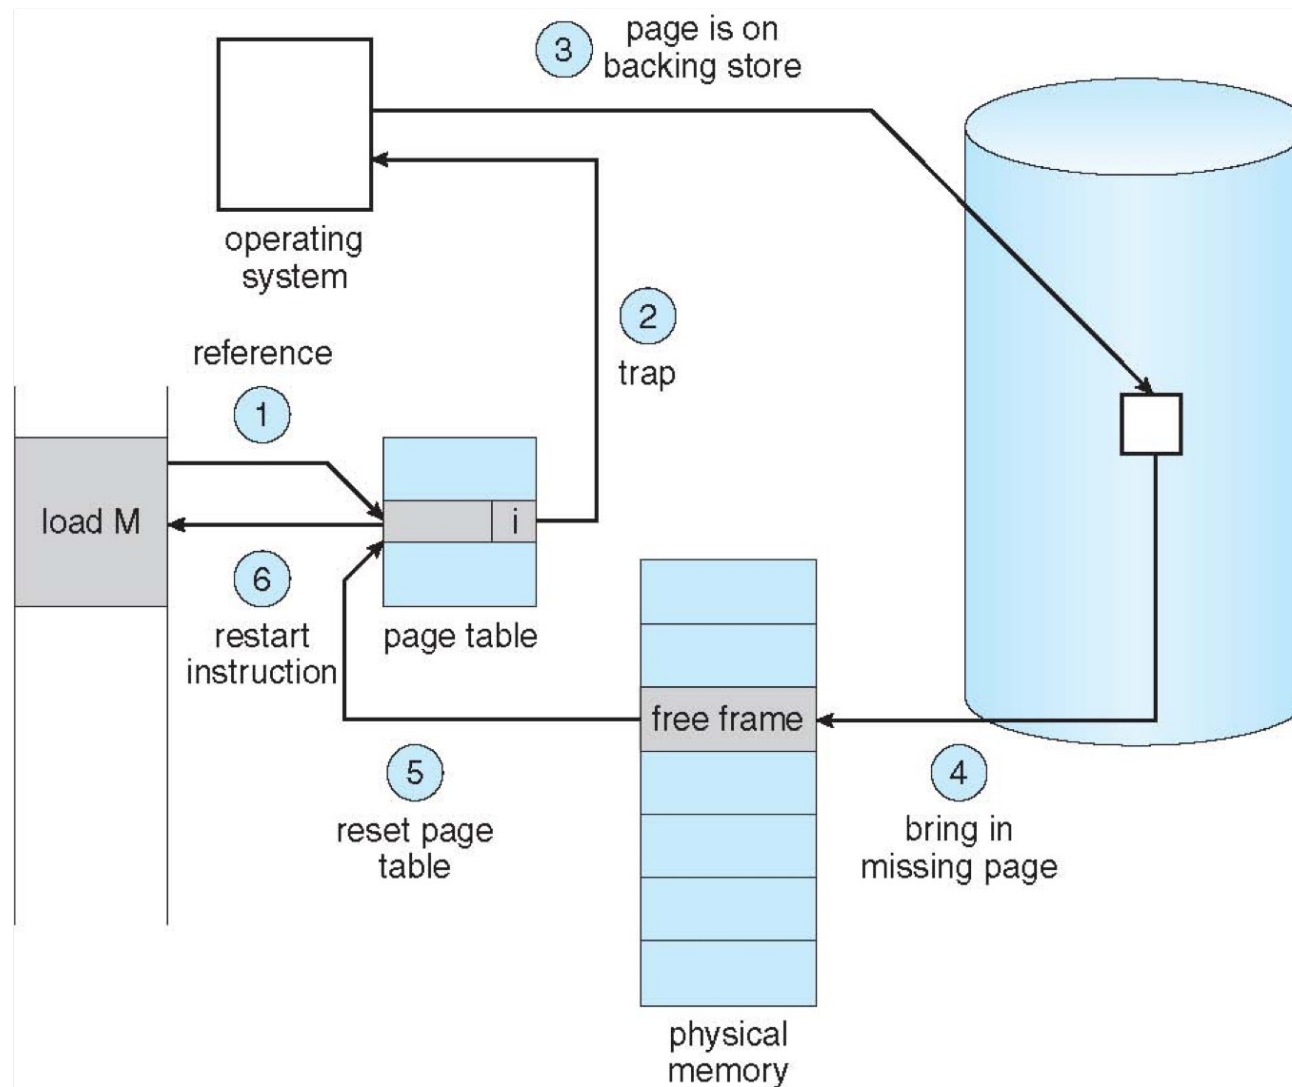
# Page Table (Mem pages are not all in memory)

- Executable file



# Page Fault Handling

- Page fault for code



# Page Fault

---

- First reference to a non-present page will trap to kernel: **page fault**
- Operating system looks at memory mapping to decide:
  - **invalid reference**  $\Rightarrow$  deliver an exception to the process
    - Via check vma in Linux
  - **valid but not in physical memory**  $\Rightarrow$  bring in
    - get an empty physical frame
    - bring page into frame via disk operation
    - set page table entry to indicate the page is now in memory
    - restart the instruction that caused the page fault

# Page Fault - swapper

---

- **Lazy swapper:** never swaps a page in memory unless it will be needed
  - the swapper that deals with pages is also called a pager
- **Pre-Paging:** pre-page all or some of pages a process will need, before they are referenced
  - it can reduce the number of page faults during execution
  - if pre-paged pages are **unused**, I/O and memory was wasted
    - although it reduces page faults, total I/O# likely is higher

# Page Fault

---

- Extreme case: start process with no frames in memory (aka. **pure demand paging**)
  - OS sets instruction pointer to first instruction of process
    - invalid page  $\Rightarrow$  page fault
  - every page is paged in on first access
    - **program locality** reduces the overhead
  - an instruction could access multiple pages  $\Rightarrow$  multiple page faults
    - e.g., instruction, data, and page table entries for them
- Demand paging needs hardware support
  - page table entries with **valid / invalid bit**
  - **backing storage** (usually disks)
  - **instruction restart**

# Page Fault - Get Free Frame

---

- When a page fault occurs, the operating system must bring the desired page from secondary storage into main memory.
- Most operating systems maintain a free-frame list -- a pool of free frames for satisfying such requests.



- Operating system typically allocate free frames using a technique known as **zero-fill-on-demand** -- the content of the frames zeroed-out before being allocated.
- When a system starts up, all available memory is placed on the free-frame list.

# Stages in Demand Paging – Worse Case

---

- 1. Trap to the operating system
- 2. Save the user registers and process state
- 3. Determine that the interrupt was a page fault
  - 3.1 Check that the page reference was legal
- 4. Find a free frame
- 5. Determine the location of the page on the disk, issue a read from the disk to the free frame:
  - 5.1 Wait in a queue for this device until the read request is serviced
  - 5.2 Wait for the device seek and/or latency time
  - 5.3 Begin the transfer of the page to a free frame

# Stages in Demand Paging – Worse Case

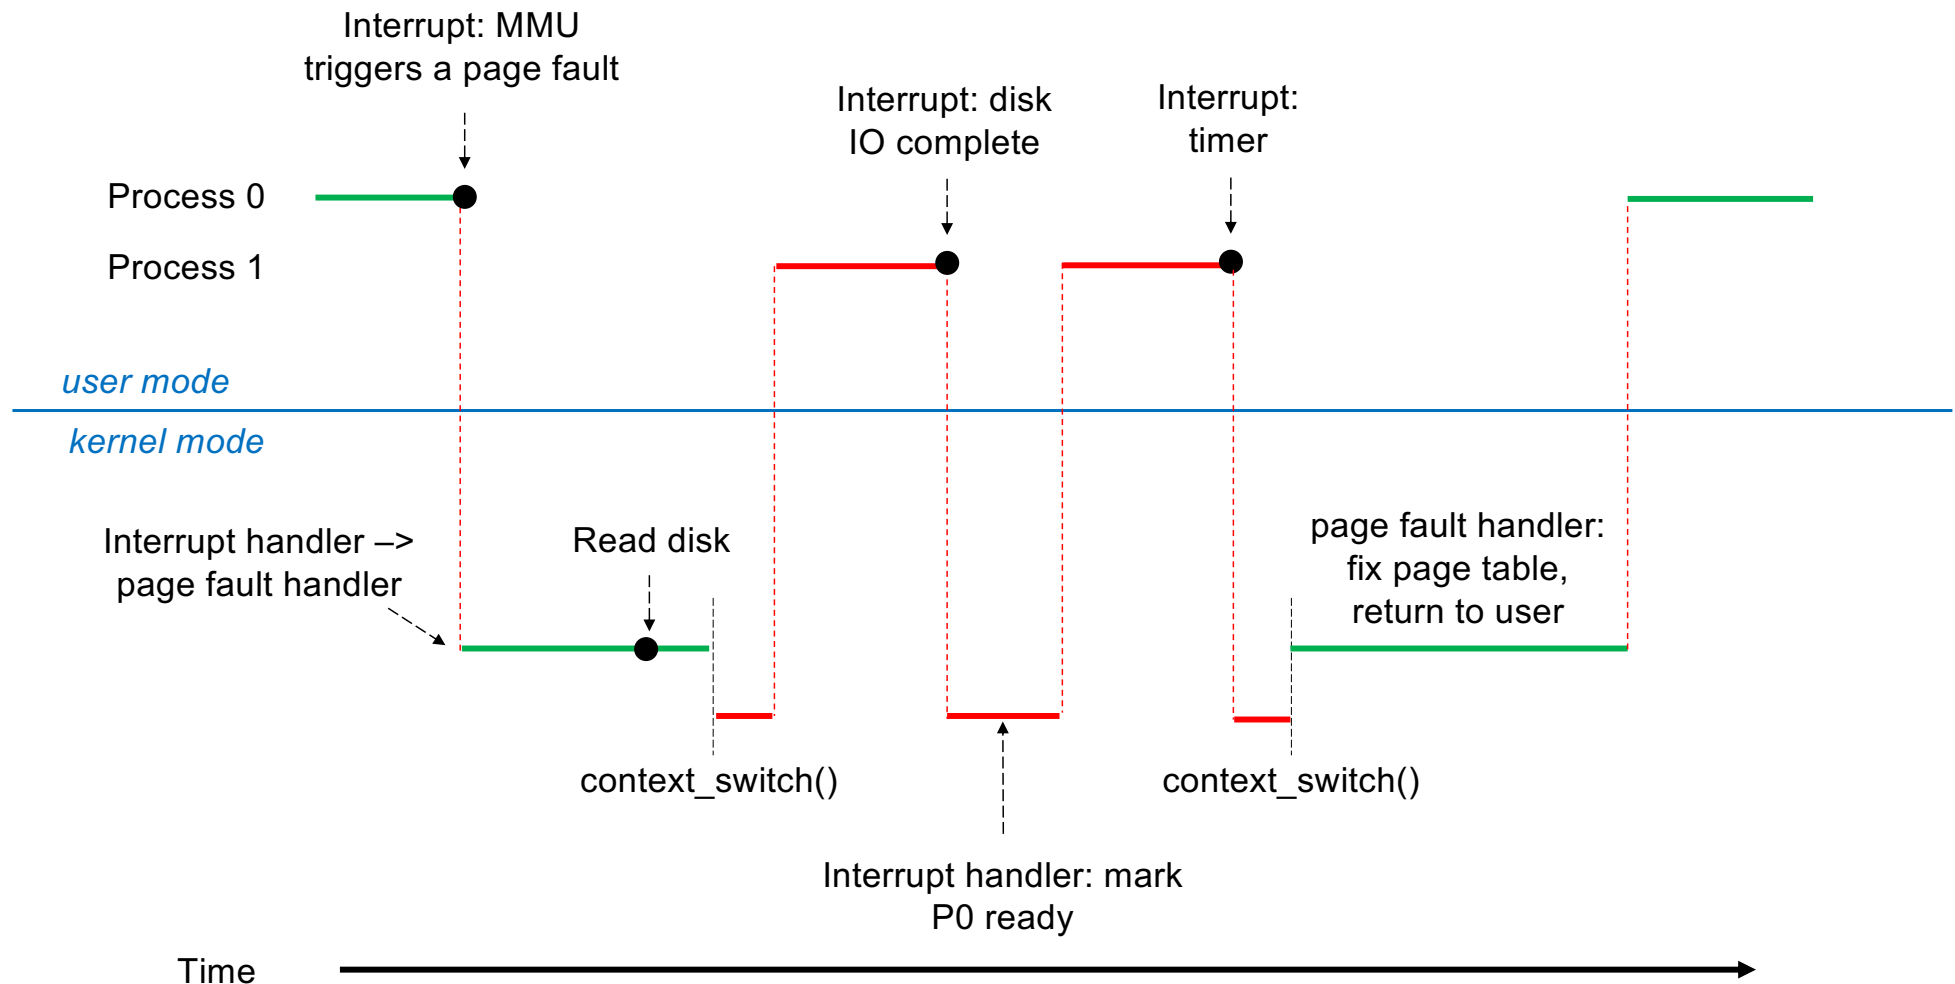
---

- 6. While waiting, allocate the CPU to other process
- 7. Receive an interrupt from the disk I/O subsystem (I/O completed)
  - 7.1 Determine that the interrupt was from the disk
  - 7.2 Mark page fault process ready
- 8. Wait for the CPU to be allocated to this process again
  - 8.1 Save registers and process state for other process
  - 8.2 Context switch to page fault process
- 9. Correct the page table, mapping new frame
- 10. Return to user: restore the user registers, process state, and new page table, and then resume the interrupted instruction



# Stages in Demand Paging - Worse Case

- Assume a page fault happen in Process 0
- While waiting for disk, the CPU is allocated to Process 1



# Demand Paging: EAT

---

- Page fault rate:  $0 \leq p \leq 1$ 
  - if  $p = 0$  no page faults
  - if  $p = 1$ , every reference is a fault
- Effective Access Time (EAT):  
 $(1 - p) \times \text{memory access} + p \times (\text{page fault overhead} + \text{swap page out} + \text{swap page in} + \text{instruction restart overhead})$

# Demand Paging Example

---

- Assume memory access time: 200 nanoseconds, average page-fault service time: 8 milliseconds
- $EAT = (1 - p) \times 200 + p \times (8 \text{ milliseconds})$   
 $= (1 - p) \times 200 + p \times 8,000,000$   
 $= 200 + p \times 7,999,800$
- if one out of 1,000 causes a page fault, then  $EAT = 8.2$  microseconds
  - a slowdown by a factor of 41000
    - $8.2 \text{ ms} / 0.2 \text{ us} = 41000$
- if want < 10 percent, less than one page fault in every 400,000 accesses

# Demand Paging Optimizations

---

- Swap space I/O faster than file system I/O even if on the same device
  - Swap allocated in larger chunks, less management needed than file system
- Copy entire process image from **disk** to **swap space** at process load time
  - Then page in and out of swap space
  - Used in older BSD Unix

# Demand Paging Optimizations

---

- Demand page in from program binary on disk, but **discard** rather than paging out when freeing frame (and reload from disk next time)
  - Following cases still need to **write** to swap space
    - Pages not associated with a file (like stack and heap) - anonymous memory
    - Pages modified in memory but not yet written back to the file system
- Mobile systems
  - Typically don't support swapping
  - Instead, demand page from **file system** and reclaim read-only pages (such as code)


# Copy-on-Write

---

- **Copy-on-write (COW)** allows parent and child processes to initially share the same pages in memory
  - the page is shared as long as no process modifies it
  - if either process modifies a shared page, only then is the page copied
- COW allows more efficient **process creation**
  - no need to copy the parent memory during fork
  - only changed memory will be copied later
- vfork syscall optimizes the case that child calls **exec** immediately after fork
  - parent is suspend until child exits or calls exec
  - child shares the parent resource, including the heap and the stack
    - child cannot return from the function or call exit, should call `_exit`
  - vfork could be fragile, **it is invented when COW has not been implemented**

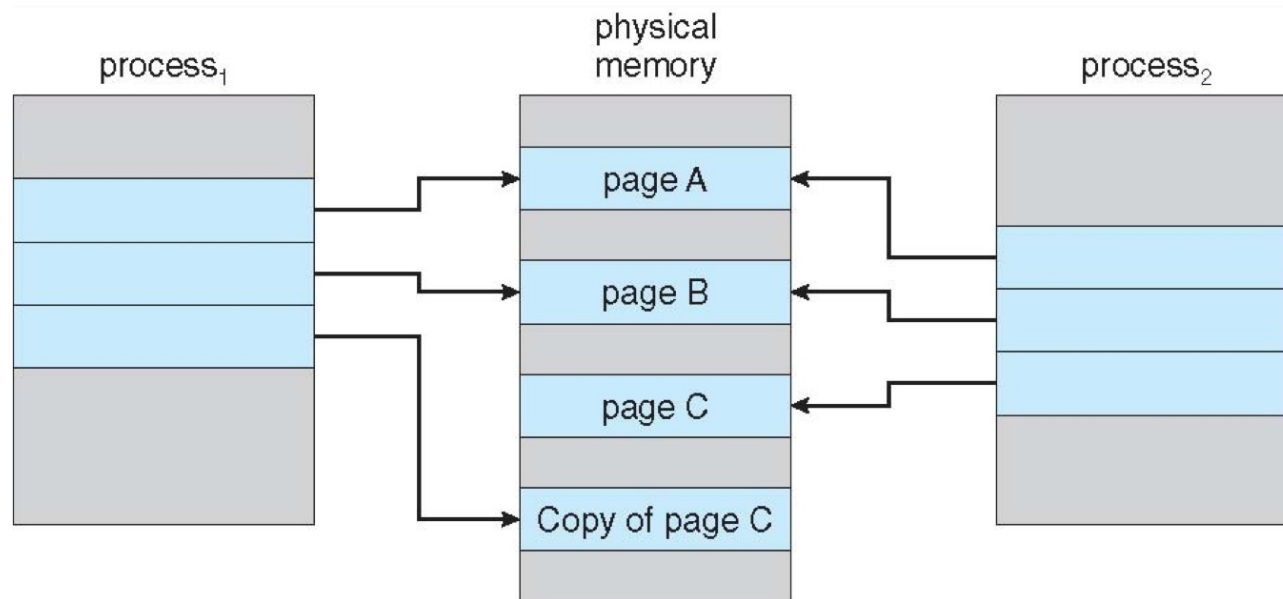
# Before Process 1 Modifies Page C

---



# After Process 1 Modifies Page C

---





# What Happens if There is no Free Frame?

---


- Used up by process pages
- Also in demand from the kernel, I/O buffers, etc
- How much to allocate to each?
- **Page replacement** - find some page in memory, but not really in use, page it out
  - Algorithm - terminate? swap out? replace the page?
  - Performance - want an algorithm which will result in minimum number of page faults
- Same page may be brought into memory several times

# Page Replacement

---

- Memory is an important resource, system may run out of memory
- To prevent out-of-memory, swap out some pages
  - page replacement usually is a part of the page fault handler
  - policies to select victim page require careful design
    - need to reduce overhead and avoid **thrashing**
  - use modified (dirty) bit to reduce number of pages to swap out
    - only modified pages are written to disk
  - select some processes to kill (last resort)
- **Page replacement completes separation between logical memory and physical memory** - large virtual memory can be provided on a smaller physical memory

# Need For Page Replacement




# Page Fault Handler (with Page Replacement)

---

- To **page in** a page:
  - find the location of the desired page on disk
  - find a free frame:
    - if there is a free frame, use it
    - if there is none, use a page replacement policy to pick a victim frame, write victim frame to disk if dirty
  - bring the desired page into the free frame; update the page tables
  - restart the instruction that caused the trap
- Note now potentially **2 page I/O** for one page fault ➡ increase EAT

# Page Replacement




# Page Replacement Algorithms

---

- Page-replacement algorithm should have lowest page-fault rate on both first access and re-access
  - FIFO, optimal, LRU, LFU, MFU...
- To evaluate a page replacement algorithm:
  - run it on a particular string of memory references (reference string)
    - string is just page numbers, not full addresses
  - compute the number of page faults on that string
    - repeated access to the same page does not cause a page fault
  - in all our examples, the reference string is  
7,0,1,2,0,3,0,4,2,3,0,3,0,3,2,1,2,0,1,7,0,1

# Page Faults v.s. Number of Frames


---



# First-In-First-Out (FIFO)

---

- **FIFO**: replace the first page loaded
  - Similar to sliding a window of  $n$  in the reference string
  - Our reference string 7,0,1,2,0,3,0,4,2,3,0,3,0,3,2,1,2,0,1,7,0,1
    - will cause 15 page faults with 3 frames
  - How about reference string of 1,2,3,4,1,2,5,1,2,3,4,5 /w 3 or 4 frames?
- For FIFO, adding **more frames** can cause **more page faults**!
- **Belady's Anomaly**




15 page faults




# FIFO Illustrating Belady's Anomaly

---




1 2 3 4 1 2 5 1 2 3 4 5

# Belady's Anomaly




9 page faults



10 page faults!


# Optimal Algorithm

- **Optimal** : replace page that will not be used for the longest time
  - 9-page-fault is optimal for the example on previous slide
- How do you know which page will not be used for the longest time?
  - can't read the future
  - used for measuring how well your algorithm performs



# Least Recently Used (LRU)

- **LRU** replaces pages that have not been used for the longest time
  - associate time of last use with each page, select pages w/ oldest timestamp
  - generally good algorithm and frequently used
  - 12 faults for our example, better than FIFO but worse than OPT
- LRU and OPT do **NOT** have Belady's Anomaly
- How to implement LRU?
  - **counter-based**
  - **stack-based**




# LRU Implementation

---

- **Counter-based** implementation
  - every page table entry has a counter
  - every time page is referenced, copy the **clock** into the counter
  - when a page needs to be replaced, search for page with smallest counter
    - min-heap can be used
- **Stack-based** implementation
  - keep a stack of page numbers (in double linked list)
  - when a page is referenced, move it to the top of the stack
  - each update is more expensive, but no need to search for replacement

# Stack-based LRU

---



# LRU Approximation Implementation

---

- Counter-based and stack-based LRU have high performance overhead
- Hardware provides a **reference bit**
- LRU approximation with a **reference bit**
  - associate with each page a reference bit, initially set to 0
  - when page is referenced, set the bit to 1 (done by the hardware)
  - replace any page with reference bit = 0 (if one exists)
    - We do not know the **order**, however

# Additional-Reference-Bits Algorithm

---

- Reordering the bits at regular intervals
  - Suppose we have 8-bits byte for each page
  - During a time interval (100ms), shifts bit rights by 1 bit, sets the high bit if used, and then discards the low-order bits
  - 00000000 => has not been used in 8 time intervals
  - 11111111 => has been used in all time intervals
  - 11000100 vs 01110111 : which one is used more recently?





# LRU Implementation

---

- **Second-chance** algorithm
  - Generally FIFO, plus hardware-provided **reference bit**
  - **Clock** replacement
  - If page to be replaced has
    - Reference bit = 0 -> replace it
    - reference bit = 1 then:
      - set reference bit 0, leave page in memory
      - replace next page, subject to same rules

# Second-chance (clock) Page-replacement Algorithm



# Enhanced Second-Chance Algorithm

---

- Improve algorithm by using **reference bit** and **modify bit** (if available) in concert
- Take ordered pair (reference, modify):
  - (0, 0) neither recently used nor modified - best page to replace
  - (0, 1) not recently used but modified - not quite as good, must write out before replacement
  - (1, 0) recently used but clean - probably will be used again soon
  - (1, 1) recently used and modified - probably will be used again soon and need to write out before replacement
- When page replacement called for, use the clock scheme but use the four classes replace page in lowest non-empty class
  - Might need to search circular queue several times

# Counting-based Page Replacement

---

- Keep the number of references made to each page
- **Least Frequently Used (LFU)** replaces page with the smallest counter
  - A page is heavily used during process initialization and then never used
- **Most Frequently Used (MFU)** replaces page with the largest counter
  - based on the argument that page with the smallest count was probably just brought in and has yet to be used
- LFU and MFU are not common

# Page-Buffering Algorithms

---

- **Keep a pool of free frames, always**
  - **frame available when needed**, no need to find at fault time
  - Read page into free frames without waiting for victims to write out
    - Restart as soon as possible
  - When convenient, evict victim
- Possibly, keep list of **modified pages**
  - When disk **idles**, write pages there and set to non-dirty: this page can be replaced **without** writing pages to backing store
- Possibly, keep free frame contents intact and note what is in them - a kind of cache
  - If **referenced again before reused**, no need to load contents again from disk
    - **cache hit**

# Applications and Page Replacement

---

- All of these algorithms have OS guessing about future page access
- Some applications have better knowledge - i.e. databases
- Memory intensive applications can cause **double buffering - a waste of memory**
  - OS keeps copy of page in memory as I/O buffer
  - Application keeps page in memory for its own work
- Operating system can given direct access to the disk, getting out of the way of the applications
  - **Raw disk mode**
- Bypasses buffering, locking, etc

# Allocation of Frames

---

- Each process needs **minimum number** of frames -according to instructions semantics
- Example: IBM 370 - **6 pages to handle SS MOVE** instruction:
  - instruction is 6 bytes, might span 2 pages
  - 2 pages to handle from
  - 2 pages to handle to
- **Maximum** of course is total frames in the system
- Two major allocation schemes for **process** memory allocation
  - Equal allocation
  - Proportional allocation

# Frame Allocation

---

- Equal allocation
  - For example, if there are 100 frames (after allocating frames for the OS) and 5 processes, give each process 20 frames
  - Keep some as free frame buffer pool
- Proportional allocation
  - Allocate according to the size of process
  - Dynamic as degree of multiprogramming, process sizes change

$s_i$  = size of process  $p_i$

$S = \sum s_i$

$m$  = total number of frames

$a_i$  = allocation for  $p_i = \frac{s_i}{S} \times m$

$m = 62$

$s_1 = 10$

$s_2 = 127$

$a_1 = \frac{10}{137} \times 62 \approx 4$

$a_2 = \frac{127}{137} \times 62 \approx 57$



# Global vs. Local Allocation

---

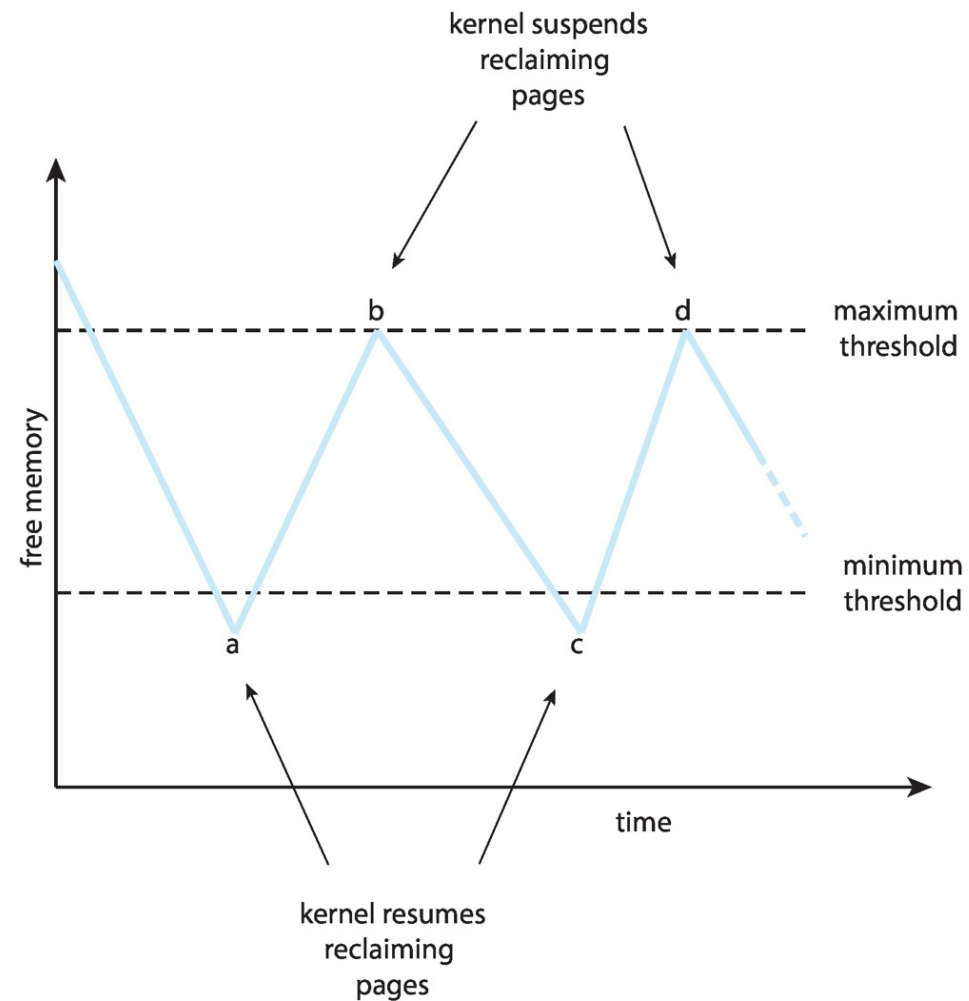
- **Global replacement** - process selects a replacement frame from the set of all frames; one process can take a frame from another
  - But then process execution time can **vary** greatly - depends on others
  - But greater throughput so more common
- **Local replacement** - each process selects from only its own set of allocated frames
  - More consistent per-process performance
  - But possibly underutilized memory

# Reclaiming Pages

---

- A strategy to implement global page-replacement policy
- All memory requests are satisfied from the free-frame list
- Rather than waiting for the list to drop to zero before we begin selecting pages for replacement, page replacement is triggered when the list falls below a **certain threshold**.
- This strategy attempts to ensure **there is always sufficient** free memory to satisfy new requests.

# Reclaiming Pages Example



## What happens if memory is below the minimum threshold

- Reclaim pages aggressively
  - Kill some processes
    - According to OOM score
    - how likely it is to be terminated in case of low available memory

# Major and minor page faults

---

- Major: page is referenced but not in memory
  - Can only be satisfied by disk
  - `do_anonymous_page` is not major
- Minor: mapping does not exist, but the page is in memory
  - Shared library
  - Reclaimed and not freed yet


```
wenbo@parallels:~$ ps eo minflt,majflt,cmd
MINFL  MAJFL  CMD
    704      1 /usr/lib/gdm3/gdm-x-session --run-script env
59284    17 /usr/lib/xorg/Xorg vt2 -displayfd 3 -auth /ru
 8439    12 /usr/lib/gnome-session/gnome-session-binary -
73769    25 /usr/bin/gnome-shell USER=wenbo LC_TIME=zh_CN
 1346      0 ibus-daemon --xim --panel disable USER=wenbo
   376      1 /usr/lib/ibus/ibus-dconf USER=wenbo LC_TIME=z
 1659      0 /usr/lib/ibus/ibus-x11 --kill-daemon USER=wen
```

Thanks to shared libraries!

# Non-Uniform Memory Access

---

- So far all memory accessed equally
- Many systems are NUMA - speed of access to memory varies
  - Consider system boards containing CPUs and memory, interconnected over a system bus
- NUMA multiprocessing architecture



# Non-Uniform Memory Access (Cont.)

---

- Optimal performance comes from allocating memory “close to” the CPU on which the thread is scheduled
  - And modifying the **scheduler** to schedule the thread on the same system board when possible
  - Linux
    - Kernel maintains scheduling domains: does not allow threads to migrate across domains
    - A separate free-frame list for each NUMA node - allocating memory from the node it is running

# Thrashing


---

- If a process doesn't have "enough" pages(frames), page-fault rate may be high
  - page fault to get page, replace some existing frame
  - but quickly need replaced frame back
  - this leads to:
    - low CPU utilization ➡
    - kernel thinks it needs to increase the degree of multiprogramming to maximize CPU utilization ➡
    - another process added to the system
- **Thrashing:** a process is busy swapping pages in and out



# Thrashing

---



# Demand Paging and Thrashing

---

- Why does demand paging work?
  - process memory access has **high locality**
  - process migrates from one locality to another, localities may overlap
- Why does thrashing occur?
  - total memory size < total size of locality


# Resolve thrashing - Option I

---

- Limit thrashing effects by using local page replacement
  - One process starts thrashing does not affect others -> it cannot cause other processes thrashing
    - Select swap out page from the same process

# Resolve thrashing - Option II

- Provide a process with as many frames as it needs. How?
  - Time (a), the locality is the set of pages {18-24, 29-33}.
  - Time (b), the locality changes to {18-20, 24-29, 31-33}
  - 18, 19, and 20 are part of both localities




# Working-Set Model

---

- **Working-set window( $\Delta$ ):** a fixed number of page references
  - if  $\Delta$  too small  $\Rightarrow$  will not include entire locality
  - if  $\Delta$  too large  $\Rightarrow$  will include several localities
  - if  $\Delta = \infty \Rightarrow$  will include entire program
- **Working set of process  $p_i$  ( $WSS_i$ ):** total number of pages referenced in the most recent  $\Delta$  (varies in time)
- **Total working sets:**  $D = \sum WSS_i$ 
  - approximation of total locality
  - if  $D > m \Rightarrow$  possibility of thrashing
  - to avoid thrashing: if  $D > m$ , suspend or swap out some processes

# Working-Set Model

- Working-set window  $\Delta = 10$




# Challenge: Keeping Track of the Working Set

---

- Approximate with interval timer + a reference bit (in register)
  - The reference bit is set to 1 by hardware when accessing a page
- Example:  $\Delta = 10,000$  time units
  - Timer interrupts after every 5000 time units
  - Keep 2 bits in memory for each page
  - Whenever a timer interrupts
    - OS copies the reference bit (to memory bits [0] or [1]) and sets the values of all reference bits to 0
  - If one of the bits in memory = 1  $\Rightarrow$  page in working set
- Why is this not completely accurate?
  - 2 memory bits can only record two interrupts
  - Can not tell **when (in 5000 time units)** the access occurs
- Improvement = 10 bits and interrupt every **1000** time units

# Working Sets and Page Fault Rates

- Assumes there is no thrashing
- Direct relationship between working set of a process and its page-fault rate
- Working set changes over time
- Peaks and valleys over time




Page fault increases due to new locality



# Page-Fault Frequency

---

- More direct approach than WSS
- Establish “acceptable” page-fault frequency (PFF) rate
  - If actual rate too low, process loses frame
  - If actual rate too high, process gains frame
- Need to swap out a process if no free frames are available



# Kernel Memory Allocation

---

- Kernel memory allocation is treated differently from user memory, it is often allocated from a free-memory pool
  - kernel requests memory for structures of varying sizes -> minimize waste due to fragmentation
  - Some kernel memory needs to be **physically contiguous**
    - e.g., for device I/O

# Other Considerations

---

- Prepaging
- Page size
- TLB reach
- Inverted page table
- Program structure
- I/O interlock and page locking

# Prepaging

---

- To reduce the large number of page faults that occurs at process startup
  - Prepage all or some of the pages a process will need, before they are referenced
  - But if prepaged pages are unused, I/O and memory was wasted
- Assume  $s$  pages are prepaged and a fraction  $a$  of these pages is used
  - Is cost of  $s * a$  save pages faults  $>$  or  $<$  than the cost of prepaging
  - $s * (1 - a)$  unnecessary pages?
  - $a$  near zero  $\rightarrow$  prepaging loses

# Page Size

---

- Sometimes OS designers have a choice
  - Especially if running on custom-built CPU
- Page size selection must take into consideration:
  - Fragmentation -> small page size
  - Page table size -> large page size
  - Resolution -> small page size
  - I/O overhead -> large page size
  - Number of page faults -> large page size
  - Locality -> small page size
  - TLB size and effectiveness -> large page size
- Always power of 2, usually in the range  $2^{12}$  (4,096 bytes) to  $2^{22}$  (4,194,304 bytes)
- On average, **growing over time**

# TLB Reach

---

- **TLB reach:** the amount of memory accessible from the TLB
  - $\text{TLB reach} = (\text{TLB size}) \times (\text{page size})$
- Ideally, the working set of each process is stored in the TLB
  - otherwise there is a high degree of TLB misses
- **Increase the page size to reduce TLB pressure**
  - it may increase fragmentation as not all applications require large page sizes
  - multiple page sizes allow applications that require larger page sizes to use them without an increase in fragmentation

# Other Issues: Program Structure

---

- Program structure can affect page faults
  - `int[128,128] data;`
  - Assume page size is 512B; each row is stored in one page;
  - Assume system has less than 127 physical frames
  - Program 1:

```
for (j = 0; j < 128; j++)  
    for (i = 0; i < 128; i++)  
        data[i,j] = 0;
```

128 × 128 = 16,384 page faults

- Program 2:


```
for (i = 0; i < 128; i++)  
    for (j = 0; j < 128; j++)  
        data[i,j] = 0;
```

128 page faults

# I/O interlock

---

- I/O Interlock - Pages must sometimes be locked into memory
- Consider I/O - Pages that are used for copying a file from a device must be **locked from being selected for eviction by a page replacement algorithm**





# Windows XP

---

- Uses demand paging with clustering
  - clustering brings in **pages surrounding the faulting page**: locality
- Processes are assigned **working set** minimum and set maximum
  - *wsmin*: minimum number of pages the process is guaranteed to have
  - *wsmax*: a process may be assigned as many pages up to its *wsmax*
- When the amount of free memory in the system falls below a threshold:
  - automatic working set trimming to restore the amount of free memory
  - it removes pages from processes that have more pages than the *wsmin*

# Memory management in Linux

---

- mm\_struct
  - [https://elixir.bootlin.com/linux/v5.10/source/include/linux/mm\\_types.h#L388](https://elixir.bootlin.com/linux/v5.10/source/include/linux/mm_types.h#L388)
  - Where is page table?
  - Code walk through

# Page fault handling in Linux

---

- `do_page_fault`
  - <https://elixir.bootlin.com/linux/v5.10/source/arch/riscv/mm/fault.c#L189>
  - Code walk through

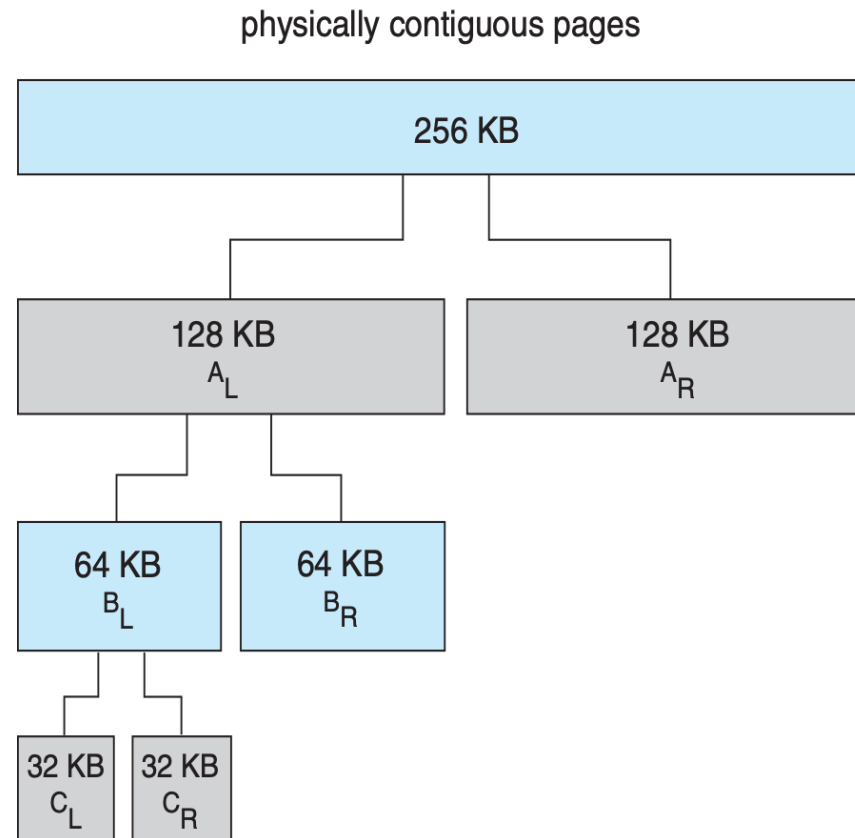
# Linux Buddy System

---

- Memory allocated using power-of-2 allocator
  - memory is allocated in units of the size of **power of 2**
    - round up a request to the closest allocation unit
    - split the unit into two "**buddies**" until a proper sized chunk is available
  - e.g., assume only 256KB chunk is available, kernel requests 21KB
    - split it into  $A_l$  and  $A_r$  of 128KB each
    - further split an 128KB chunk into  $B_l$  and  $B_r$  of 64KB
    - again, split a 64KB chunk into  $C_l$  and  $C_r$  of 32KB each
    - give one chunk for the request
- advantage: it can quickly merge unused chunks into larger chunk
- disadvantage: **internal fragmentation**
  - 21k request -> 32k segment

# Buddy System Allocator

---




# Slab Allocator

---

- Slab allocator is a **cache of objects**
  - a **cache** in a slab allocator consists of one or more slabs
  - a slab contains **one or more pages**, divided into **equal-sized objects**
  - kernel uses one cache for each unique kernel data structure
    - when cache created, allocate a slab, divided the slab into free objects
    - objects for the data structure is allocated from free objects in the slab
    - if a slab is full of used objects, next object comes from an empty/new slab
- Benefits: **no fragmentation** and fast memory allocation
  - some of the object fields may be reusable; no need to initialize again

# Slab Allocation

---



A 12k slab (3 pages) can store 4 3k objects.

# Slab Allocator in Linux

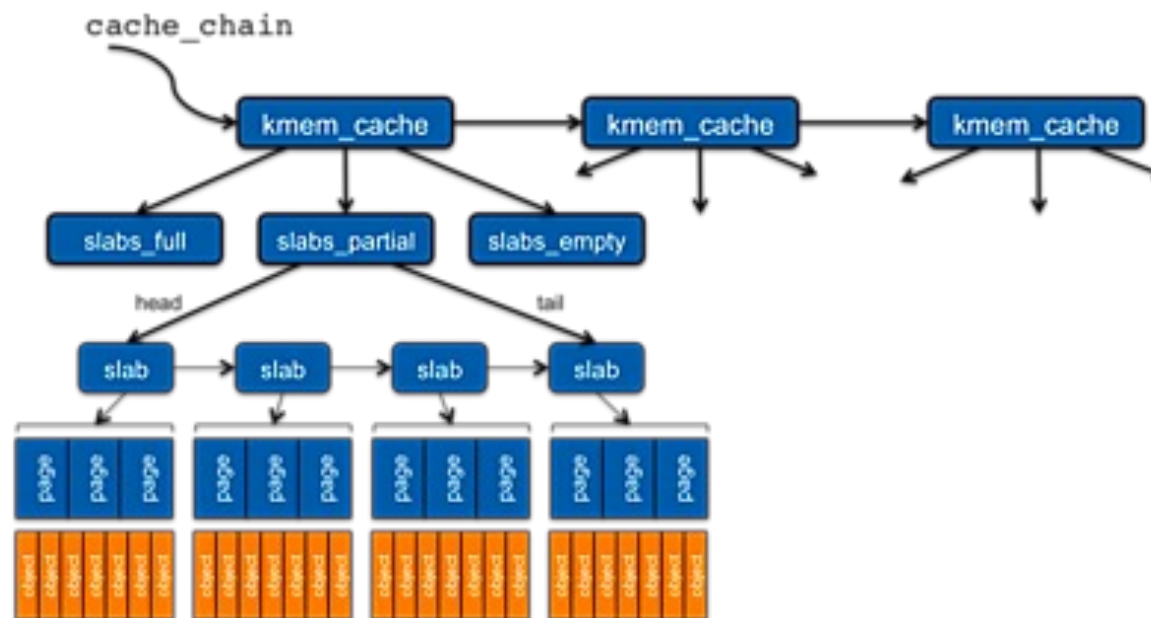
---

- For example process descriptor is of type *struct task\_struct*
  - Approx. 1.7KB of memory (some old linux version)
- New task -> allocate new struct from cache
  - Will use existing free *struct task\_struct*
- A Slab can be in three possible states
  - **Full** - all used
  - **Empty** - all free
  - **Partial** - mix of free and used
- Upon request, slab allocator
  - Uses free struct in **partial** slab
  - If none, takes one from **empty** slab
  - If no empty slab, create new empty



# Slab in Linux

---



# Slab Allocator in Linux (Cont.)

---

- Slab started in Solaris, now wide-spread for both kernel mode and user memory in various OSes
- Linux 2.2 had SLAB, now has both SLOB and SLUB allocators
  - SLOB for systems with limited memory
    - Simple List of Blocks - maintains 3 list objects for small, medium, large objects
- SLUB is performance-optimized SLAB removes per-CPU queues, metadata stored in page structure

# Takeaway

---

- Page fault
  - Valid virtual address, invalid physical address
- Page replacement
  - FIFO, Optimal, LRU, 2nd chance
- Thrashing and working set
- Buddy system
- slab

# Memory management summary

---

- Partition
- Segmentation
- Paging
  - Page table
  - Demand paging
  - Page replacement
  - Working set
- Linux memory mapping
  - 32-bit
  - 64-bit