


Virtual Memory and Linux

Based on Alan Ott's slides on
Embedded Linux Conference
April 4-6, 2016



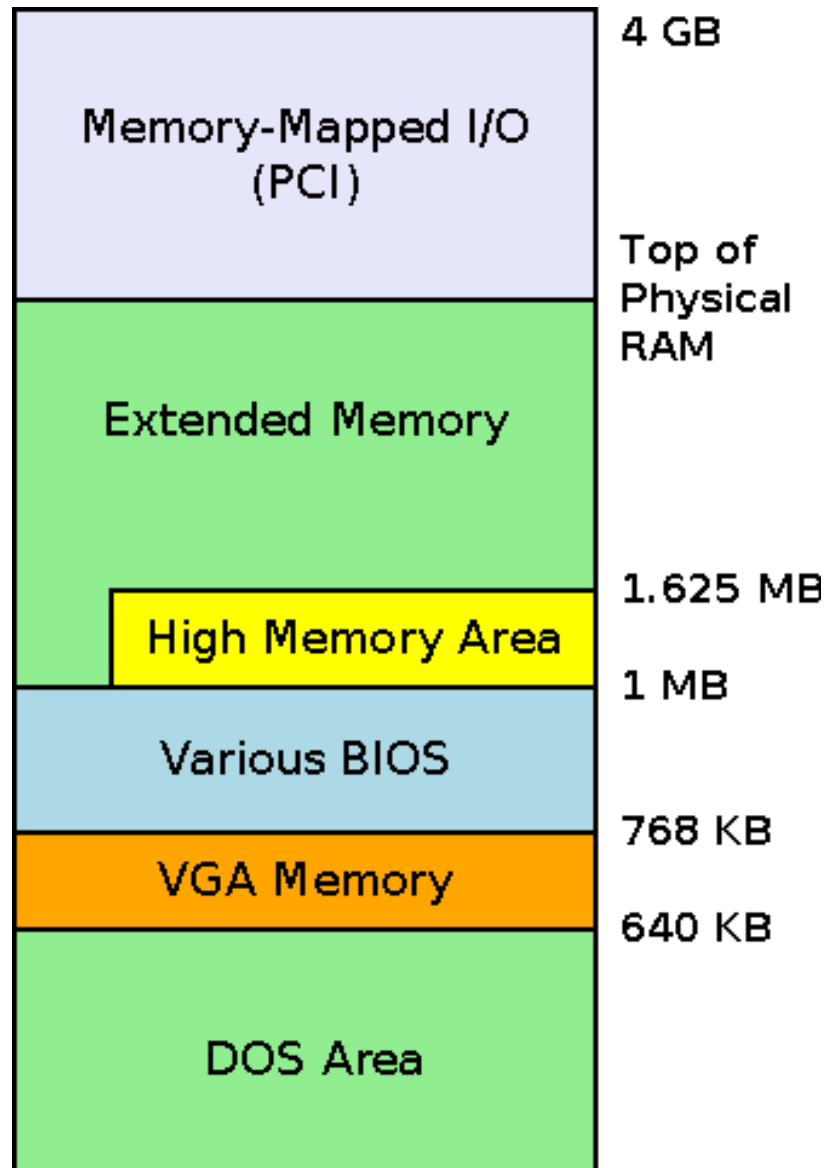
Flat Memory

- Older and modern, but simple systems have a single address space
 - Memory and peripherals share
 - Memory will be mapped to one part
 - Peripherals will be mapped to another
 - All processes and OS share the same memory space
 - No memory protection!
 - User space can stomp kernel mem!

Flat Memory

- CPUs with flat memory
 - 8086-80206
 - ARM Cortex-M
 - IoT chips - most 8- and 16-bit systems

Old x86 Memory Map




- Lots of Legacy
- RAM is split (DOS Area and Extended)
- Hardware mapped between RAM areas
- High and Extended area accessed differently

Limitations

- Portable C programs expect flat memory
 - Accessing memory by segments limits portability
- Management is tricky
 - Need to keep processes separated
- No protection
 - Rogue programs can corrupt the entire system

Virtual Memory



What is Virtual Memory?

- Virtual Memory (logical memory) is an address mapping
- Maps virtual address space to physical address space
 - Maps virtual addresses to physical RAM
 - Maps virtual addresses to hardware devices
 - PCI devices
 - GPU RAM
 - On-SoC IP blocks

What is Virtual Memory?

- Advantages
 - Each processes can have a different memory mapping
 - One process's RAM is inaccessible (and invisible) to other processes
 - Built-in memory protection
 - Kernel RAM is invisible to user-space processes
 - Memory can be moved
 - Memory can be swapped to disk

What is Virtual Memory?

- Advantages (cont)
 - Hardware device memory can be mapped into a process's address space
 - Requires kernel perform the mapping
 - Physical RAM can be mapped into multiple processes at once
 - Shared memory, shared libraries
 - Memory regions can have access permissions
 - Read, write, execute

Virtual Memory Details

- Two address spaces
 - Physical addresses
 - Addresses as used by the hardware (except CPU)
 - DMA, peripherals
 - Virtual addresses
 - Addresses as used by software (CPU generated)
 - Load/Store instructions (RISC)
 - Any instruction accessing RAM (CISC)

Virtual Memory Details

- Mapping is performed in hardware
 - No performance penalty for accessing already- mapped RAM regions
 - Permissions are handled without penalty
 - The same CPU instructions are used for accessing RAM and mapped hardware
 - Software, during its normal operation, will only use virtual addresses
 - Includes kernel and user space

Memory-Management Unit

- The memory-management unit (MMU) is the hardware responsible for implementing virtual memory
 - Sits between the CPU core and memory
 - Most often part of the physical CPU itself.
 - On ARM, it's part of the licensed core
 - Separate from the RAM controller
 - DDR controller is a separate IP block

Memory-Management Unit

- MMU (cont)
 - Transparently handles all memory accesses from Load/Store instructions
 - Maps accesses using virtual addresses to system RAM
 - Maps accesses using virtual addresses to memory-mapped peripheral hardware
 - Handles permissions
 - Generates an exception (page fault) on an invalid access
 - Unmapped address or insufficient permissions

Translation Lookaside Buffer

- The TLB stores the mappings from virtual to physical address space in hardware
 - Also holds permission bits
- TLB is a part of the MMU
- TLB is consulted by the MMU when the CPU accesses a virtual address

Page Faults

- A page fault is a CPU exception, generated when software attempts to use an invalid virtual address. There are three cases:
 - The virtual address is not mapped for the process requesting it.
 - The processes has insufficient permissions for the address requested.
 - The virtual address is valid, but swapped out
 - This is a software condition

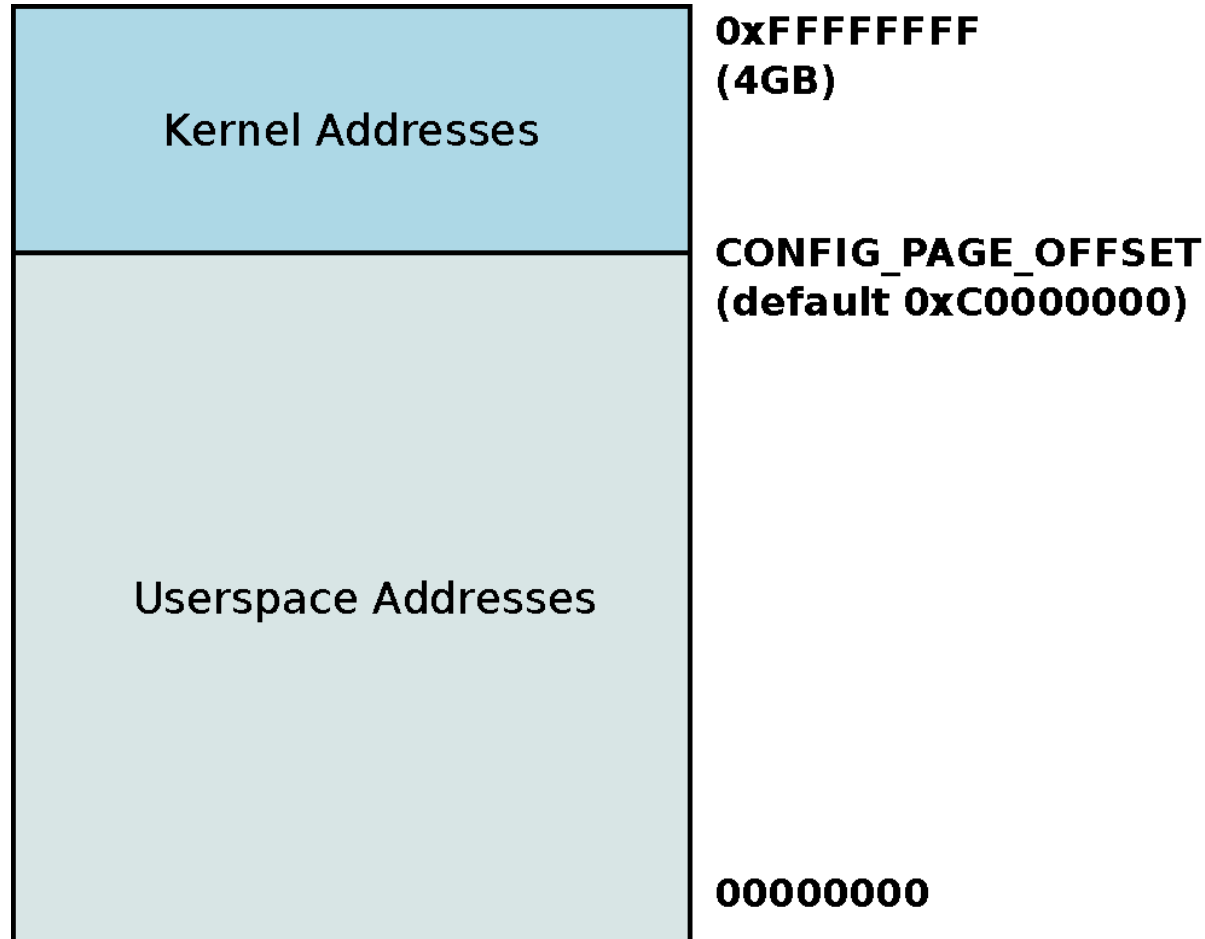
Lazy Allocation

- The kernel uses lazy allocation of physical memory.
 - When memory is requested by userspace, physical memory is not allocated until it's touched.
 - This is an optimization, knowing that many userspace programs allocate more RAM than they ever touch.
 - Buffers, etc.

Virtual Addresses

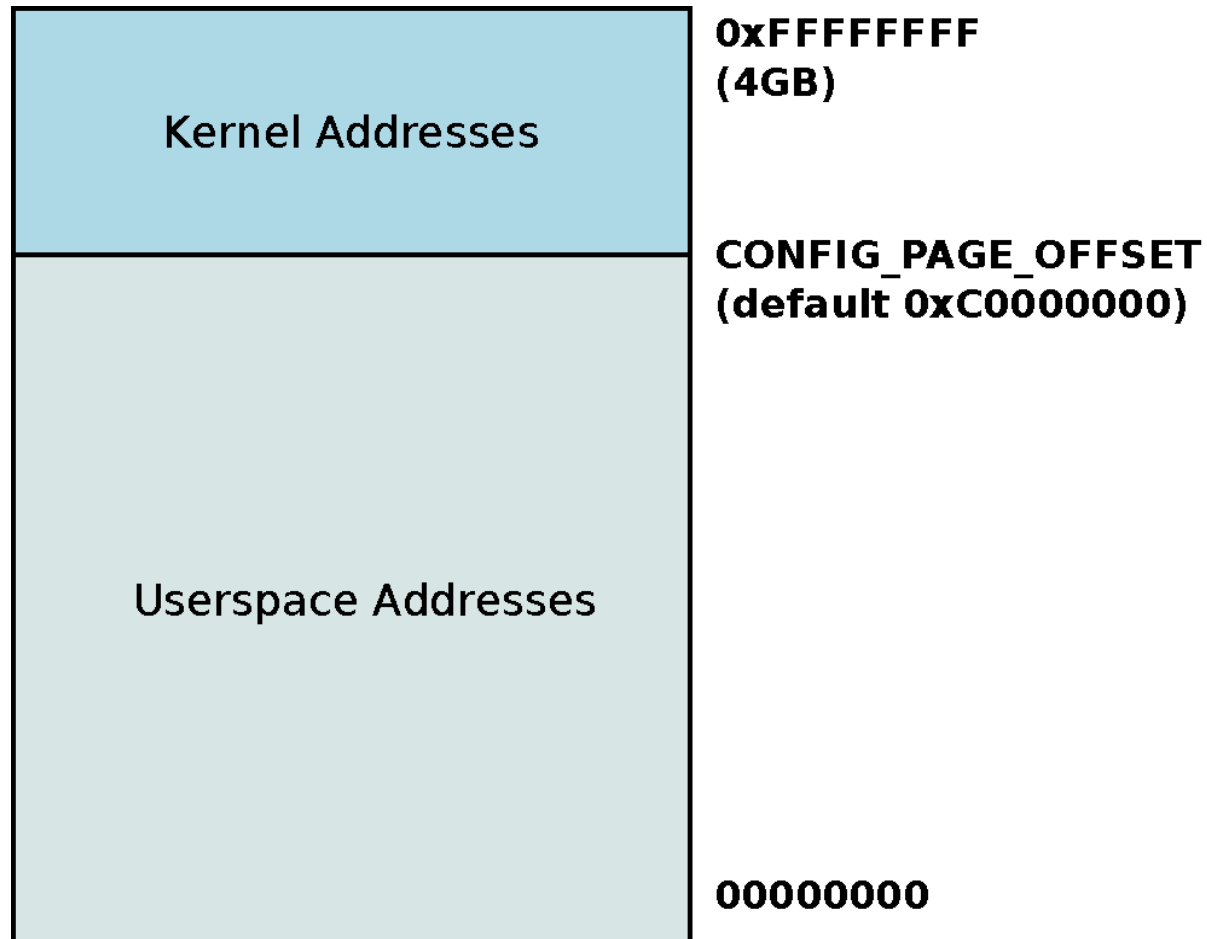
- In Linux, the kernel uses virtual addresses, as userspace processes do.
- Virtual address space is split.
 - The upper part is used for the kernel
 - The lower part is used for userspace
- On 32-bit, the split is at
0xC0000000

Virtual Addresses - Linux



- By default, the kernel uses the top 1GB of virtual address space.
- Each userspace processes get the lower 3GB of virtual address space.

User Virtual Addresses



- Each process will have its own mapping for user virtual addresses
- The mapping is changed during context switch

Show me the code

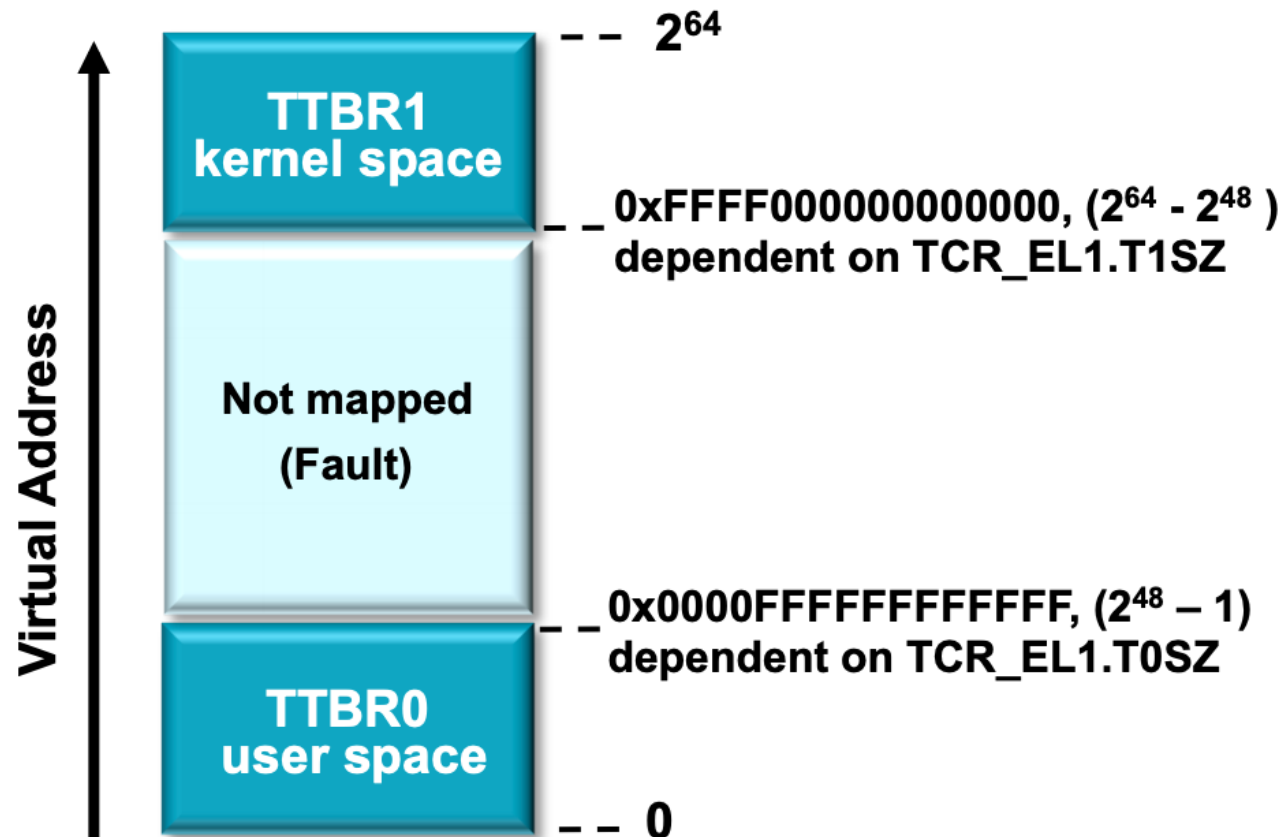
- Talk is cheap, show me the code
 - switch_mm()

Virtual Addresses - Linux

- Kernel address space is the area above `CONFIG_PAGE_OFFSET`.
 - For 32-bit, this is configurable at kernel build time.
 - The kernel can be given a different amount of address space as desired
 - See `CONFIG_VMSPLIT_1G`, `CONFIG_VMSPLIT_2G`, etc.
 - For 64-bit, the split varies by architecture, but it's high enough
 - `0xffff000000000000`– ARM
 - `0xffff000000000000`– x86_64

Virtual Addresses 64-bit

- ARM 64
 - 48-bit page table
 - 4-level: 9+9+9+9+12



Virtual Addresses 64-bit

- ARM 64
 - 39-bit page table
 - 3-level: 9+9+9+12
- 39-bit virtual address for both user and kernel
 - 00000000000000000–0000007fffffffffff (512GB): user
 - [architectural gap]
 - ffffffff8000000000–ffffffffffbffffffffff (~240MB): vmalloc
 - ffffffffbbffff0000–ffffffffffbcfffffffffff (64KB): [guard]
 - ffffffffbc00000000–ffffffffffbdfffffffffff (8GB): vmemmap
 - ffffffffbe00000000–ffffffffffbfbffffffffff (~8GB): [guard]
 - ffffffffbfcc000000–ffffffffffbffffffffffff (64MB): modules
 - ffffffffcc00000000–fffffffffffffffffffff (256GB): mapped RAM
- 4KB page configuration
 - 3 levels of page tables (pgtable-nopud.h)
 - Linear mapping using 4KB, 2MB or 1GB blocks
 - AArch32 (compat) supported

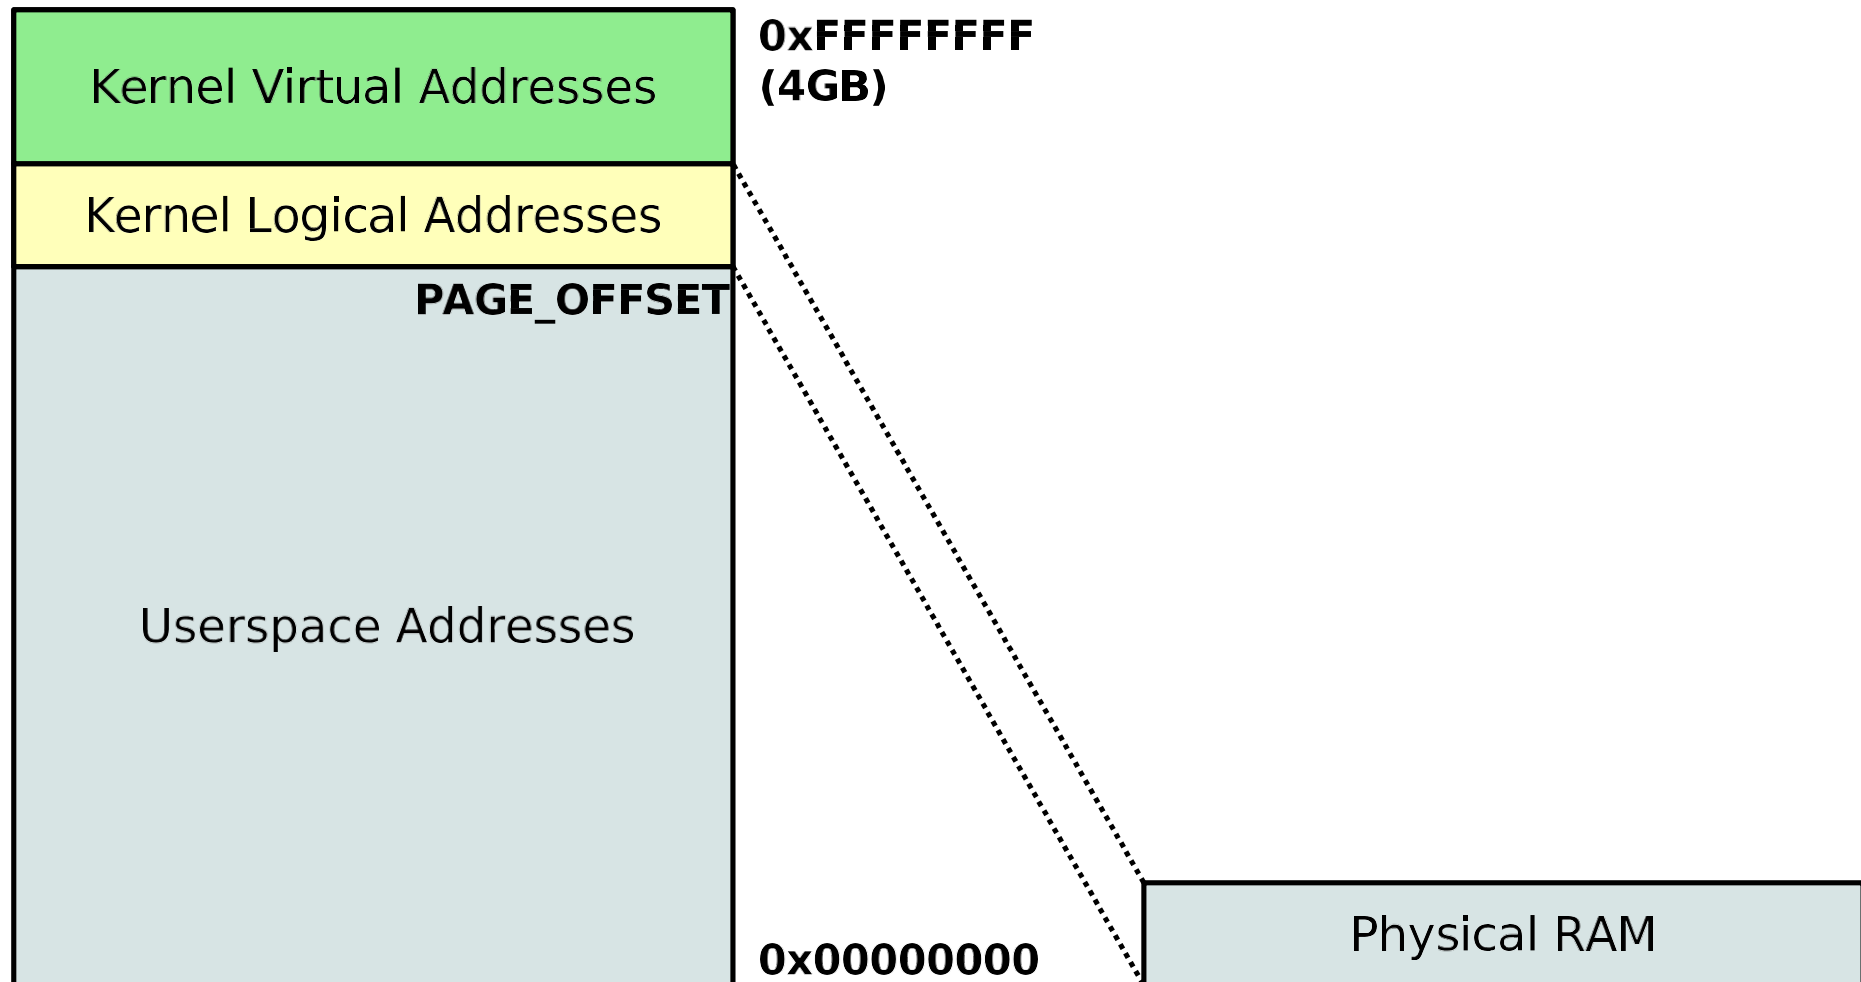
Virtual address and physical RAM mapping

- How kernel manage physical RAM?
 - Such as allocate a physical frame to user process
- Must map physical RAM to kernel AS
 - Must have an address in kernel AS
 - For 32-bit, kernel address space (**AS**) is 1GB
 - No problem to handle small RAM ($\leq 896\text{MB}$)

Kernel Virtual Addresses

Virtual Address Space

Physical Address Space




Kernel Logical Addresses

- Kernel Logical Addresses
 - Addresses above PAGE_OFFSET
 - First 896MB of the kernel AS
 - Virtual addresses are a linear offset from their physical addresses.
 - Eg: Virt: 0xc0000000 → Phys: 0x00000000
 - This makes converting between physical and virtual addresses easy

How to do VA to PA translation?

- Given a virtual address v , what is the physical address?
- Given a physical address p , what is the virtual address?



```
#define __pa_to_va_nodebug(x) ((void *)((unsigned long) (x) + va_pa_offset))
#define __va_to_pa_nodebug(x) ((unsigned long)(x) - va_pa_offset)
```

Kernel Logical Addresses

- Kernel Logical addresses can be converted to and from physical addresses using the macros:

 __pa(x)

 __va(x)

- For low-memory systems (below ~1G of RAM) Kernel Logical address space starts at PAGE_OFFSET and goes through the end of *physical* memory.

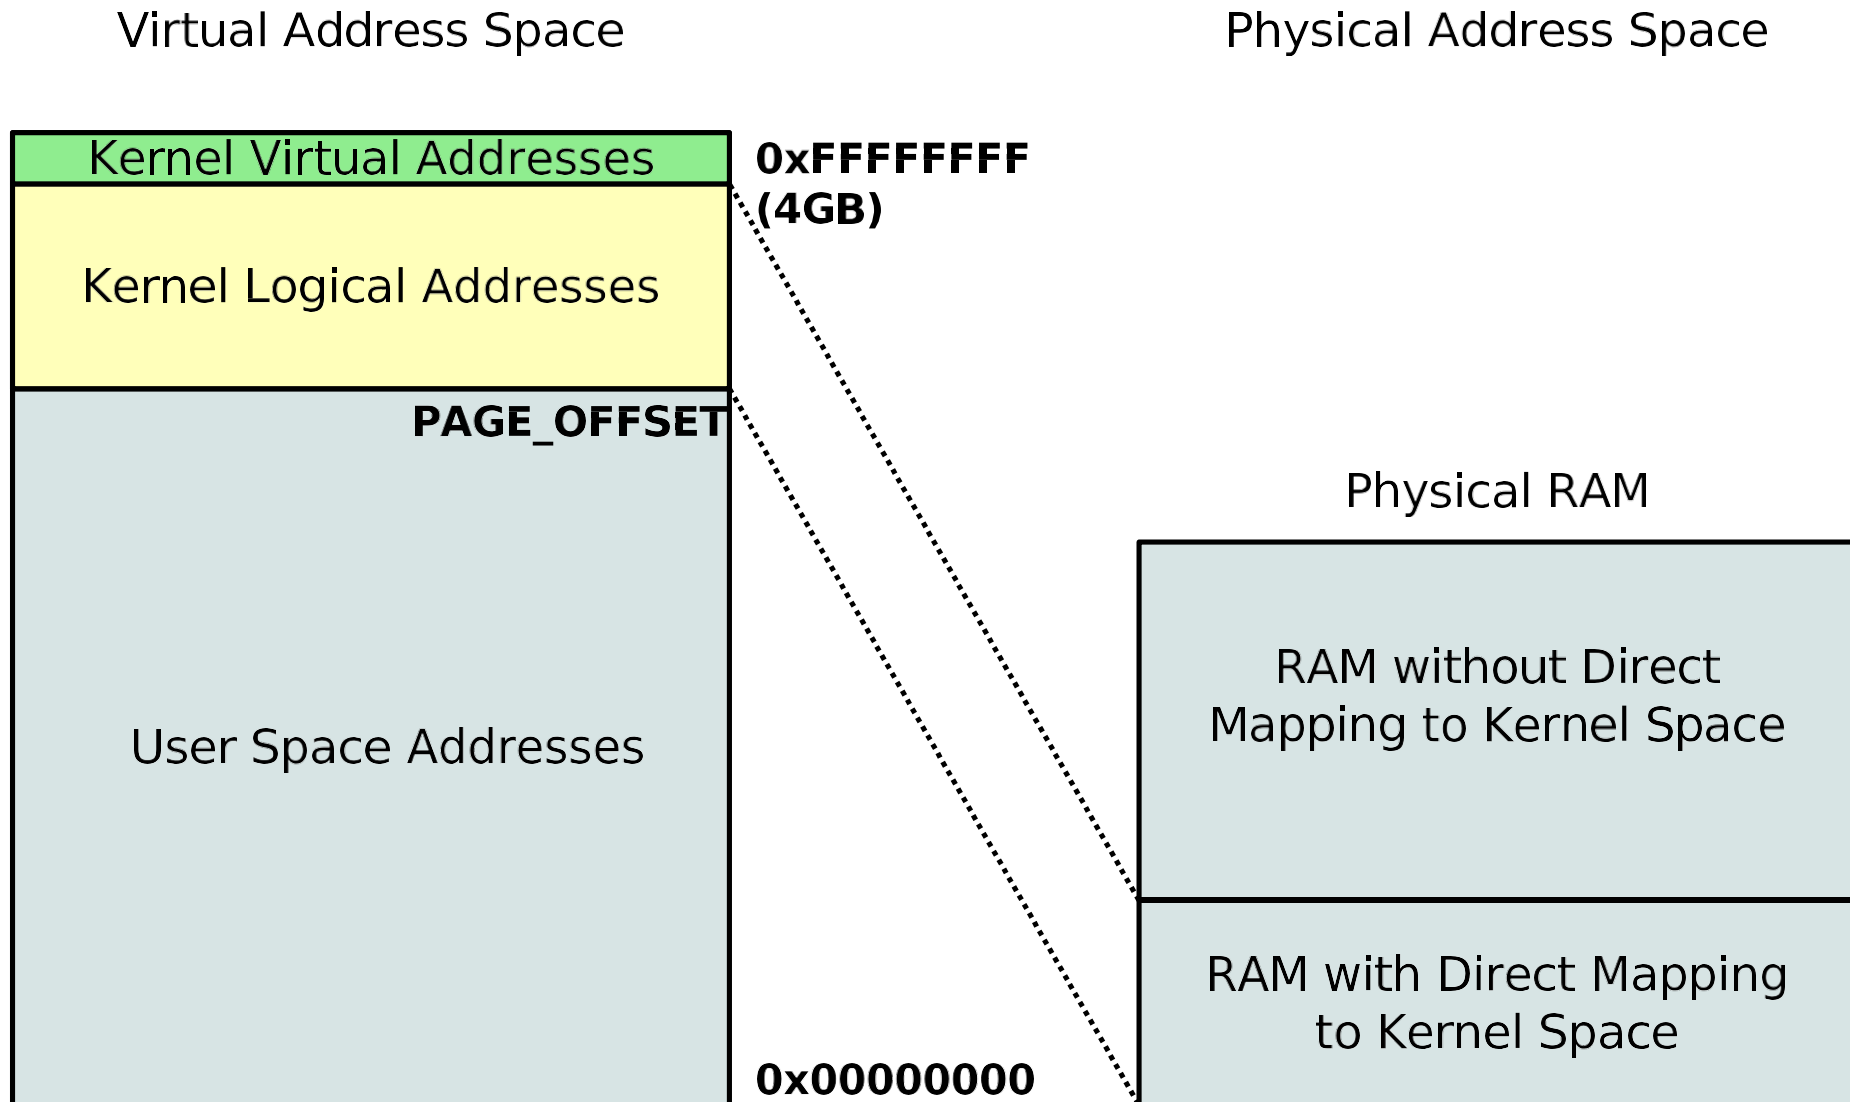
Large physical RAM

- How kernel manage physical RAM?
 - Such as allocate a physical frame to user process
- Must map physical RAM to kernel AS
 - For 32-bit, kernel address space (**AS**) is 1GB
 - No problem to handle small RAM ($\leq 896\text{MB}$)
 - How about **large** physical RAM $\geq 1\text{GB}$
 - Even **larger** than the kernel AS

Large physical RAM

- For large memory systems (more than ~1GB RAM), **not** all of the physical RAM can be **linearly** mapped into the kernel's address space.
 - Kernel address space is the top 1GB of virtual address space, by default.
 - Further, 128 MB is reserved at the top of the kernel's memory space for non-contiguous allocations
 - See `vmalloc()` described later

Kernel Virtual Addresses (Large Mem)



Large physical RAM

- As a result, only the bottom part of physical RAM has a kernel logical address
 - Note that on 64-bit systems, kernel AS is much larger than physical RAM, this case never happens
 - 39-bit kernel AS is 512GB
 - 48-bit kernel AS is 256TB
- Thus, in a large memory situation, only the bottom part of physical RAM is mapped linearly into kernel logical address space
 - Top 128MB guarantees that there is always enough kernel address space to accommodate all the RAM

Kernel Virtual Addresses

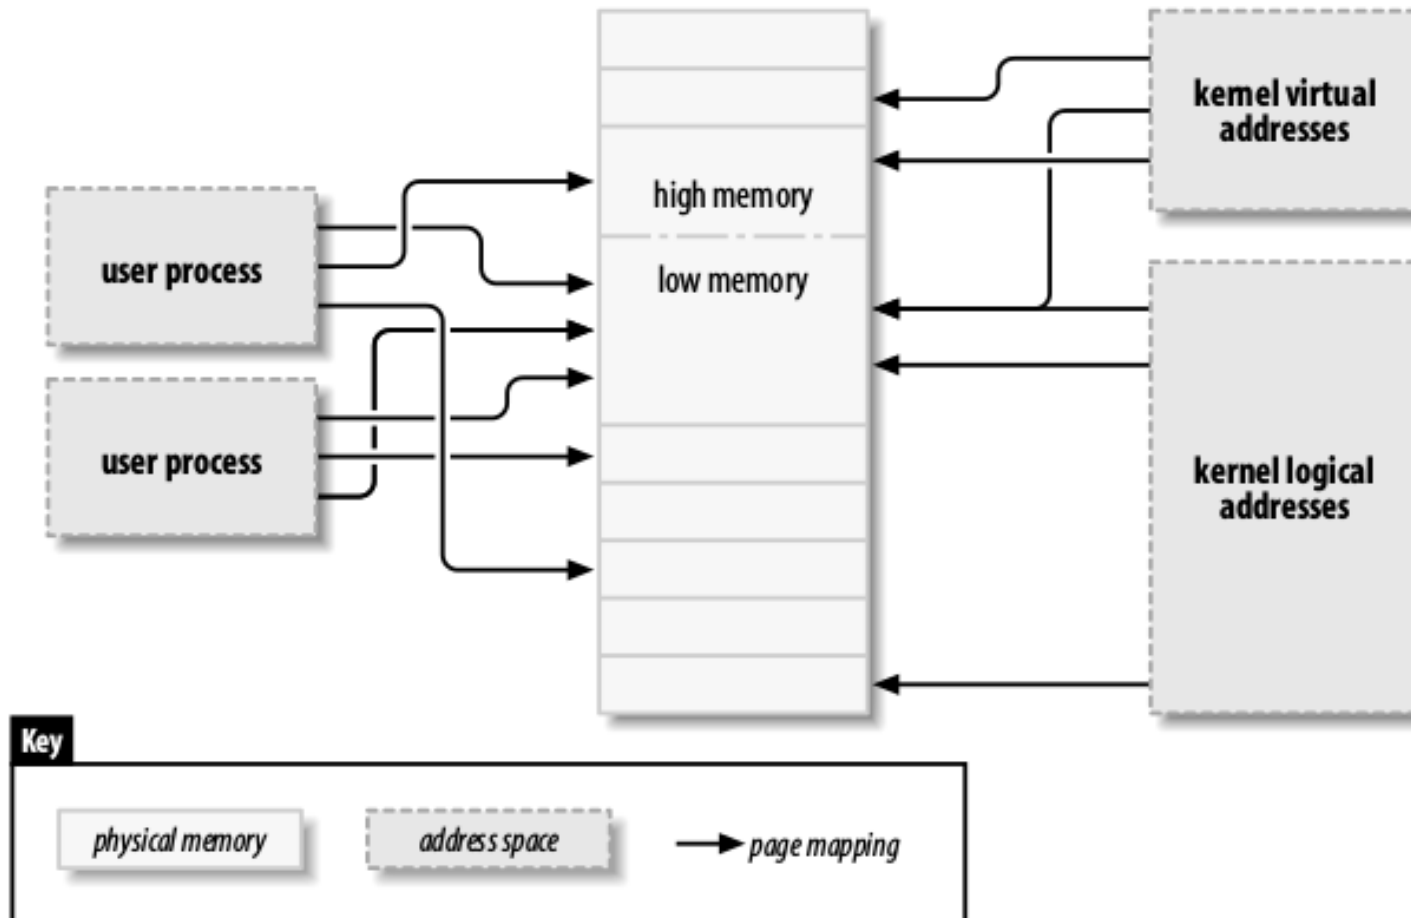
- Kernel Virtual Addresses are addresses in the region above the kernel logical address mapping.
- Kernel Virtual Addresses are used for non-contiguous memory mappings
 - Often for large buffers which could potentially be unable to get physically contiguous regions allocated.
- Also referred to as the `vmalloc()` area

Virtual Addresses - Summary

- There are three kinds of virtual addresses in Linux.
 - The terminology varies, even in the kernel source, but the definitions in *Linux Device Drivers*, 3rd Edition, chapter 15, are somewhat standard.
 - LDD 3 can be downloaded for free at:
<https://lwn.net/Kernel/LDD3/>

Virtual Addresses - Linux

- User virtual address
- Kernel virtual address
- Kernel logical address

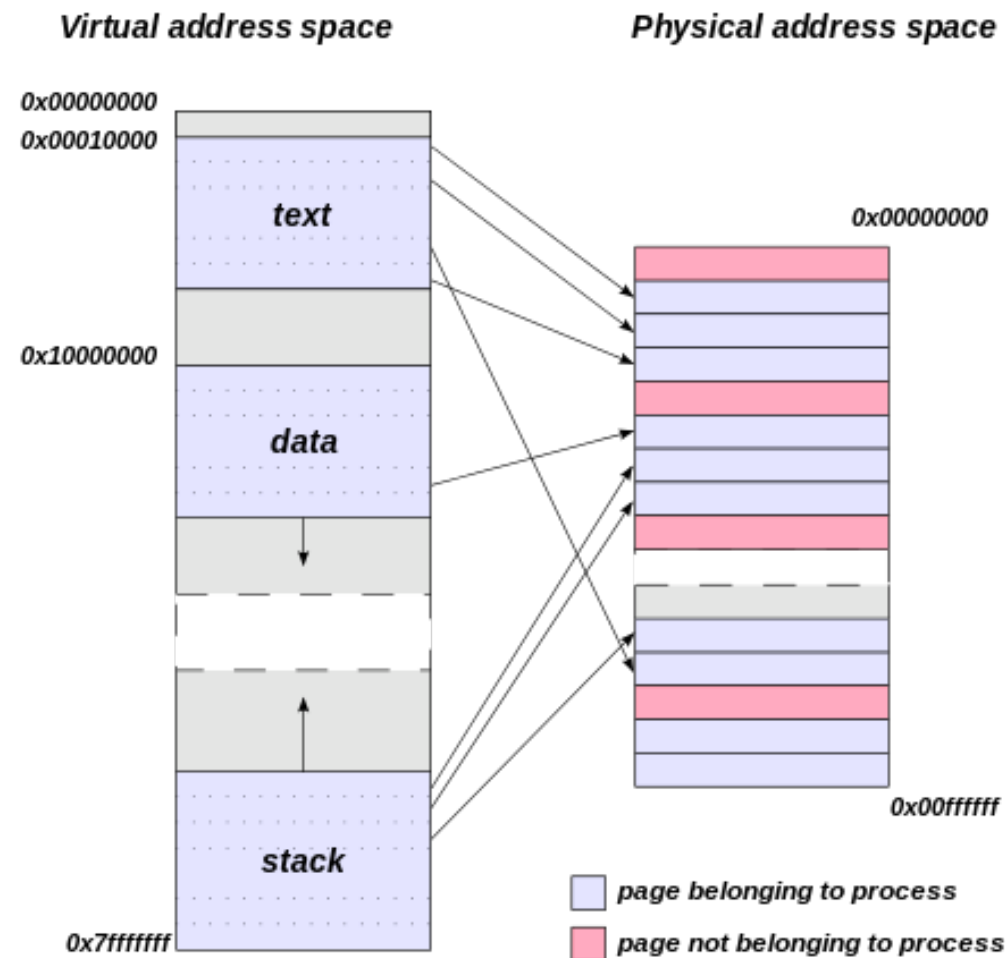


Virtual Addresses - Linux

- User virtual address, Kernel virtual address, Kernel logical address
 - All belong to virtual address, but has different mappings to physical memory
 - Who has contiguous mapping?
 - Linear mapping
 - Who has non-contiguous mapping?
 - Non-linear mapping
 - How to convert a VA to PA?
- All these mapping can be done by MMU
 - By walking page table
 - Review: what is page table?

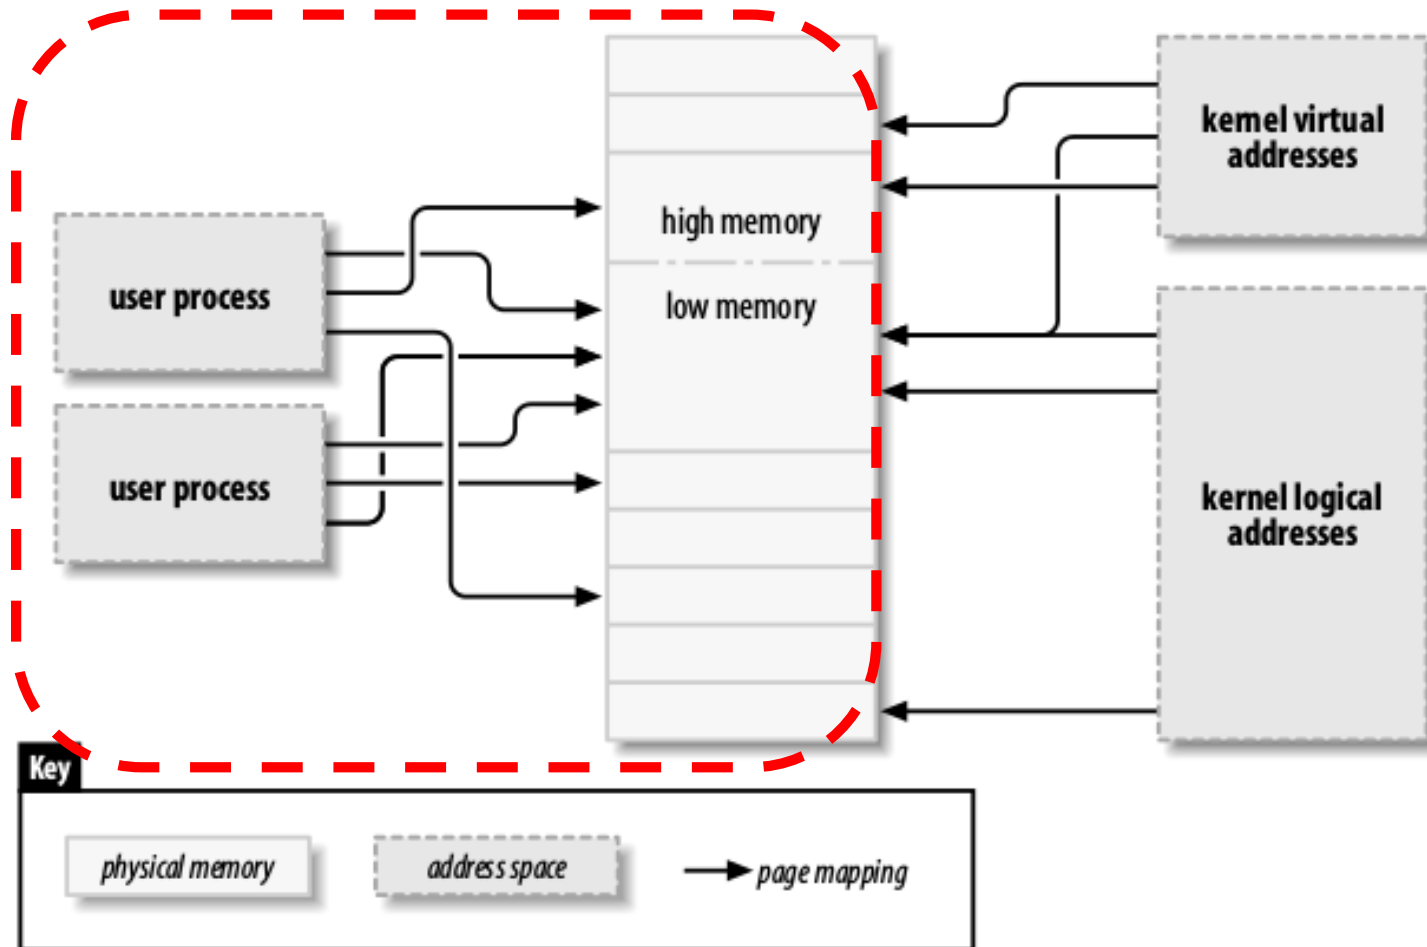
User space page table mapping

- Non-contiguous



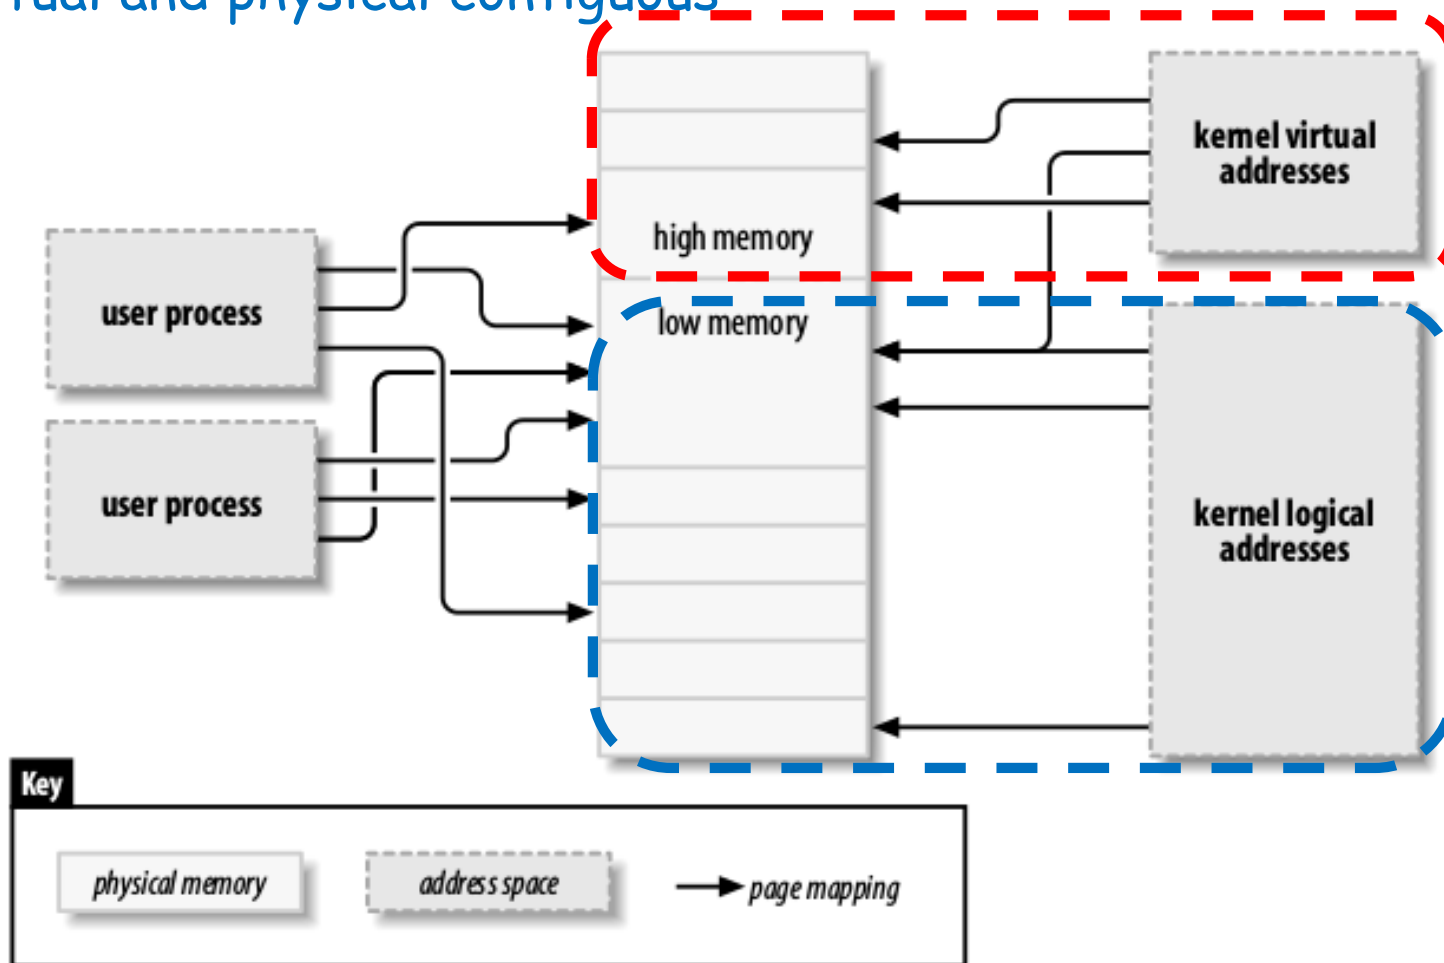
User space page table mapping

- Virtual contiguous, physical non-contiguous



Kernel space page table mapping

- Virtual contiguous, physical non-contiguous
- Virtual and physical contiguous




Basic Page Table Mappings

User Virtual Address Space

Page Table

Physical Address Space

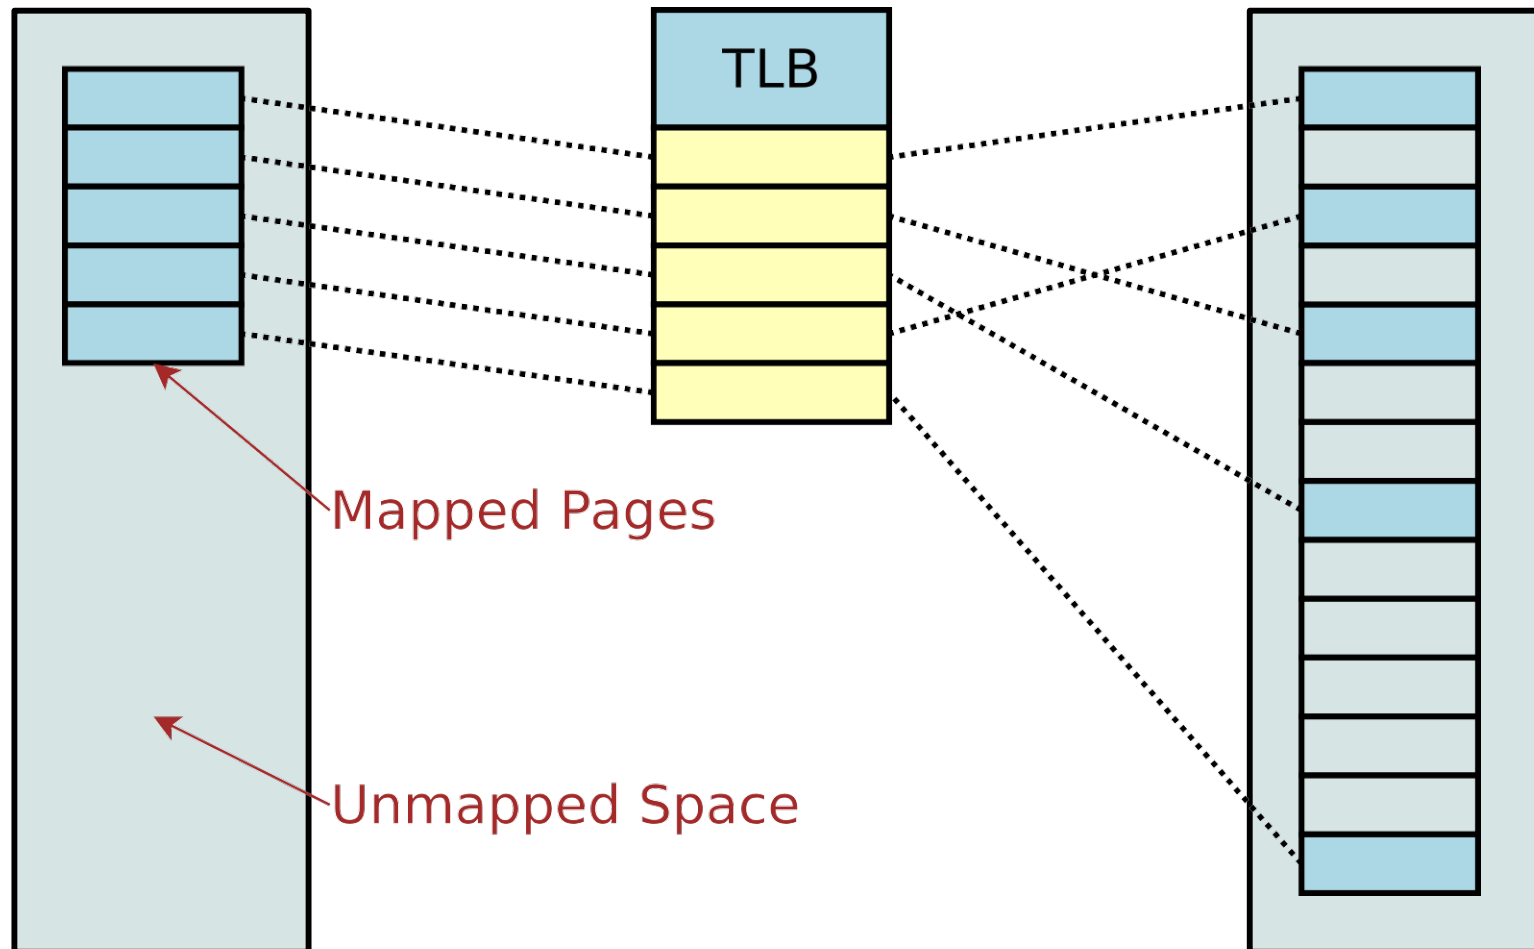


Basic Page Table Mappings

User Virtual Address Space

Page Table

Physical Address Space




Basic Page Table Mappings

User Virtual Address Space

Page Table

Physical Address Space

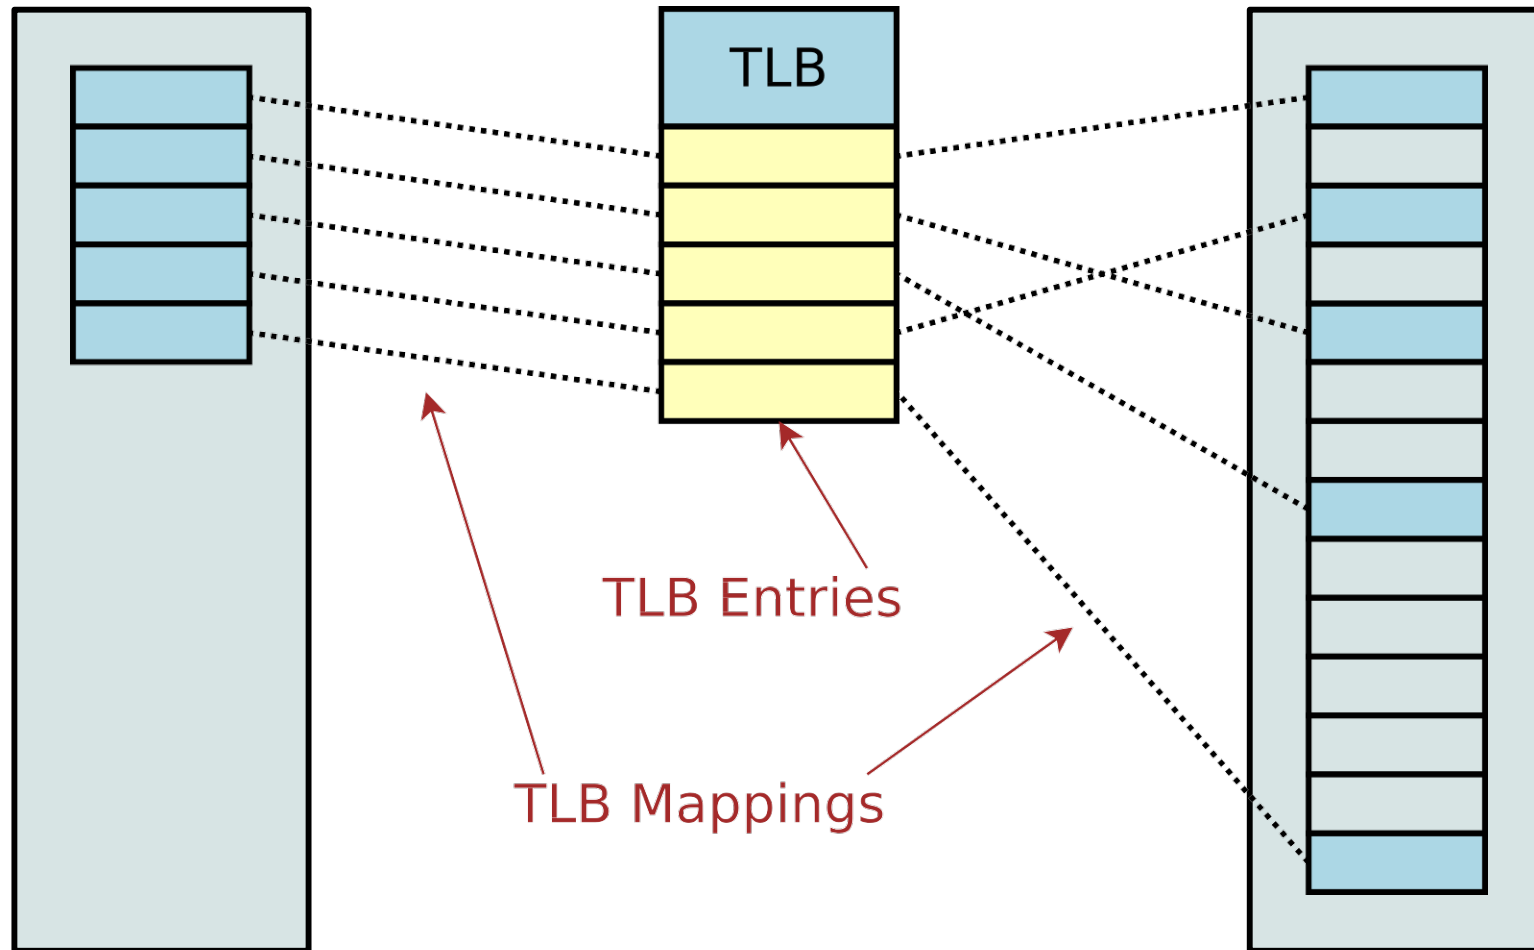


Basic Page Table Mappings

User Virtual Address Space

Page Table

Physical Address Space




Basic Page Table Mappings

User Virtual Address Space

Page Table

Physical Address Space



Page Table Mappings

- Mappings to virtually contiguous regions do not have to be physically contiguous.
 - This makes memory easier to allocate.
 - Almost all user space code does not need physically contiguous memory.

Multiple Processes


- Each process has its own mapping.
 - The same virtual addresses in different processes point to different physical addresses in other processes

Multiple Processes - Process 1

User Virtual Address Space

Page Table

Physical Address Space




Multiple Processes - Process 2

User Virtual Address Space

Page Table

Physical Address Space



Shared Memory

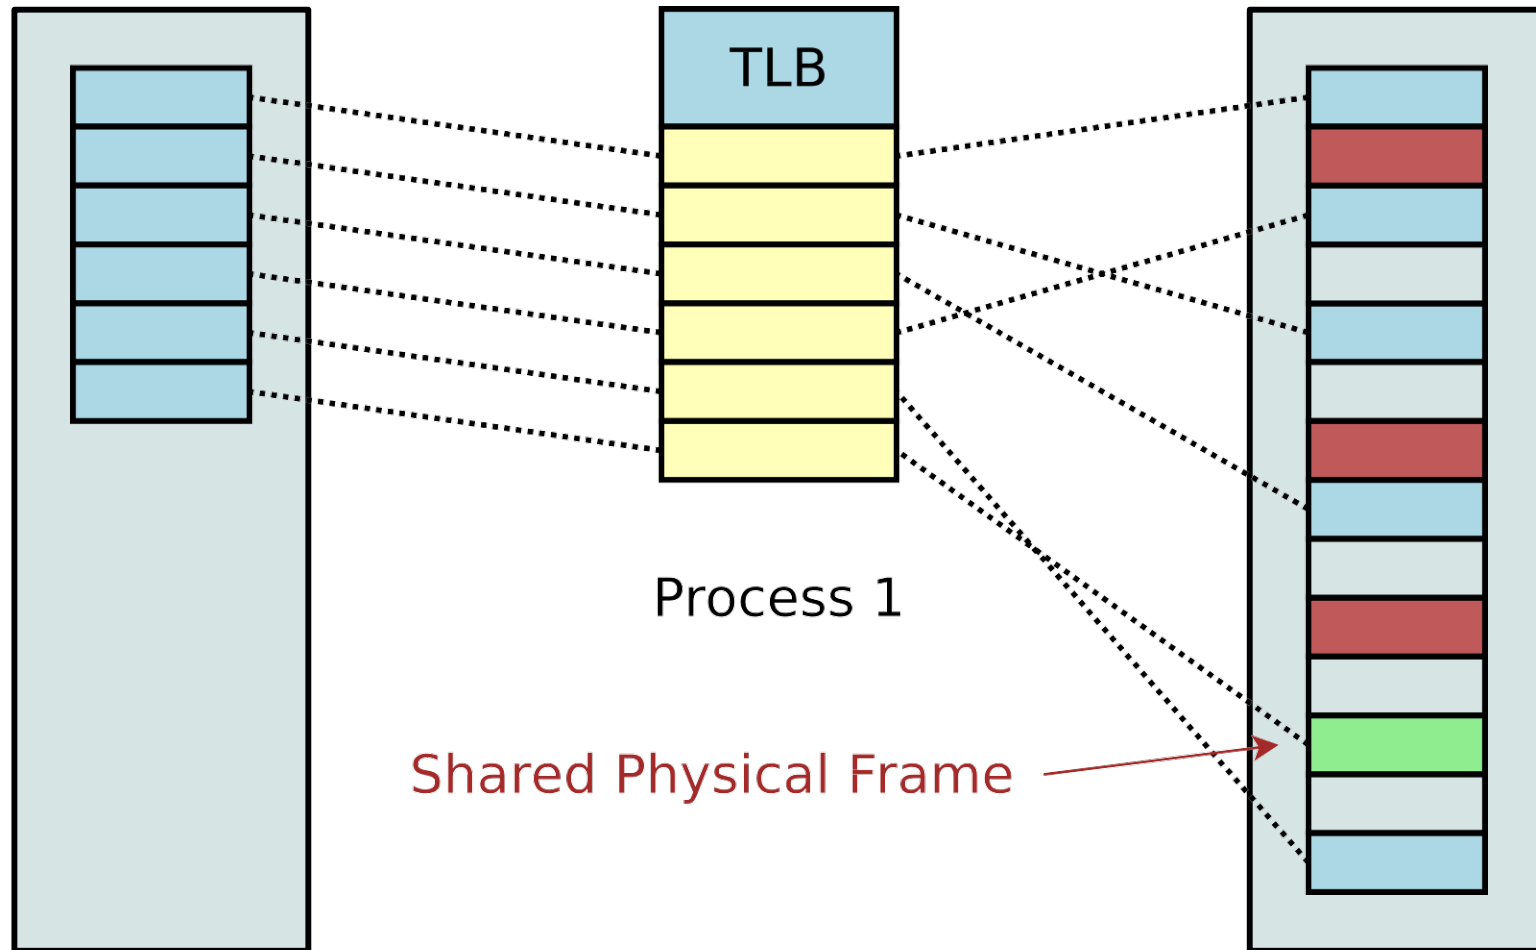
- Shared memory is easily implemented with an MMU.
 - Simply map the same physical frame into two different processes.
 - The virtual addresses need not be the same.
 - If pointers to values inside a shared memory region are used, it might be important for them to have the same virtual addresses, though.

Shared Memory - Process 1

User Virtual Address Space

Page Table

Physical Address Space

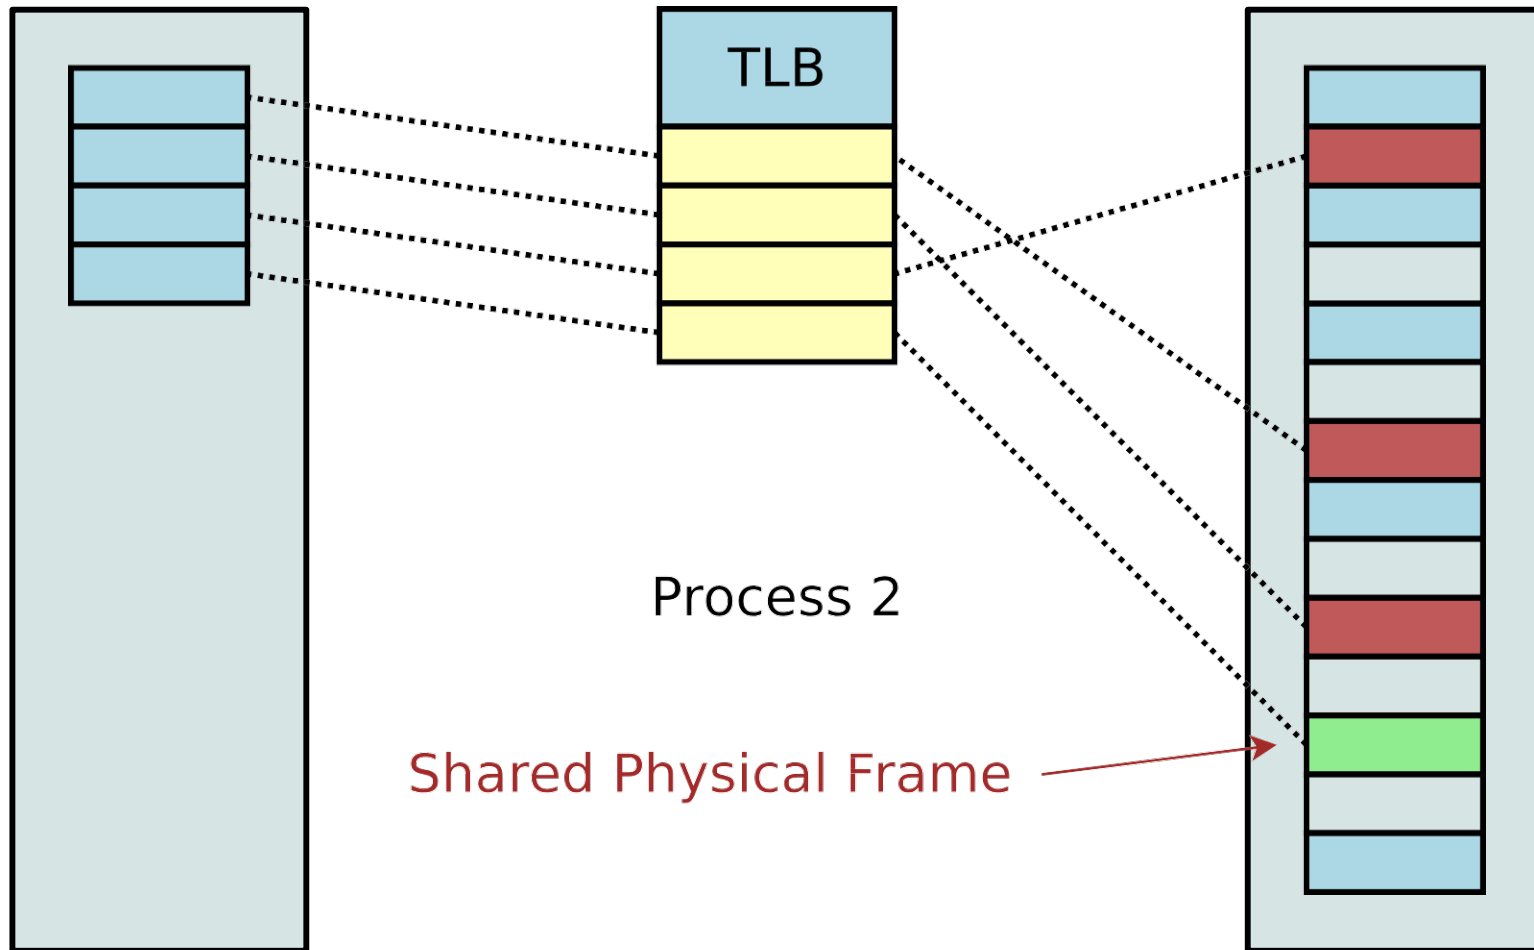


Shared Memory - Process 2

User Virtual Address Space

Page Table


Physical Address Space



Shared Memory

- Note in the previous example, the shared memory region was mapped to different virtual addresses in each process
- The `mmap()` system call allows the user space process to request a virtual address to map the shared memory region
 - The kernel may not be able to grant a mapping at this address, causing `mmap()` to return failure.

Memory Management Unit



The MMU

- The Memory Management Unit (MMU) is a hardware component which manages virtual address mappings
 - Maps virtual addresses to physical addresses
- The MMU operates on basic units of memory called pages
 - Page size varies by architecture
 - Some architectures have configurable page sizes

The MMU

- Common page sizes:
 - ARM – 4k
 - ARM64 – 4k or 64k
 - MIPS – Widely Configurable
 - x86 – 4k
 - *Architectures which are configurable are configured at kernel build time.*

The MMU

- Terminology
 - A **page** is a unit of memory sized and aligned at the page size.
 - A **frame**, or page frame, refers to a page-sized and page-aligned physical memory block.
 - *A page is somewhat abstract, where a frame is concrete*
 - *In the kernel, the abbreviation **pfn**, for **page frame number**, is often used to refer to refer to physical page frames*

The MMU

- The MMU operates in pages
 - The MMU maps physical frames to virtual addresses.
 - The TLB holds the entries of the mapping
 - Virtual address
 - Physical address
 - Permissions
 - A memory map for a process will contain many mappings

Page Faults

- When a process accesses a region of memory that is not mapped, the MMU will generate a page fault exception.
 - The kernel handles page fault exceptions regularly as part of its memory management design.

Lazy Allocation

- The kernel will not allocate pages requested by a process immediately.
 - The kernel will wait until those pages are actually used.
 - This is called lazy allocation and is a performance optimization.
 - For memory that doesn't get used, physical frame allocation never has to happen!

Lazy Allocation

- Process
 - When memory is requested, the kernel simply creates a record of the request, and then returns (quickly) to the process, without updating the TLB.
 - When that newly-allocated memory is touched, the CPU will generate a page fault, because the CPU doesn't know about the mapping

Lazy Allocation


- Process (cont)
 - In the page fault handler, the kernel determines that the mapping is valid (from the kernel's point of view).
 - The kernel updates the page table with the new mapping
 - The kernel returns from the exception handler and the user space program resumes.

Page Faults

- How The Kernel Manages Your Memory


- <https://manybutfinite.com/post/how-the-kernel-manages-your-memory/>

1. Program calls `brk()` to grow its heap




2. `brk()` enlarges heap VMA.


New pages are **not** mapped onto physical memory.



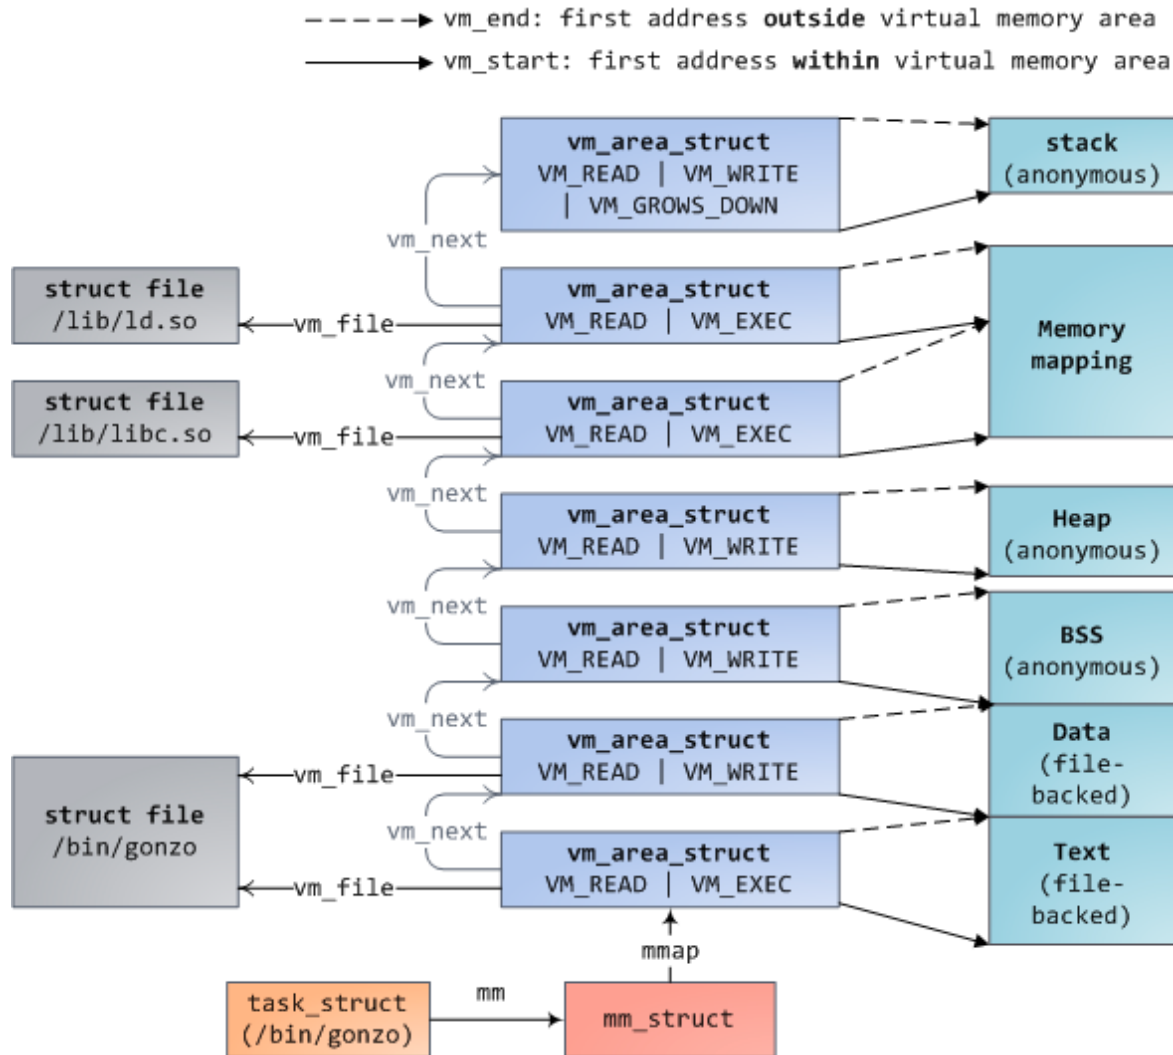
3. Program tries to access new memory.
Processor page faults.



4. Kernel assigns page frame to process,
creates PTE, resumes execution. Program is
unaware anything happened.



Page Faults



```

305 struct vm_area_struct {
306     /* The first cache line has the info for VMA tree walking. */
307
308     unsigned long vm_start;          /* Our start address within vm_mm. */
309     unsigned long vm_end;           /* The first byte after our end address
310                                     within vm mm. */

```

Lazy Allocation

- In a lazy allocation case, the user space program is never aware that the page fault happened.
 - The page fault can only be detected in the time that needs to be handled.
- For processes that are time-sensitive, data can be pre-faulted, or simply touched, at the start of execution.
 - Also see `mlock()` and `mlockall()` for pre-faulting.

Page Tables

- The entries in the TLB are a limited resource.
- Far more mappings can be made than can exist in the TLB at one time.
- The kernel must keep track of all of the mappings all of the time.
- The kernel stores all this information in the page tables.

Swapping

- When memory allocation is high, the kernel may swap some frames to disk to free up RAM.
 - Having an MMU makes this possible.
- The kernel can copy a frame to disk and remove its TLB entry.
- The frame can be re-used by another process.

Swapping

- When the frame is needed again, the CPU will generate a page fault (because the address is not in the TLB)
- The kernel can then, at page fault time:
 - Put the process to sleep
 - Copy the frame from the disk into an unused frame in RAM
 - Fix the page table entry
 - Wake the process

Swapping

- Note that when the page is restored to RAM, it's **not** necessarily restored to the same physical frame where it originally was located.
- The MMU will use the same virtual address though, so the user space program will not know the difference.
 - This is why user space memory cannot typically be used for DMA.

Takeaway

- Virtual memory in Linux
 - Kernel logical address
 - Kernel virtual address
 - User virtual address
- Contiguous vs non-contiguous
 - Virtually, physically
- Linux data struct
 - `mm_struct`, `switch_mm`, `mm_struct.pgd`
- Lazy allocation
 - Page fault