

File System Interface



Operating Systems
Wenbo Shen

Review

- Mass storage
- Disk structure
- Disk scheduling
 - FCFS, SSTF, SCAN, C-SCAN, LOOK, C-LOOK
- IO hardware
- IO access
 - polling, interrupt
- Device types
- Kernel IO subsystem

Outline

- File concept
- Access methods
- Directory structure
- Protection

How to use storage

- Now we have mass storage and IO
 - But how to use?
 - Think as a computer scientist in 1950s
- Use disk directly




How to use storage

- Now we have mass storage and IO
 - But how to use?
- **File System**
 - File system vs. Disk
 - File system presents **abstraction** of disk
 - File → Track/sector
 - To user process
 - File system provides coherent view of a group of files
 - File: a contiguous block of bytes (Unix)
 - File system provides protection

Abstraction

- CPU is abstracted to ____.
- Memory is abstracted to ____.
- Storage is abstracted to ____.



Inquiry-based Learning

- How to use file system?
 - How to use file?
 - How to use directory?
- How to implement file system?
 - How to implement file?
 - How to implement directory?

File Concept

- **File** is a contiguous logical space for storing information
 - database, audio, video, web pages...
- There are different types of file:
 - data: character, binary, and application-specific
 - program
 - special one: proc file system - use file-system interface to retrieve system information

File Attributes

- **Name** – only information kept in **human-readable form**
- **Identifier** – unique tag (number) identifies file **within file system**
- **Type** – needed for systems that support different types
- **Location** – pointer to **file location on device**
- **Size** – current file size
- **Protection** – controls who can do reading, writing, executing
- **Time, date, and user identification** – data for protection, security, and usage monitoring
- Information about files are kept in the directory structure, which is maintained on the disk
- Many variations, including **extended file attributes** such as file checksum

File info on Linux

```
wenbo@wenbo-Virtual:~/os-course$ file main.c
main.c: C source, ASCII text
wenbo@wenbo-Virtual:~/os-course$ file a.out
a.out: ELF 64-bit LSB pie executable, x86-64, version 1 (SYSV), dynamically
linked, interpreter /lib64/ld-linux-x86-64.so.2, BuildID[sha1]=e8385
a23e08804a95f2d49dee91d6de43407ffdc, for GNU/Linux 3.2.0, not stripped
wenbo@wenbo-Virtual:~/os-course$ stat a.out
  File: a.out
  Size: 15960          Blocks: 32          IO Block: 4096   regular file
Device: 803h/2051d    Inode: 3670041      Links: 1
Access: (0775/-rwxrwxr-x)  Uid: ( 1000/   wenbo)   Gid: ( 1000/   wenbo)
Access: 2022-12-15 11:29:38.536299736 +0800
Modify: 2022-11-24 14:35:14.406255682 +0800
Change: 2022-11-24 14:35:14.406255682 +0800
 Birth: 2022-11-24 14:35:14.390247501 +0800
```

File Operations

- OS provides file operations to
 - **create**:
 - space in the file system should be found
 - an entry must be allocated in the directory
 - **open**: most operations need to file to be opened first
 - return a handler for other operations
 - **read/write**: need to maintain a **pointer**
 - reposition within file – **seek**
 - **close**
 - **delete**
 - Release file space
 - Hardlink: maintain a counter - delete the file until the last link is deleted
 - **truncate**: empty a file but maintains its attributes
- Other operations can be implemented using these ones
 - Copying: create and read/write

Open Files

- Several data are needed to manage open files:
 - Open-file table: tracks open files
 - File pointer: pointer to last read/write location, per process that has the file open
 - File-open count: counter of number of times a file is open – to allow removal of data from open-file table when last processes closes it
 - Disk location of the file: cache of data access information
 - Access rights: per-process access mode information

Open Files

- Some file systems provide file lock to mediate access to a file
- Two types of lock
 - **Shared lock** - multiple processes can acquire the lock concurrently
 - **Exclusive lock** - one process can acquire such a lock
- Two locking mechanisms
 - **mandatory lock**: access is denied depending on locks held and requested
 - **advisory lock**: processes can find status of locks and decide what to do

File Types

- as part of the file names - file extension
- magic number of the file - elf

file type	usual extension	function
executable	exe, com, bin or none	ready-to-run machine- language program
object	obj, o	compiled, machine language, not linked
source code	c, cc, java, pas, asm, a	source code in various languages
batch	bat, sh	commands to the command interpreter
text	txt, doc	textual data, documents
word processor	wp, tex, rtf, doc	various word-processor formats
library	lib, a, so, dll	libraries of routines for programmers
print or view	ps, pdf, jpg	ASCII or binary file in a format for printing or viewing
archive	arc, zip, tar	related files grouped into one file, sometimes com- pressed, for archiving or storage
multimedia	mpeg, mov, rm, mp3, avi	binary file containing audio or A/V information

File Structure


- A file can have different structures, determined by OS or program
 - **No structure:** a stream of bytes or words
 - Linux dumps
 - **Simple record structure**
 - Lines of records, fixed length or variable length
 - E.G., Database
 - **Complex structures**
 - E.G., Word document, relocatable program file
- Usually user programs are responsible for identifying file structure

Outline


- File concept
- Access methods
- Directory structure
- Protection

Access Methods

- Sequential access
 - a group of elements is access **in a predetermined order**
 - for some media types, the only access mode (e.g., **tape**)
- Direct access
 - access an element at an **arbitrary position** in a sequence in (roughly) **equal time**, independent of sequence size
 - it is possible to emulate random access in a tape, but access time varies
 - sometime called random access



Sequential-access File



Sequential Access on Direct-access File

sequential access	implementation for direct access
<i>reset</i>	<i>cp = 0;</i>
<i>read next</i>	<i>read cp;</i> <i>cp = cp + 1;</i>
<i>write next</i>	<i>write cp;</i> <i>cp = cp + 1;</i>

Other methods

- Based on direct-access method
- An index for the file points to blocks
 - Find a record in the file, first search the index and then use the pointer to access the block
 - We may use multiple layers of index


Outline

- File concept
- Access methods
- Directory structure
- Protection

Disk and file system


- Disk can be subdivided into **partitions**
 - partitions also known as **minidisks, slices**
 - different partitions can have different file systems
 - a partition containing file system is known as a **volume**
 - each volume tracks file system info in the volume's table of contents
 - a file system can be general purpose or special purpose
 - disk or partition can be used **raw** (without a file system)
 - applications such as database prefer raw disks

A Typical File-system Organization



Directory Structure

- Directory is a collection of nodes containing information about all files



Both the directory structure and the files reside on disk

Operations Performed on Directory


- Create a file: new files need to be created and added to directory
- delete a file: remove a file from directory
- List a directory: list all files in directory
- Search for a file: pattern matching
- Traverse the file system: access every directory and file within a directory
- ...

Directory Organization

- Organize directories to achieve
 - **Efficiency**: to locate a file quickly
 - **Naming**: organize the directory structure to be convenient to users
 - Two users can have same name for different files
 - The same file can have several different names
 - ...

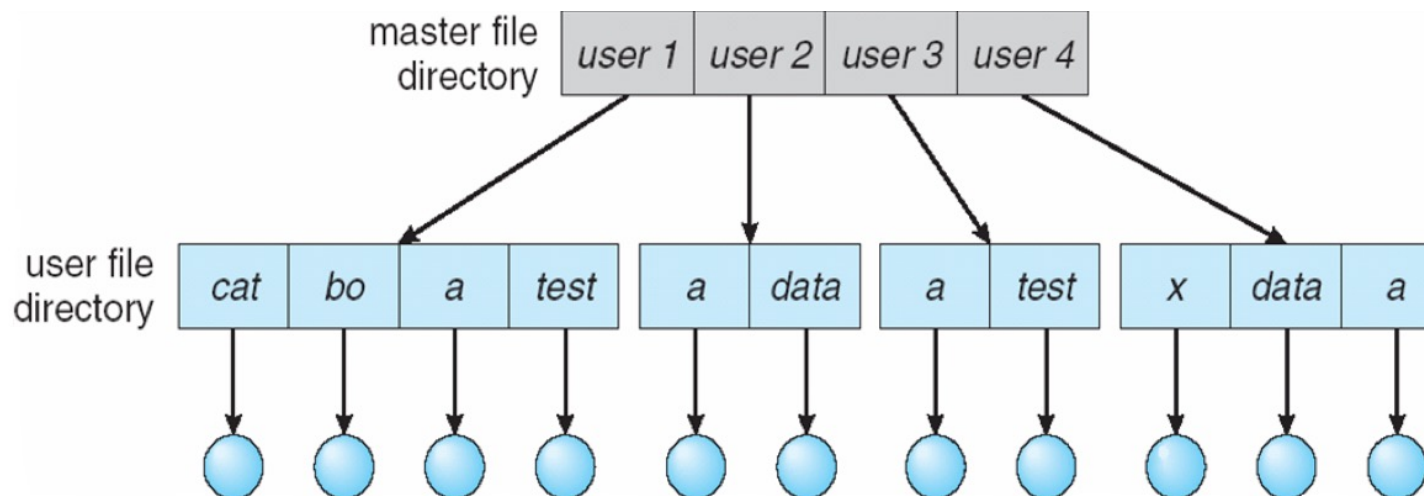
Single-Level Directory

- A single directory for all users
 - Naming problems and grouping problems
 - Two users want to have same file names



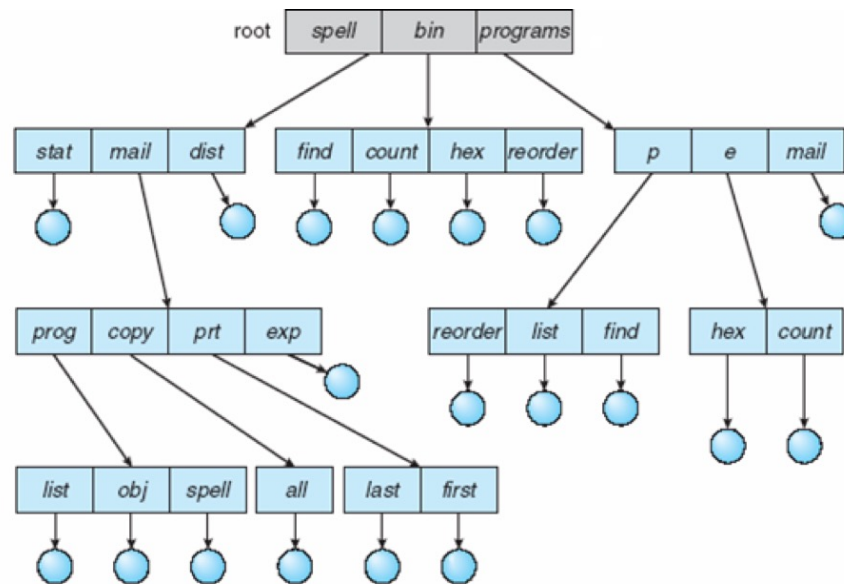
Two-Level Directory

- Separate directory for each user
 - Different user can have the same name for different files
 - Each user has his own user file directory (UFD), it is in the master file directory (MFD)
 - Efficient to search
 - How to share files between different users?
 - Path concept



Tree-Structured Directories


- Files organized into trees
 - efficient in searching, can group files, convenient naming



Tree-Structured Directories

- File can be accessed using **absolute** or **relative** path name
 - absolute path name: /home/alice/..
 - relative path is relative to the **current directory** (*pwd*)
- Creating a new file: *touch <file-name>*
- Delete a file: *rm <file-name>*
- Creating a new subdirectory
 - Example: if in current directory /mail

mkdir count



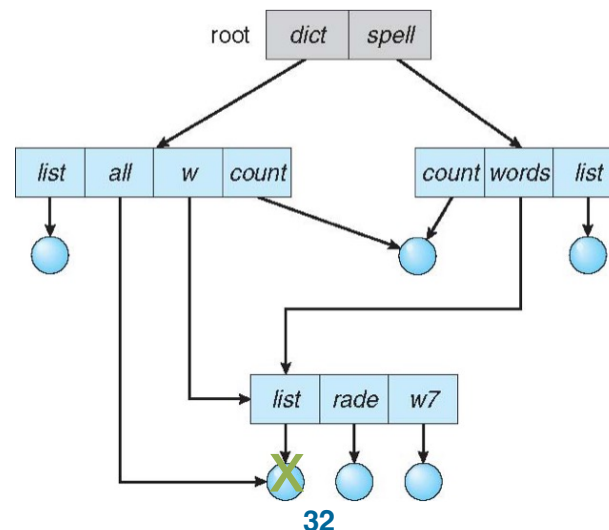
- e.g., if current directory is /mail, a *mkdir* command will create /mail/count
- How to share a file/directory? -> it's not allowed

Tree-Structured Directories

- Delete directory
 - If directory is empty, then it's easy to handle
 - If not
 - Option I: directory cannot be deleted, unless it's empty
 - Option II: delete all the files, directories and sub-directories
 - `sudo rm -rf /`



Acyclic-Graph Directories

- Organize directories into acyclic-graphs
 - allow links to a directory entry/files for **aliasing** (no longer a tree)
- Dangling pointer problem:
 - e.g., if delete **file** /dict/all, /dict/w/list and /spell/words/list are dangling pointers
 - Solution: **back pointers/reference counter**
 - Back pointers record all the pointers to the entity, a variable size record
 - Or count # of links to it and only (physically) delete it when counter is zero



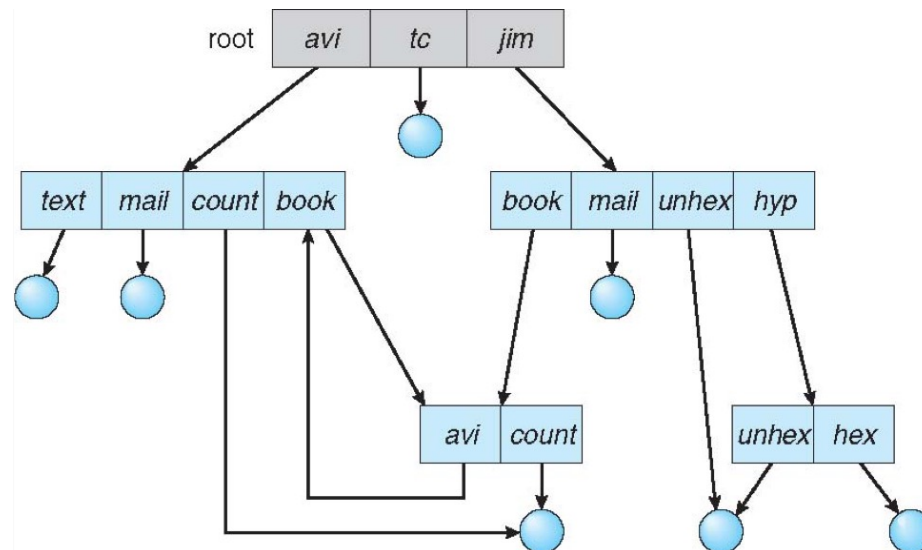
Acyclic-Graph Directories

- Share files
 - Hardlink
 - Reference count
 - Softlink



General Graph Directory

- Allowing arbitrary links may generate cycles in the directory structure
- Solution
 - allow cycles, but use **garbage collection** to reclaim disk spaces
 - every time a new link is added use a **cycle detection** algorithm






File System Mounting

- A file system must be **mounted** before it can be accessed
 - mounting links a file system to the system, usually forms a **single name space**
 - the location of the file system being mounted is call the **mount point**
 - a mounted file system makes the old directory at the mount point **invisible**

File System Mounting

- **a**: existing file system
- **b**: an unmounted partition
- **c**: the partition mounted at **/users**



Mounting a file system

```
os@os:~/temp/foo$ mount
sysfs on /sys type sysfs (rw,nosuid,nodev,noexec,relatime)
proc on /proc type proc (rw,nosuid,nodev,noexec,relatime)
udev on /dev type devtmpfs (rw,nosuid,relatime,size=475364k,nr_inodes=118841,mode=755)
devpts on /dev/pts type devpts (rw,nosuid,noexec,relatime,gid=5,mode=620,ptmxmode=000)
tmpfs on /run type tmpfs (rw,nosuid,noexec,relatime,size=100384k,mode=755)
/dev/sda1 on / type ext4 (rw,relatime,errors=remount-ro,data=ordered)
securityfs on /sys/kernel/security type securityfs (rw,nosuid,nodev,noexec,relatime)
tmpfs on /dev/shm type tmpfs (rw,nosuid,nodev)
tmpfs on /run/lock type tmpfs (rw,nosuid,nodev,noexec,relatime,size=5120k)
```

File Sharing

- Sharing of files on multi-user systems is desirable
 - Sharing must be done through a protection scheme
 - **User IDs** identify users, allowing protections to be per-user
 - **Group IDs** allow users to be in groups, permitting group access rights
- On distributed systems, files may be shared across a network
 - Network File System (NFS) is a common distributed file-sharing method

Remote File Sharing

- **Use networking** to allow file system access between systems
 - manually via programs like FTP
 - automatically, seamlessly using distributed file systems
 - semi automatically via the world wide web
- Client-server model allows clients to **mount** remote FS from servers
 - a server can serve multiple clients
 - client and user-on-client identification is complicated
 - server cannot assume the client is trusted
 - standard OS file calls are translated into remote calls
 - **NFS** is standard UNIX file sharing protocol, **CIFS** is standard for Windows

Outline

- File concept
- Access methods
- Directory structure
- Protection

Protection

- File owner/creator should be able to control
 - what can be done
 - by whom
- Types of access
 - read, write, append
 - execute
 - delete
 - list

ACL


- Assign each file and directory with an access control list (ACL)
- Advantages: fine-grained control
- Disadvantages
 - How to construct the list
 - How to store the list in directory

Unix Access Control

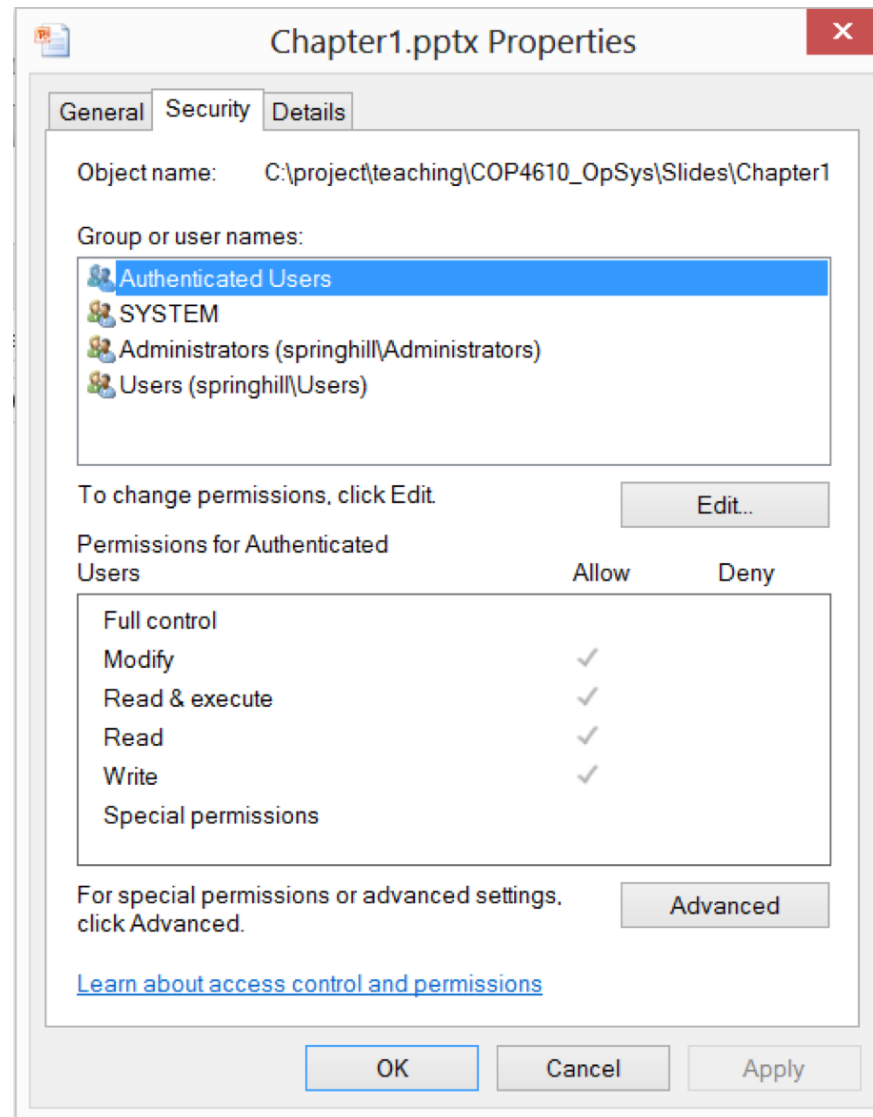
- Three modes of access: **read**, **write**, **execute** (encoded in three bits)
- Three classes of users: **owner**, **group**, and **others**

		RWX
a) owner access:	7	1 1 1
b) group access:	6	1 1 0
c) others access:	1	0 0 1

- To grant access to users, create a group and change its access mode
 - in Linux, use **chmod** and **chgrp**



Windows 8 File Access-Control



A Sample UNIX Directory Listing

-rw-rw-r--	1	pbg	staff	31200	Sep 3 08:30	intro.ps
drwx-----	5	pbg	staff	512	Jul 8 09:33	private/
drwxrwxr-x	2	pbg	staff	512	Jul 8 09:35	doc/
drwxrwx---	2	pbg	student	512	Aug 3 14:13	student-proj/
-rw-r--r--	1	pbg	staff	9423	Feb 24 2003	program.c
-rwxr-xr-x	1	pbg	staff	20471	Feb 24 2003	program
drwx--x--x	4	pbg	faculty	512	Jul 31 10:31	lib/
drwx-----	3	pbg	staff	1024	Aug 29 06:52	mail/
drwxrwxrwx	3	pbg	staff	512	Jul 8 09:35	test/

ACL in practice

```
os@os:~/os2018fall/test$ ls -l
total 0
-rw-rw-r-- 1 os os 0 Dec 18 23:21 testacl
os@os:~/os2018fall/test$ getfacl testacl
# file: testacl
# owner: os
# group: os
user::rw-
group::rw-
other::r--

os@os:~/os2018fall/test$ setfacl -m u:test:rw testacl
os@os:~/os2018fall/test$ getfacl testacl
# file: testacl
# owner: os
# group: os
user::rw-
user:test:rw-
group::rw-
mask::rw-
other::r--
```

Takeaway

- File system
- File operations
 - Create, open, read/write, close
- File type
- File structure
- File access
- Directory structure
 - Single level, two-level, tree, acyclic-graph, general graph
- Protection
 - ACL